


**RTB**  
**HOUSE =**



**Real-Time Data  
Processing  
at RTB House**

**How we have grown 10x within 2 years**  
**Bartosz Łoś, 2019**

# AGENDA

- our RTB platform

# AGENDA

- our RTB platform
- the previous iterations: three different architectures

# AGENDA

- our RTB platform
- the previous iterations: three different architectures
- the fourth iteration: multi-dc architecture

# AGENDA

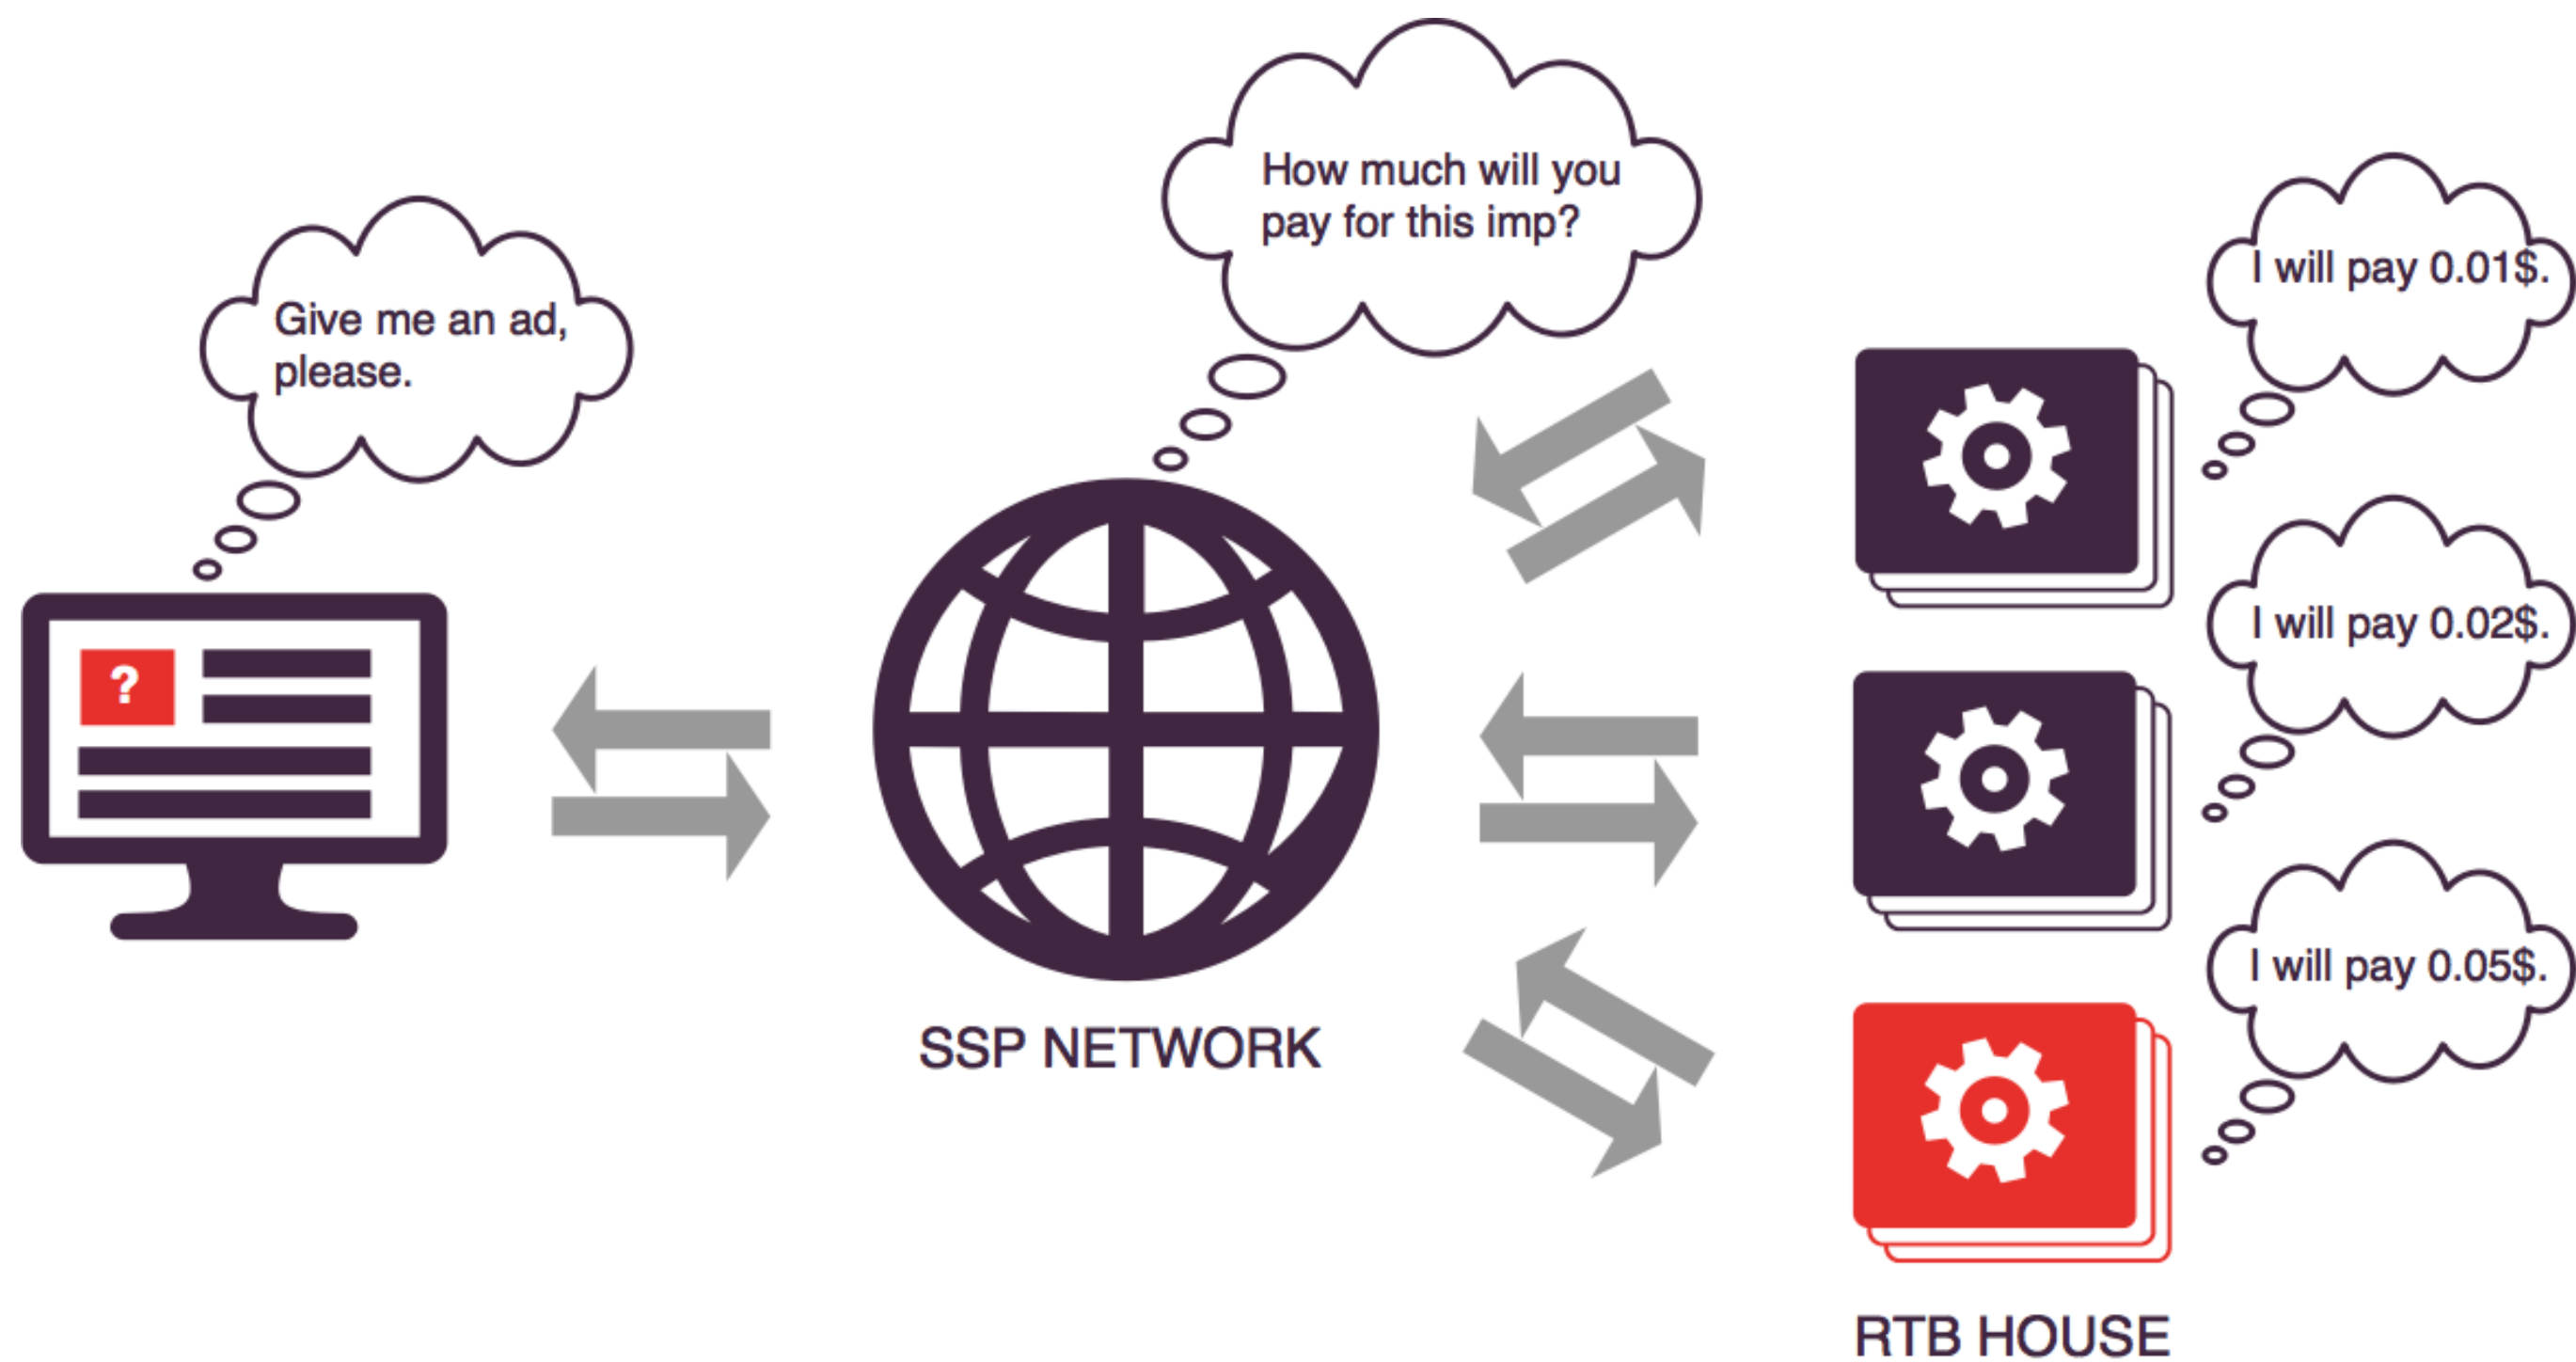
- our RTB platform
- the previous iterations: three different architectures
- the fourth iteration: multi-dc architecture
- our use cases: requirements and processing patterns

# AGENDA

- our RTB platform
- the previous iterations: three different architectures
- the fourth iteration: multi-dc architecture
- our use cases: requirements and processing patterns
- kafka workers

# OUR RTB PLATFORM

## OUR RTB PLATFORM: **THE CONTEXT**



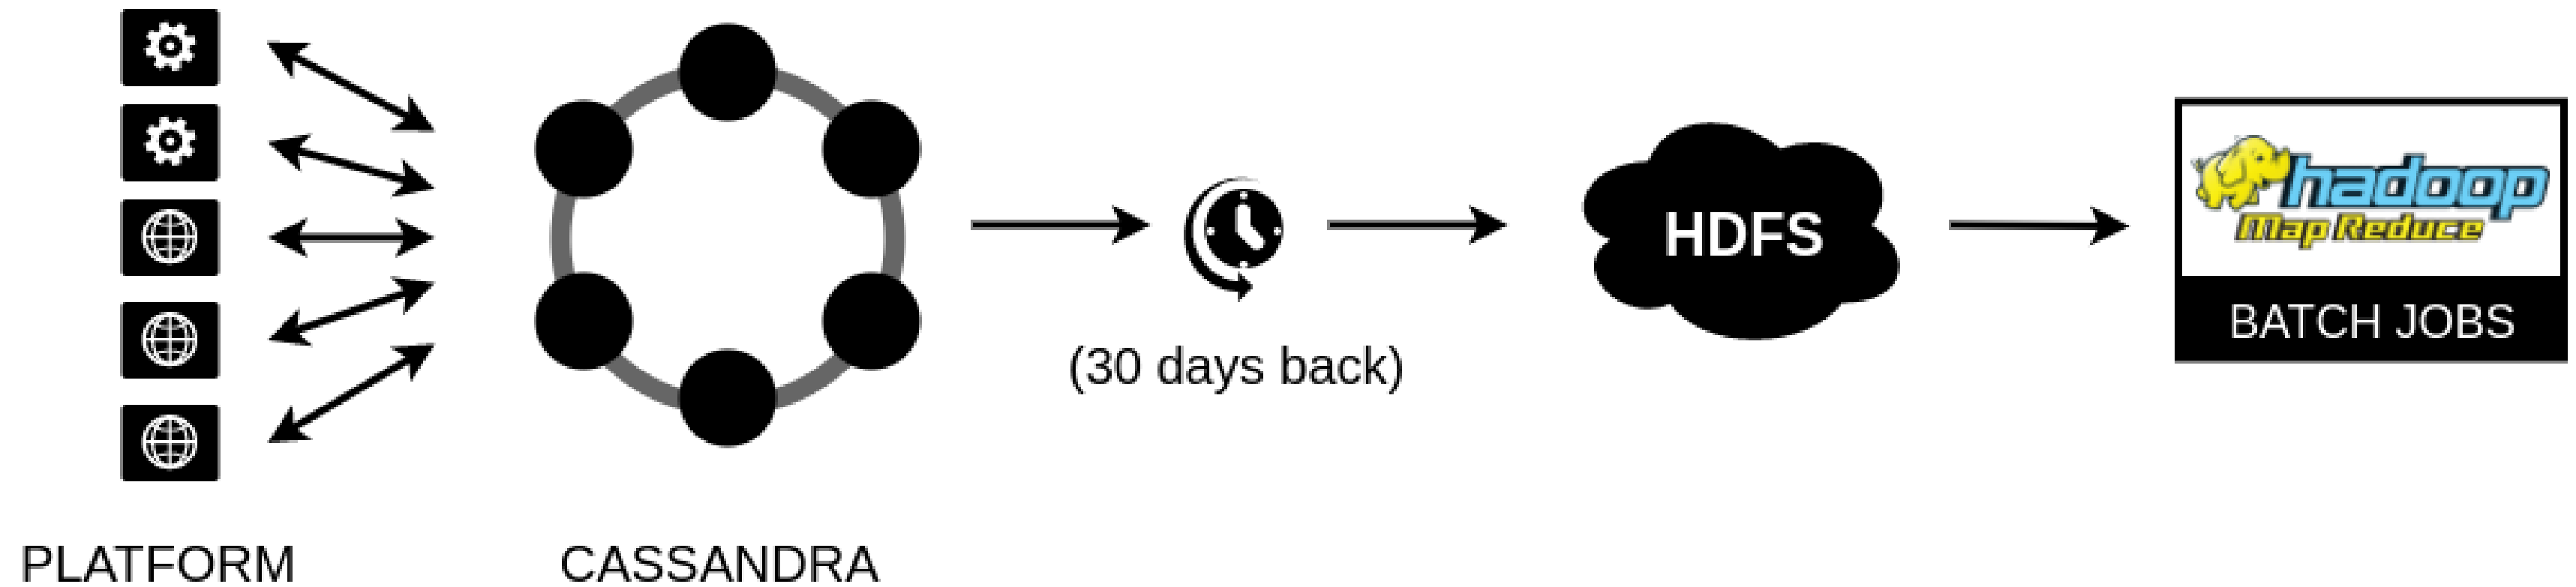


## OUR RTB PLATFORM: **THE CONTEXT**

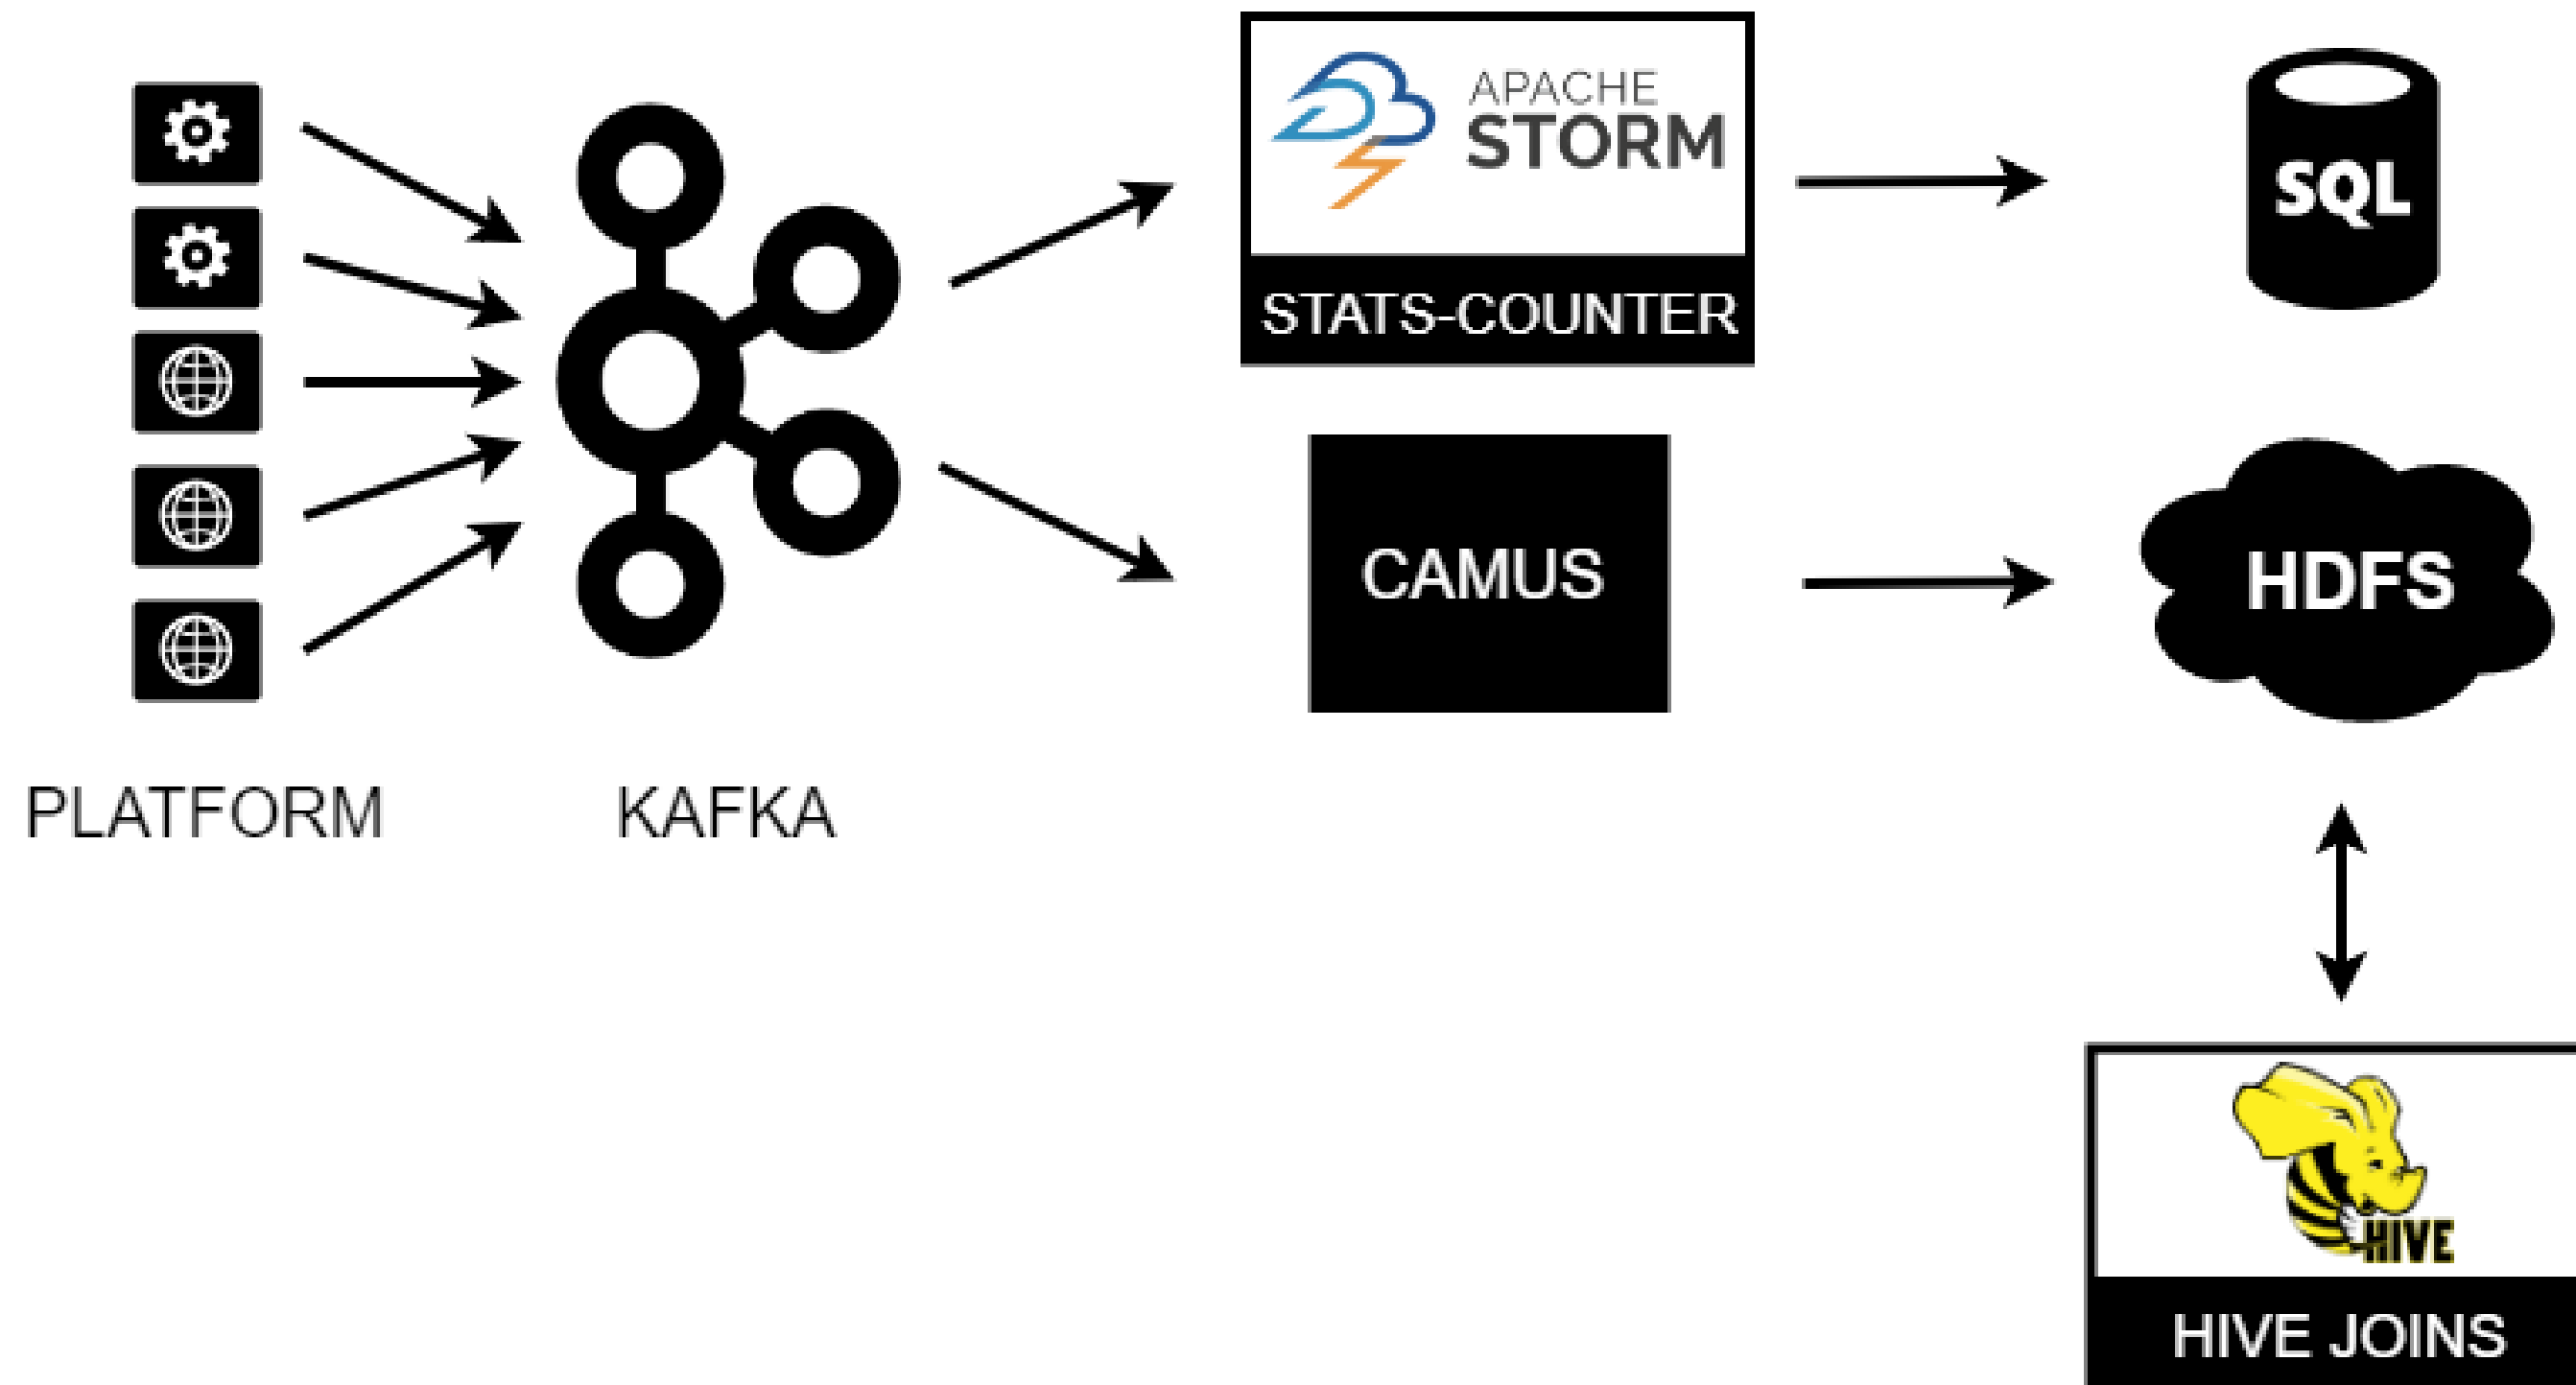


# THE PREVIOUS ITERATIONS

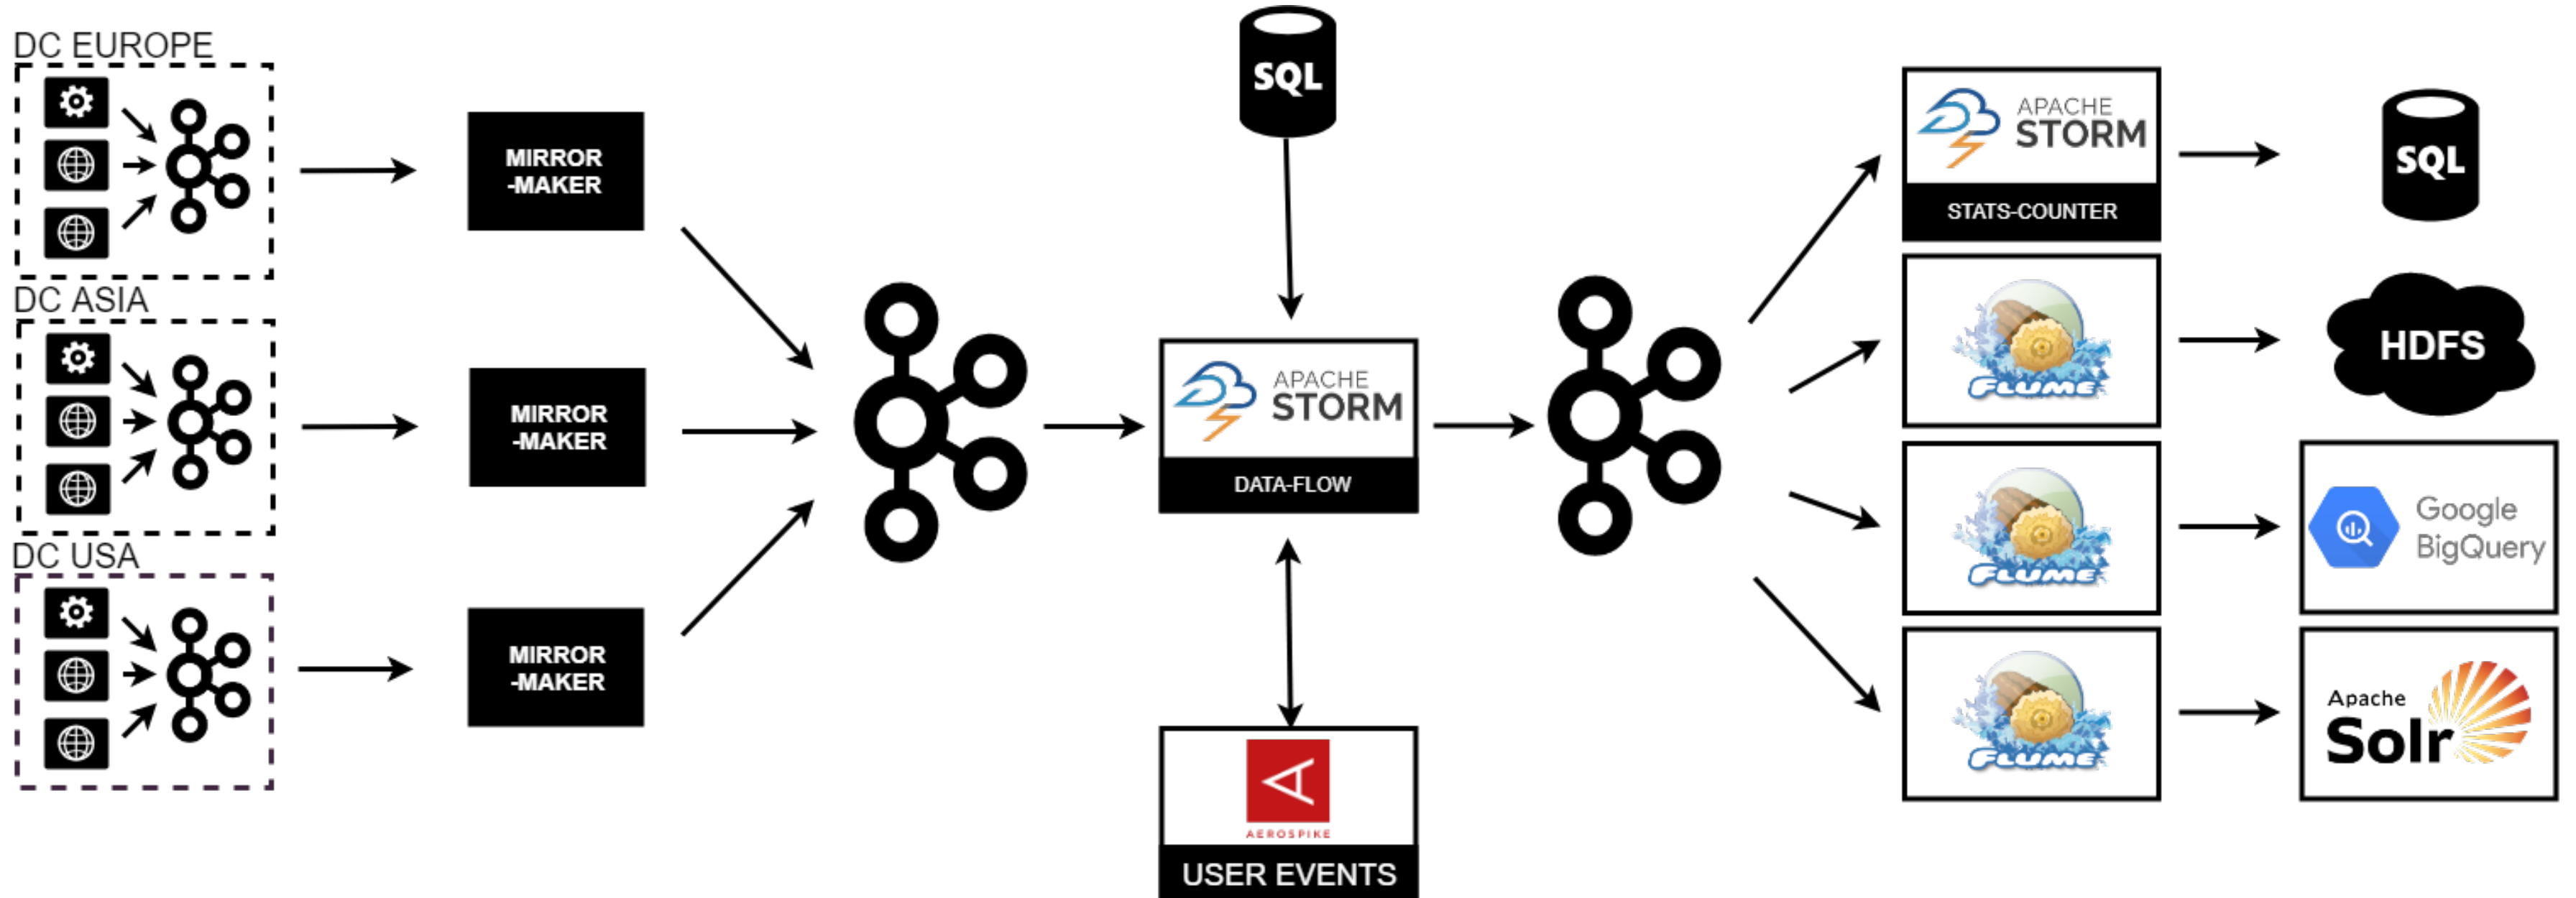
## THE 1ST ITERATION: **MUTABLE IMPRESSIONS**



## THE 2ND ITERATION: **LAMBDA ARCHITECTURE**



## THE 3RD ITERATION: IMMUTABLE STREAMS OF EVENTS



# THE FOURTH ITERATION: MULTI-DC

## THE 4TH ITERATION: **MAIN CHANGES**

- 10x larger scale:
  - from 350K to 3.5M bid requests/s within 2 years

## THE 4TH ITERATION: **MAIN CHANGES**

- 10x larger scale:
  - from 350K to 3.5M bid requests/s within 2 years
- full multi-dc architecture:
  - synchronization of user profiles
  - merging streams of events



## THE 4TH ITERATION: **MAIN CHANGES**

- 10x larger scale:
  - from 350K to 3.5M bid requests/s within 2 years
- full multi-dc architecture:
  - synchronization of user profiles
  - merging streams of events
- fixed partitioning in all DCs:
  - parallelism, merging, end-to-end lag

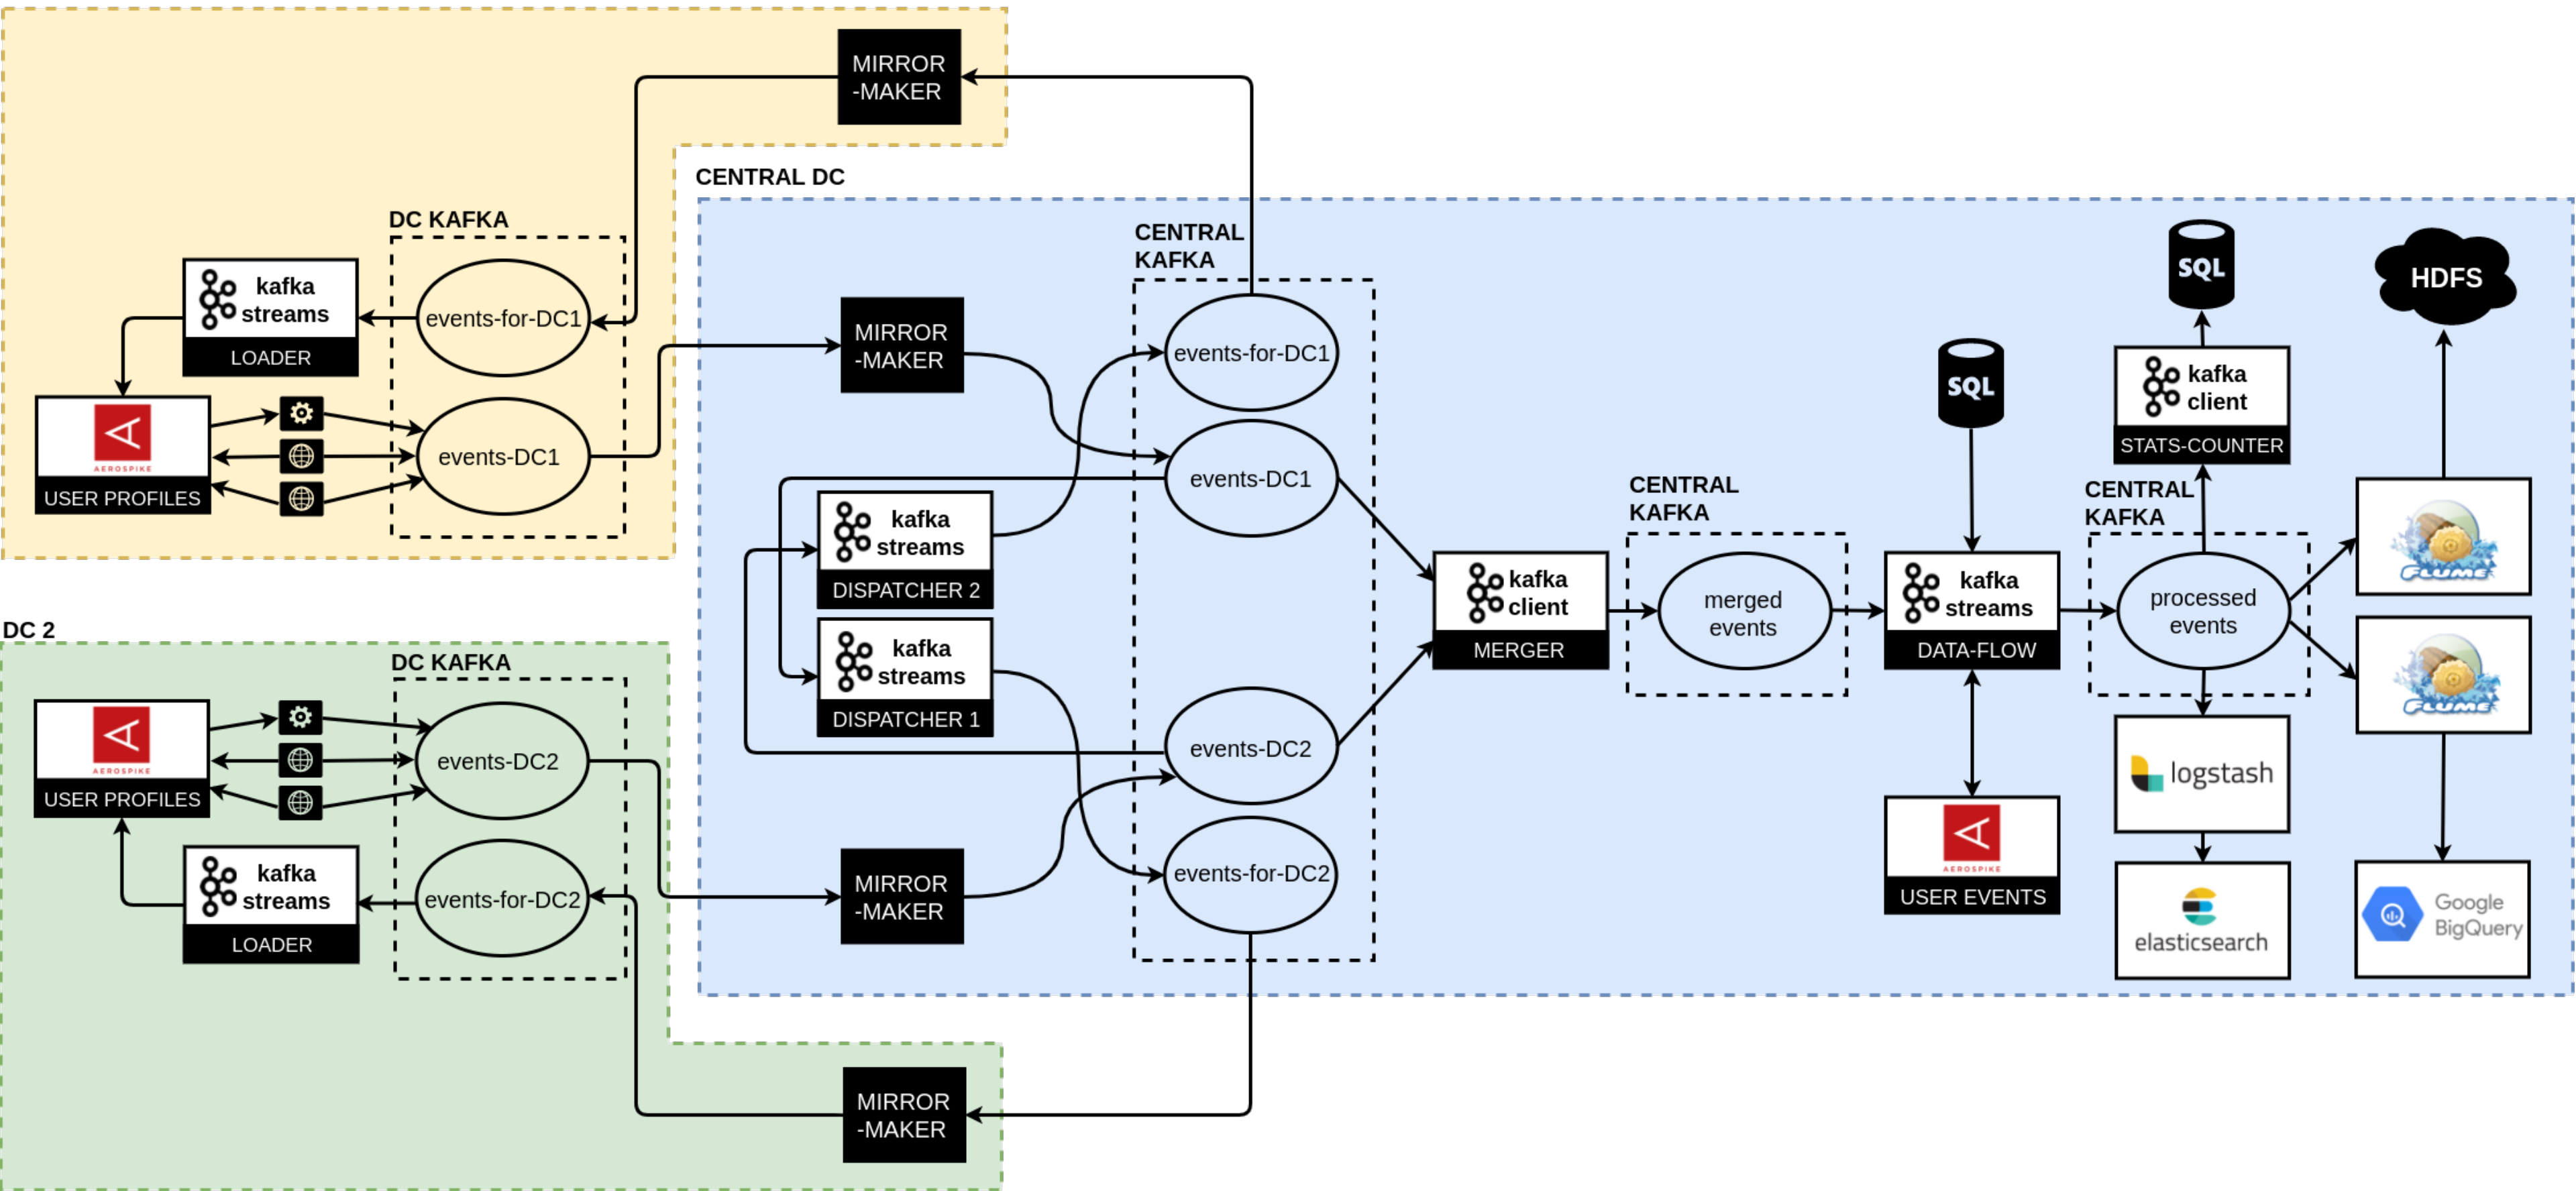
## THE 4TH ITERATION: **MAIN CHANGES**

- 10x larger scale:
  - from 350K to 3.5M bid requests/s within 2 years
- full multi-dc architecture:
  - synchronization of user profiles
  - merging streams of events
- fixed partitioning in all DCs:
  - parallelism, merging, end-to-end lag
- end-to-end exactly-once processing:
  - at-least-once output semantics & deduplication

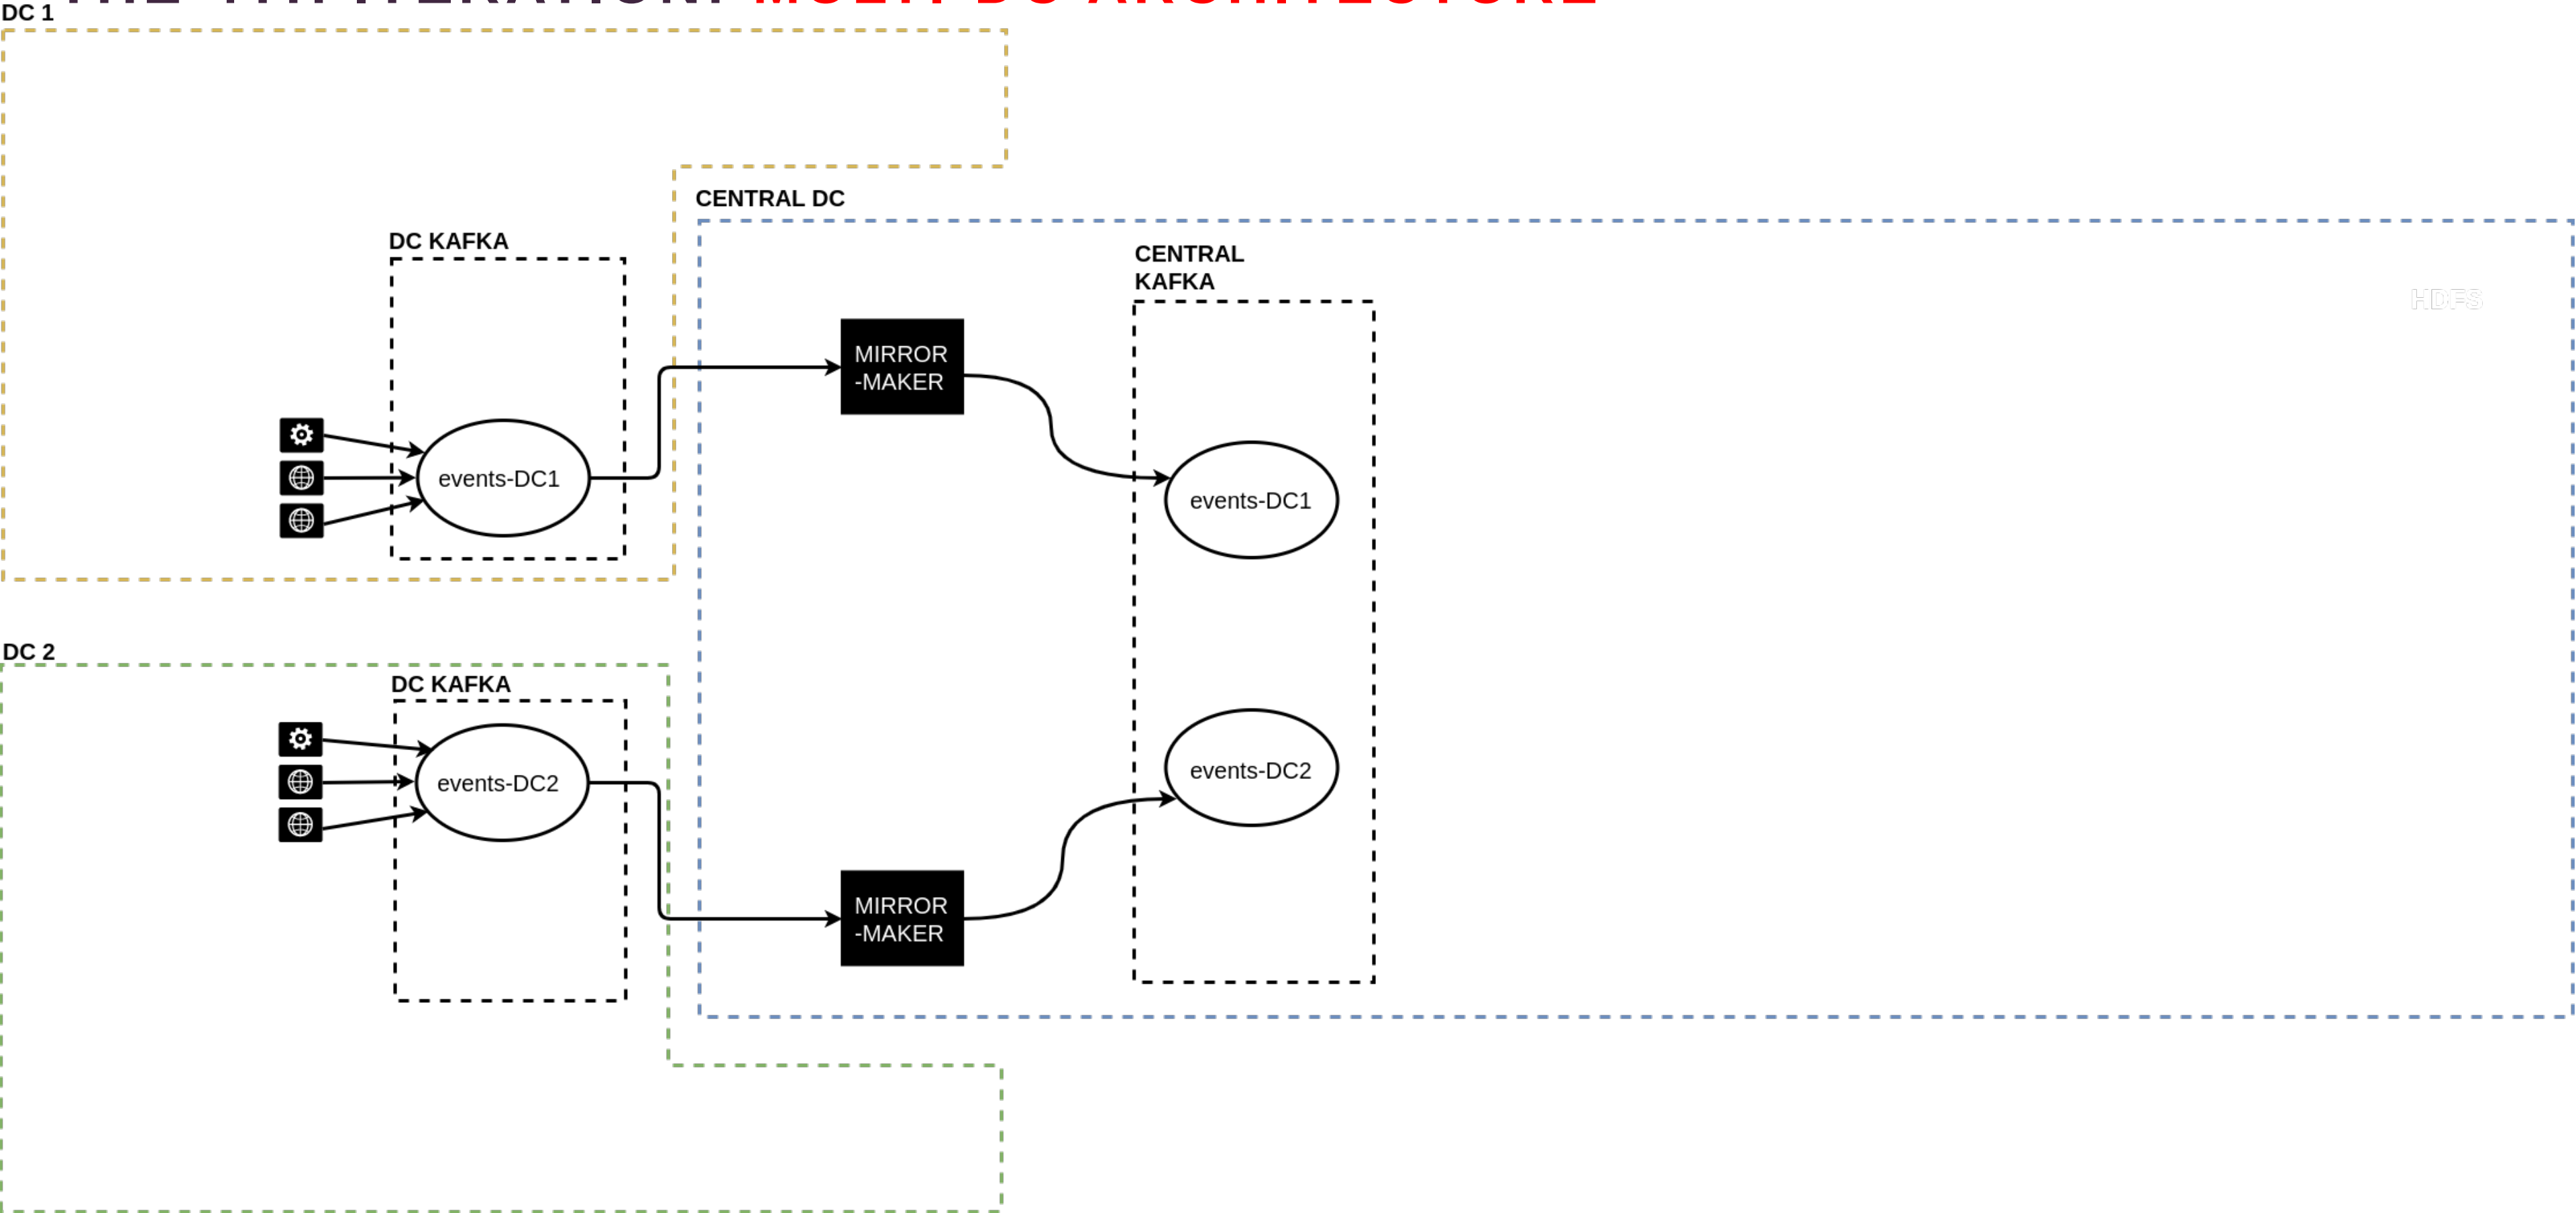
## THE 4TH ITERATION: **MAIN CHANGES**

- 10x larger scale:
  - from 350K to 3.5M bid requests/s within 2 years
- full multi-dc architecture:
  - synchronization of user profiles
  - merging streams of events
- fixed partitioning in all DCs:
  - parallelism, merging, end-to-end lag
- end-to-end exactly-once processing:
  - at-least-once output semantics & deduplication
- a few better components:
  - new stats-counter, new data-flow
  - logstash
  - merger, dispatcher & loader

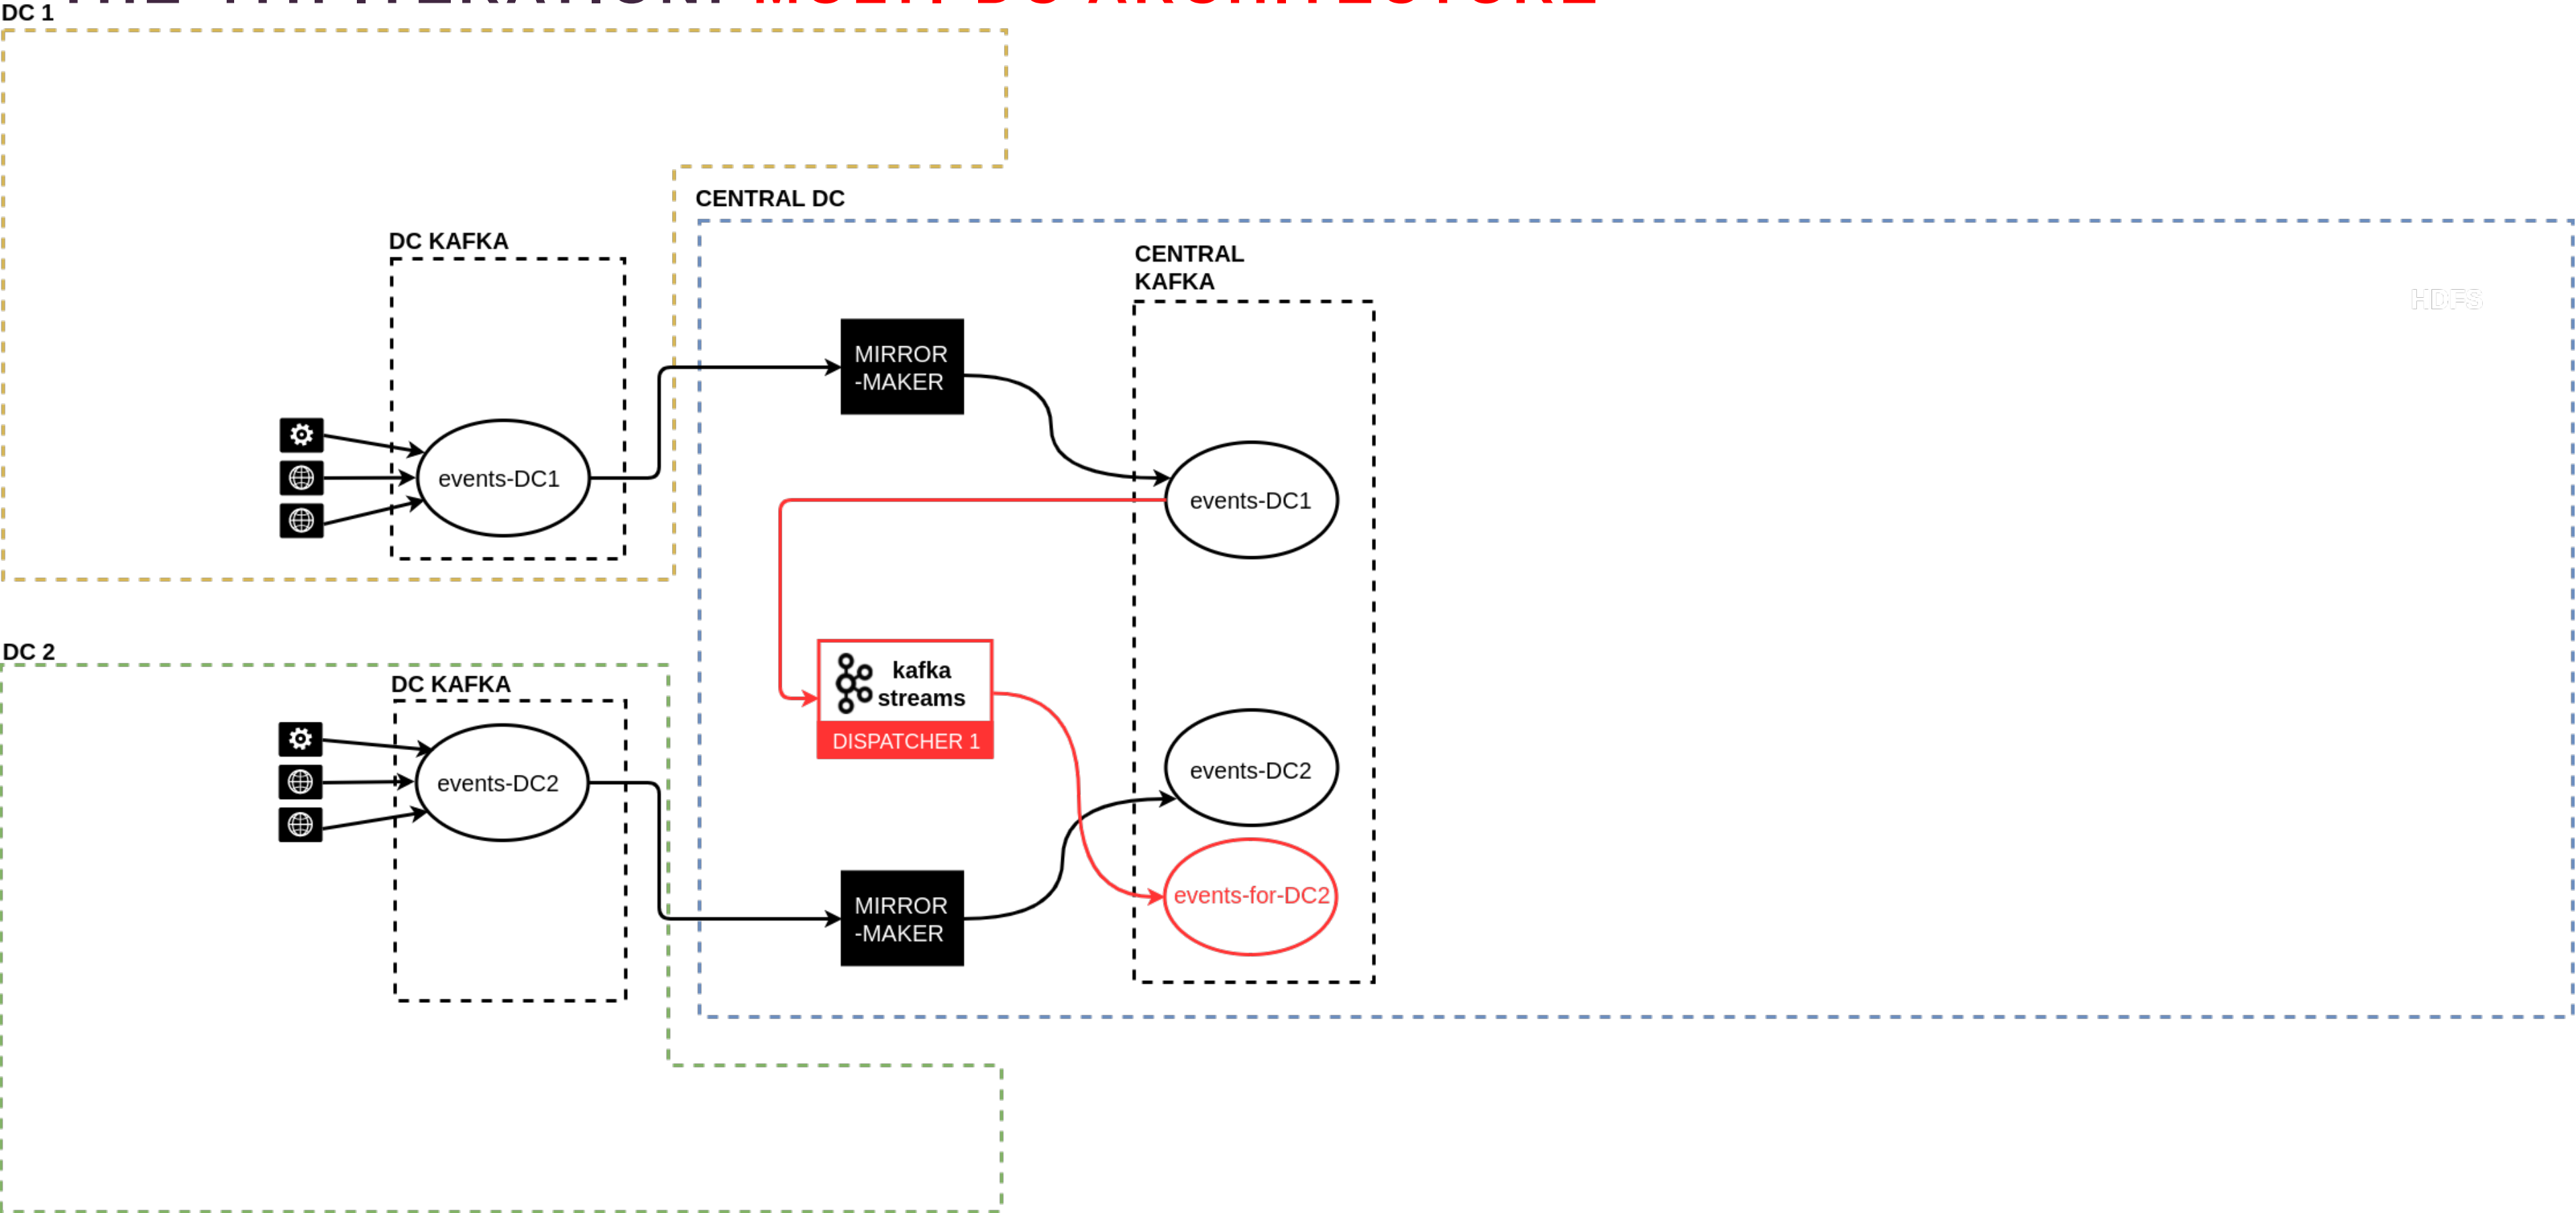
# THE 4TH ITERATION: MULTI-DC ARCHITECTURE



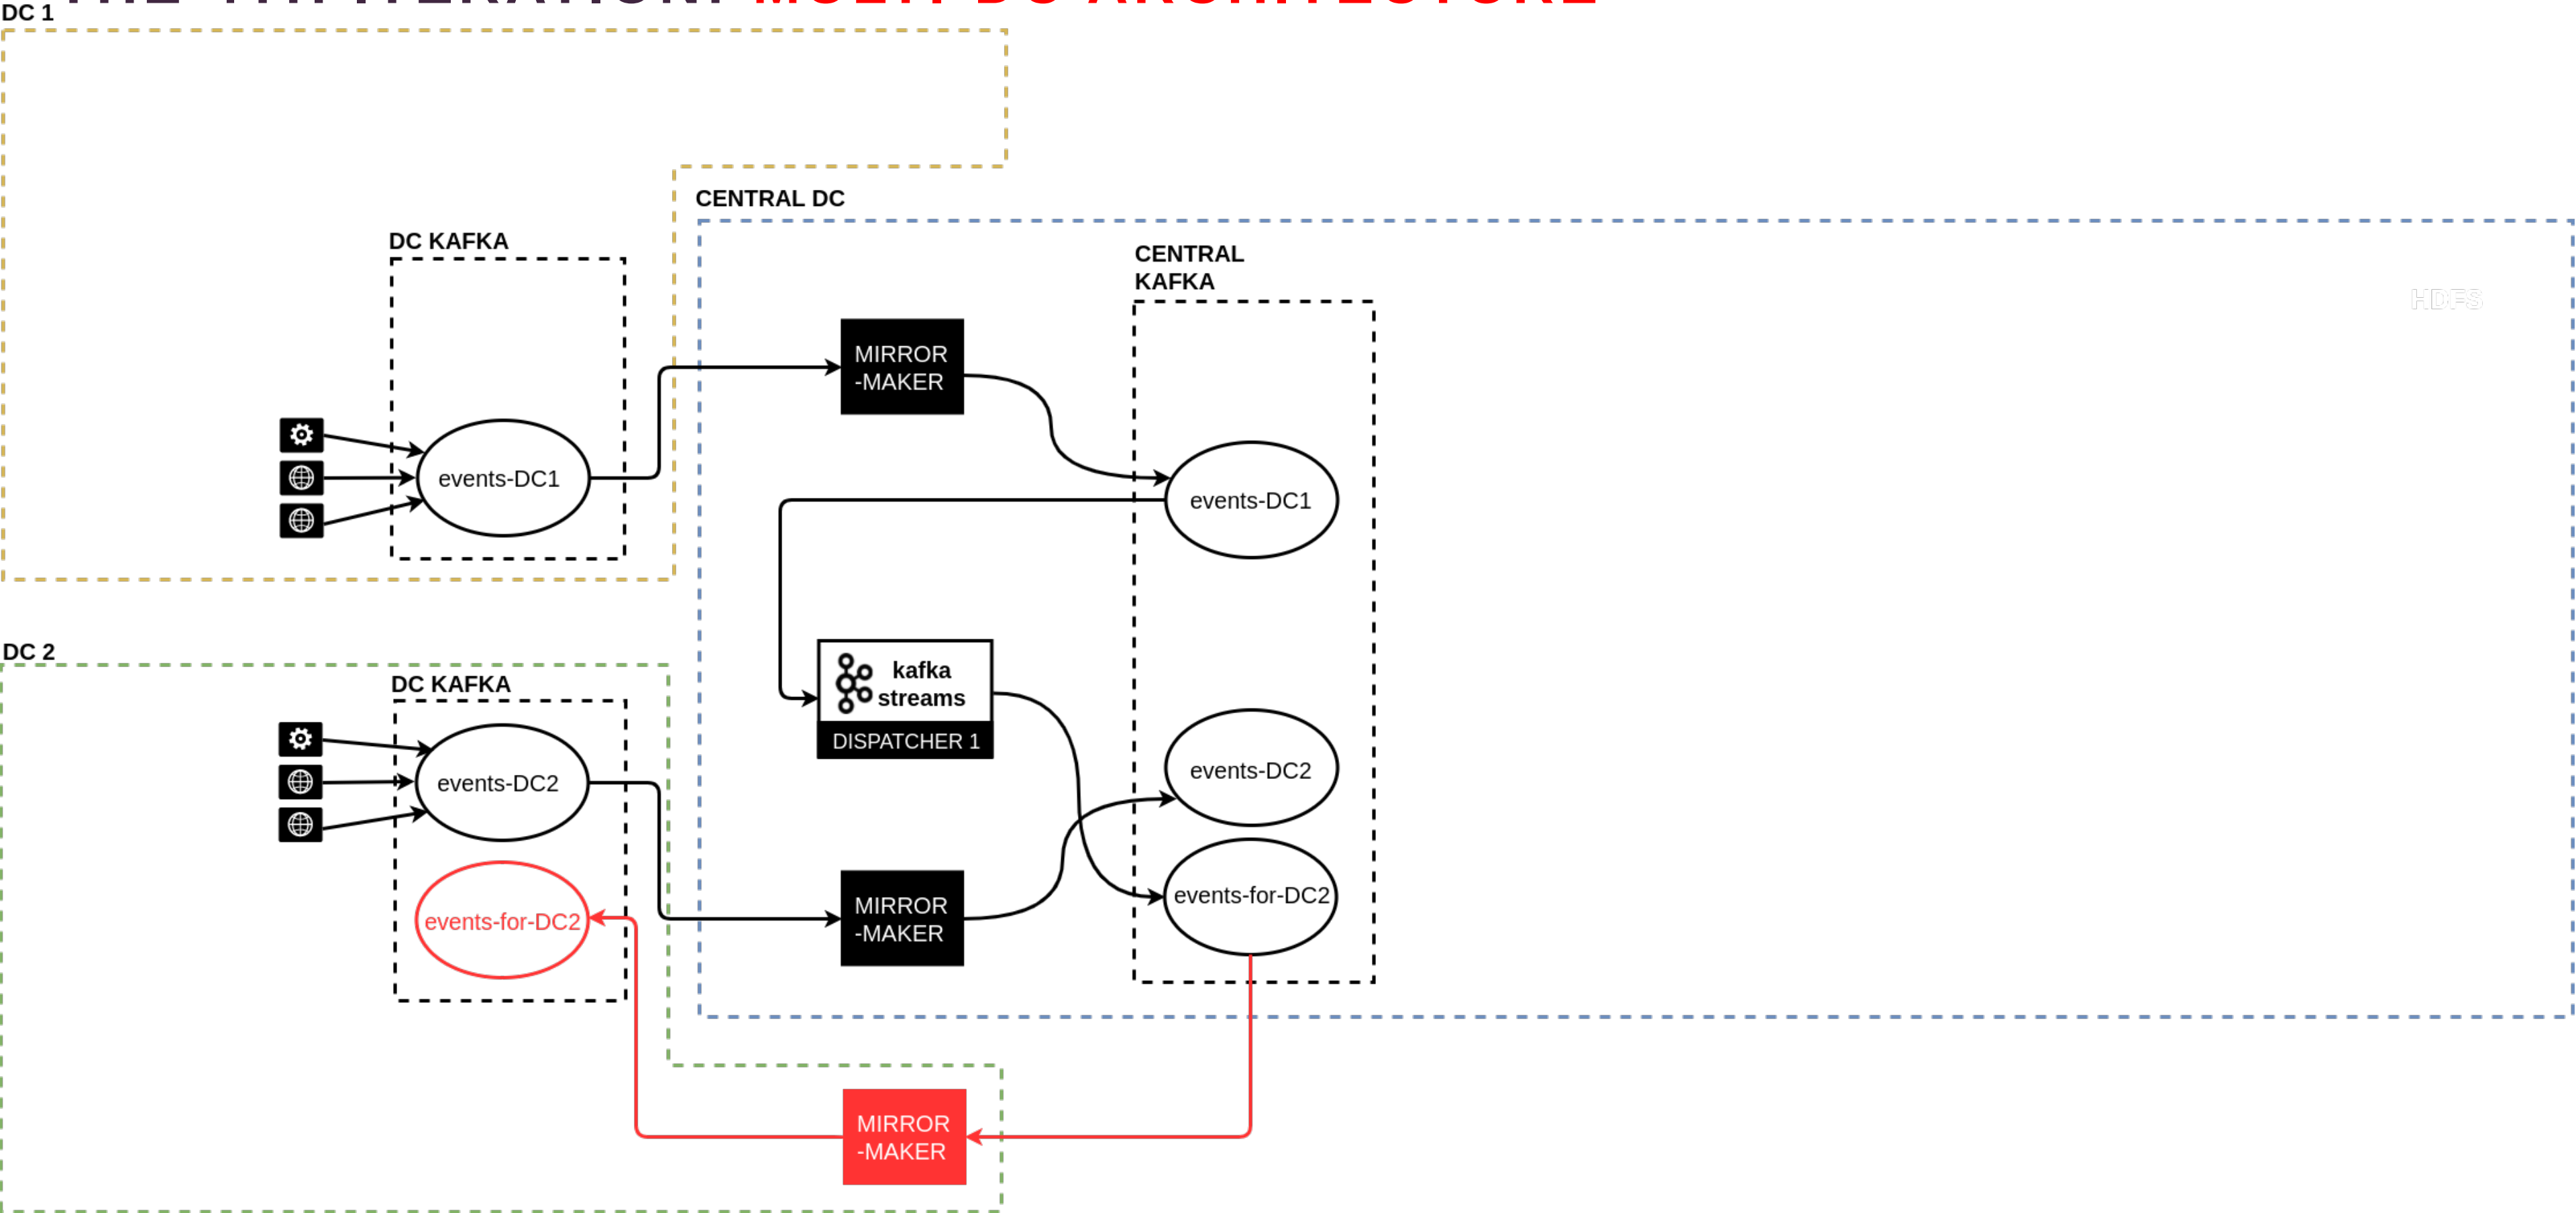
# THE 4TH ITERATION: MULTI-DC ARCHITECTURE



# THE 4TH ITERATION: MULTI-DC ARCHITECTURE




# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

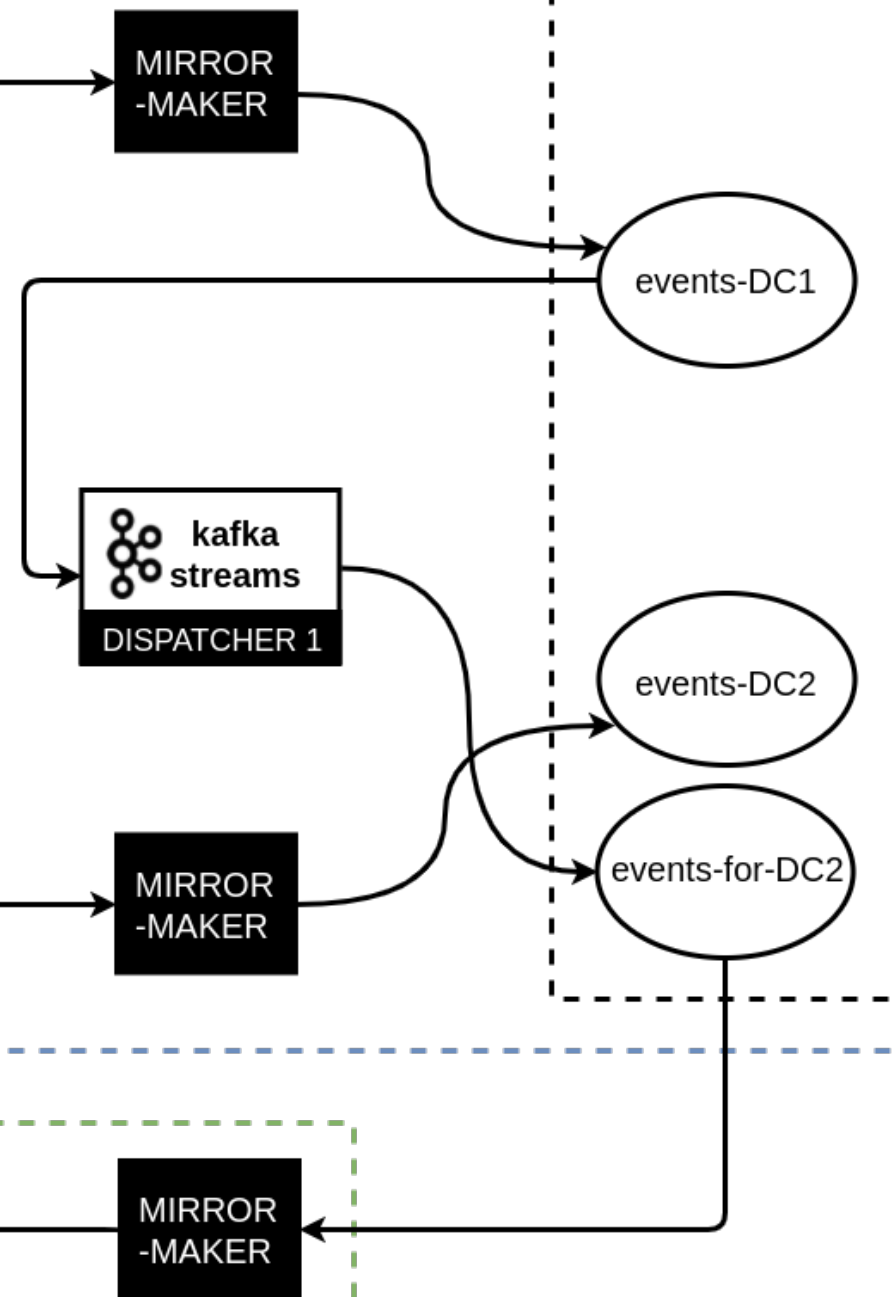


# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

DC 1

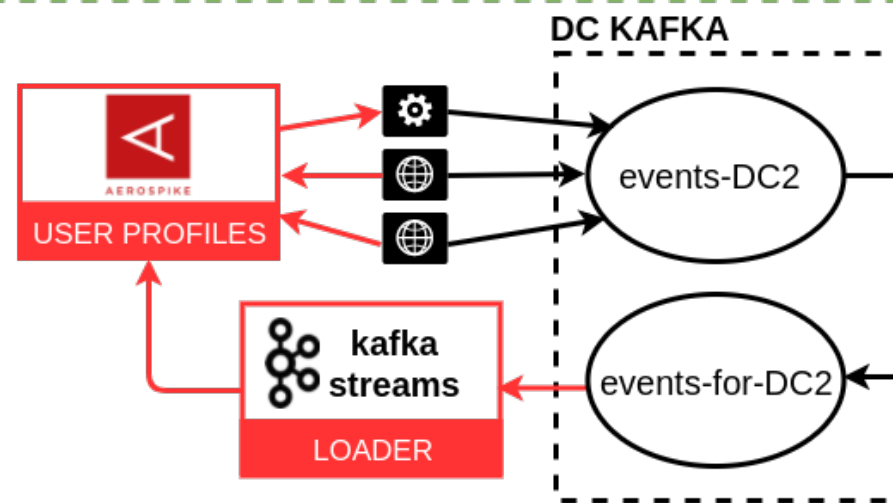


CENTRAL DC



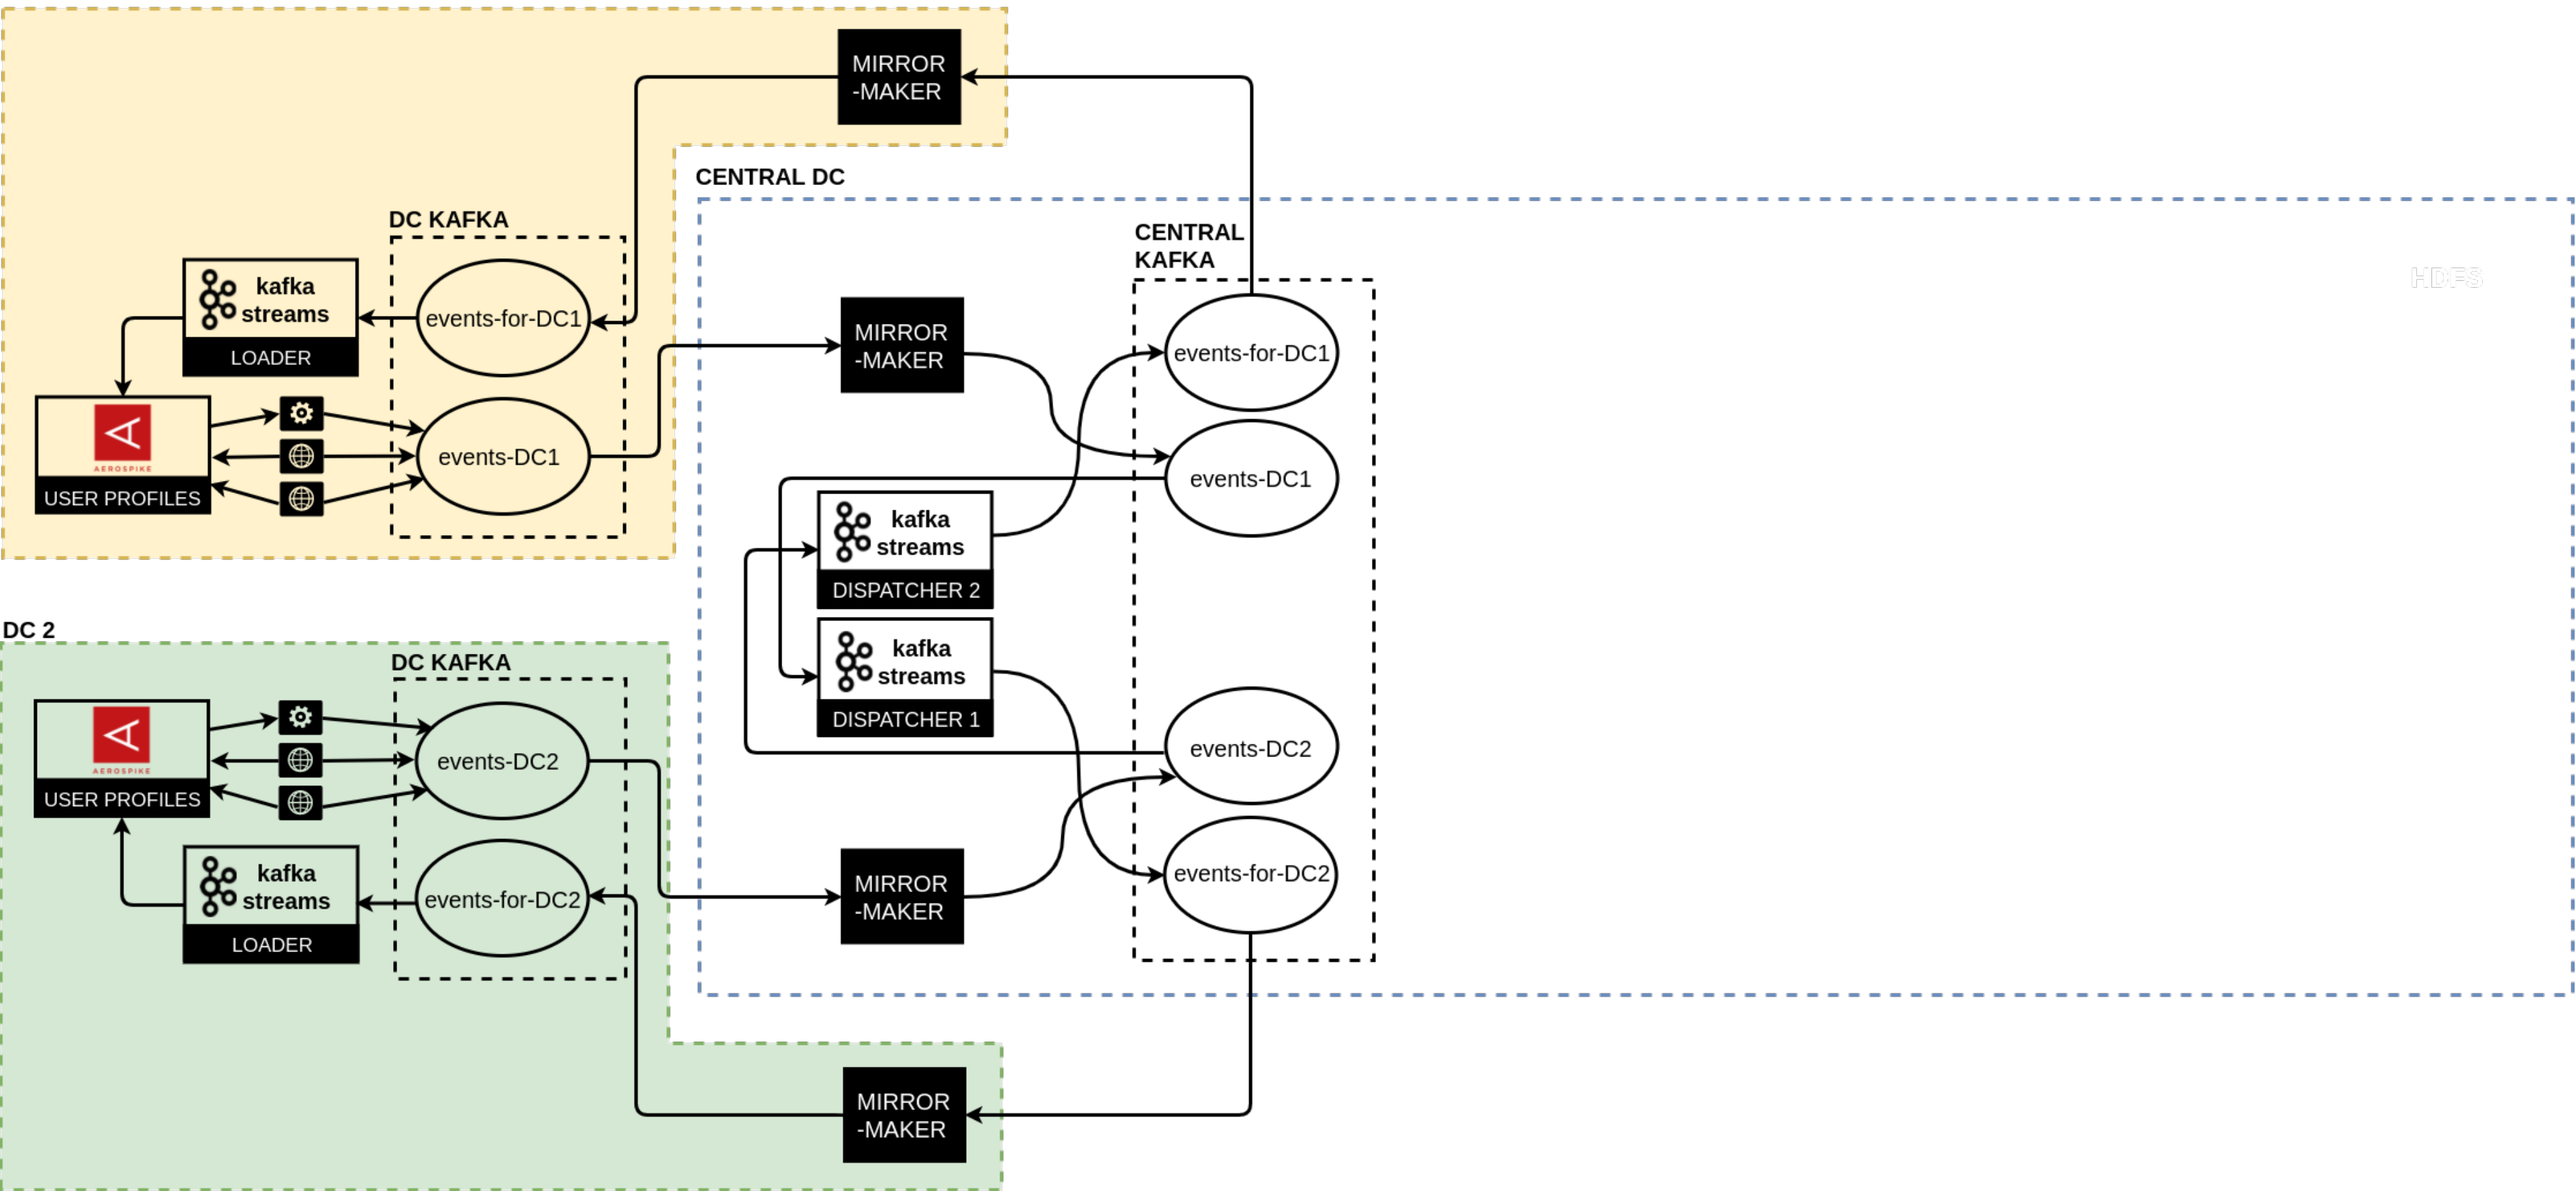
HDFS

DC 2

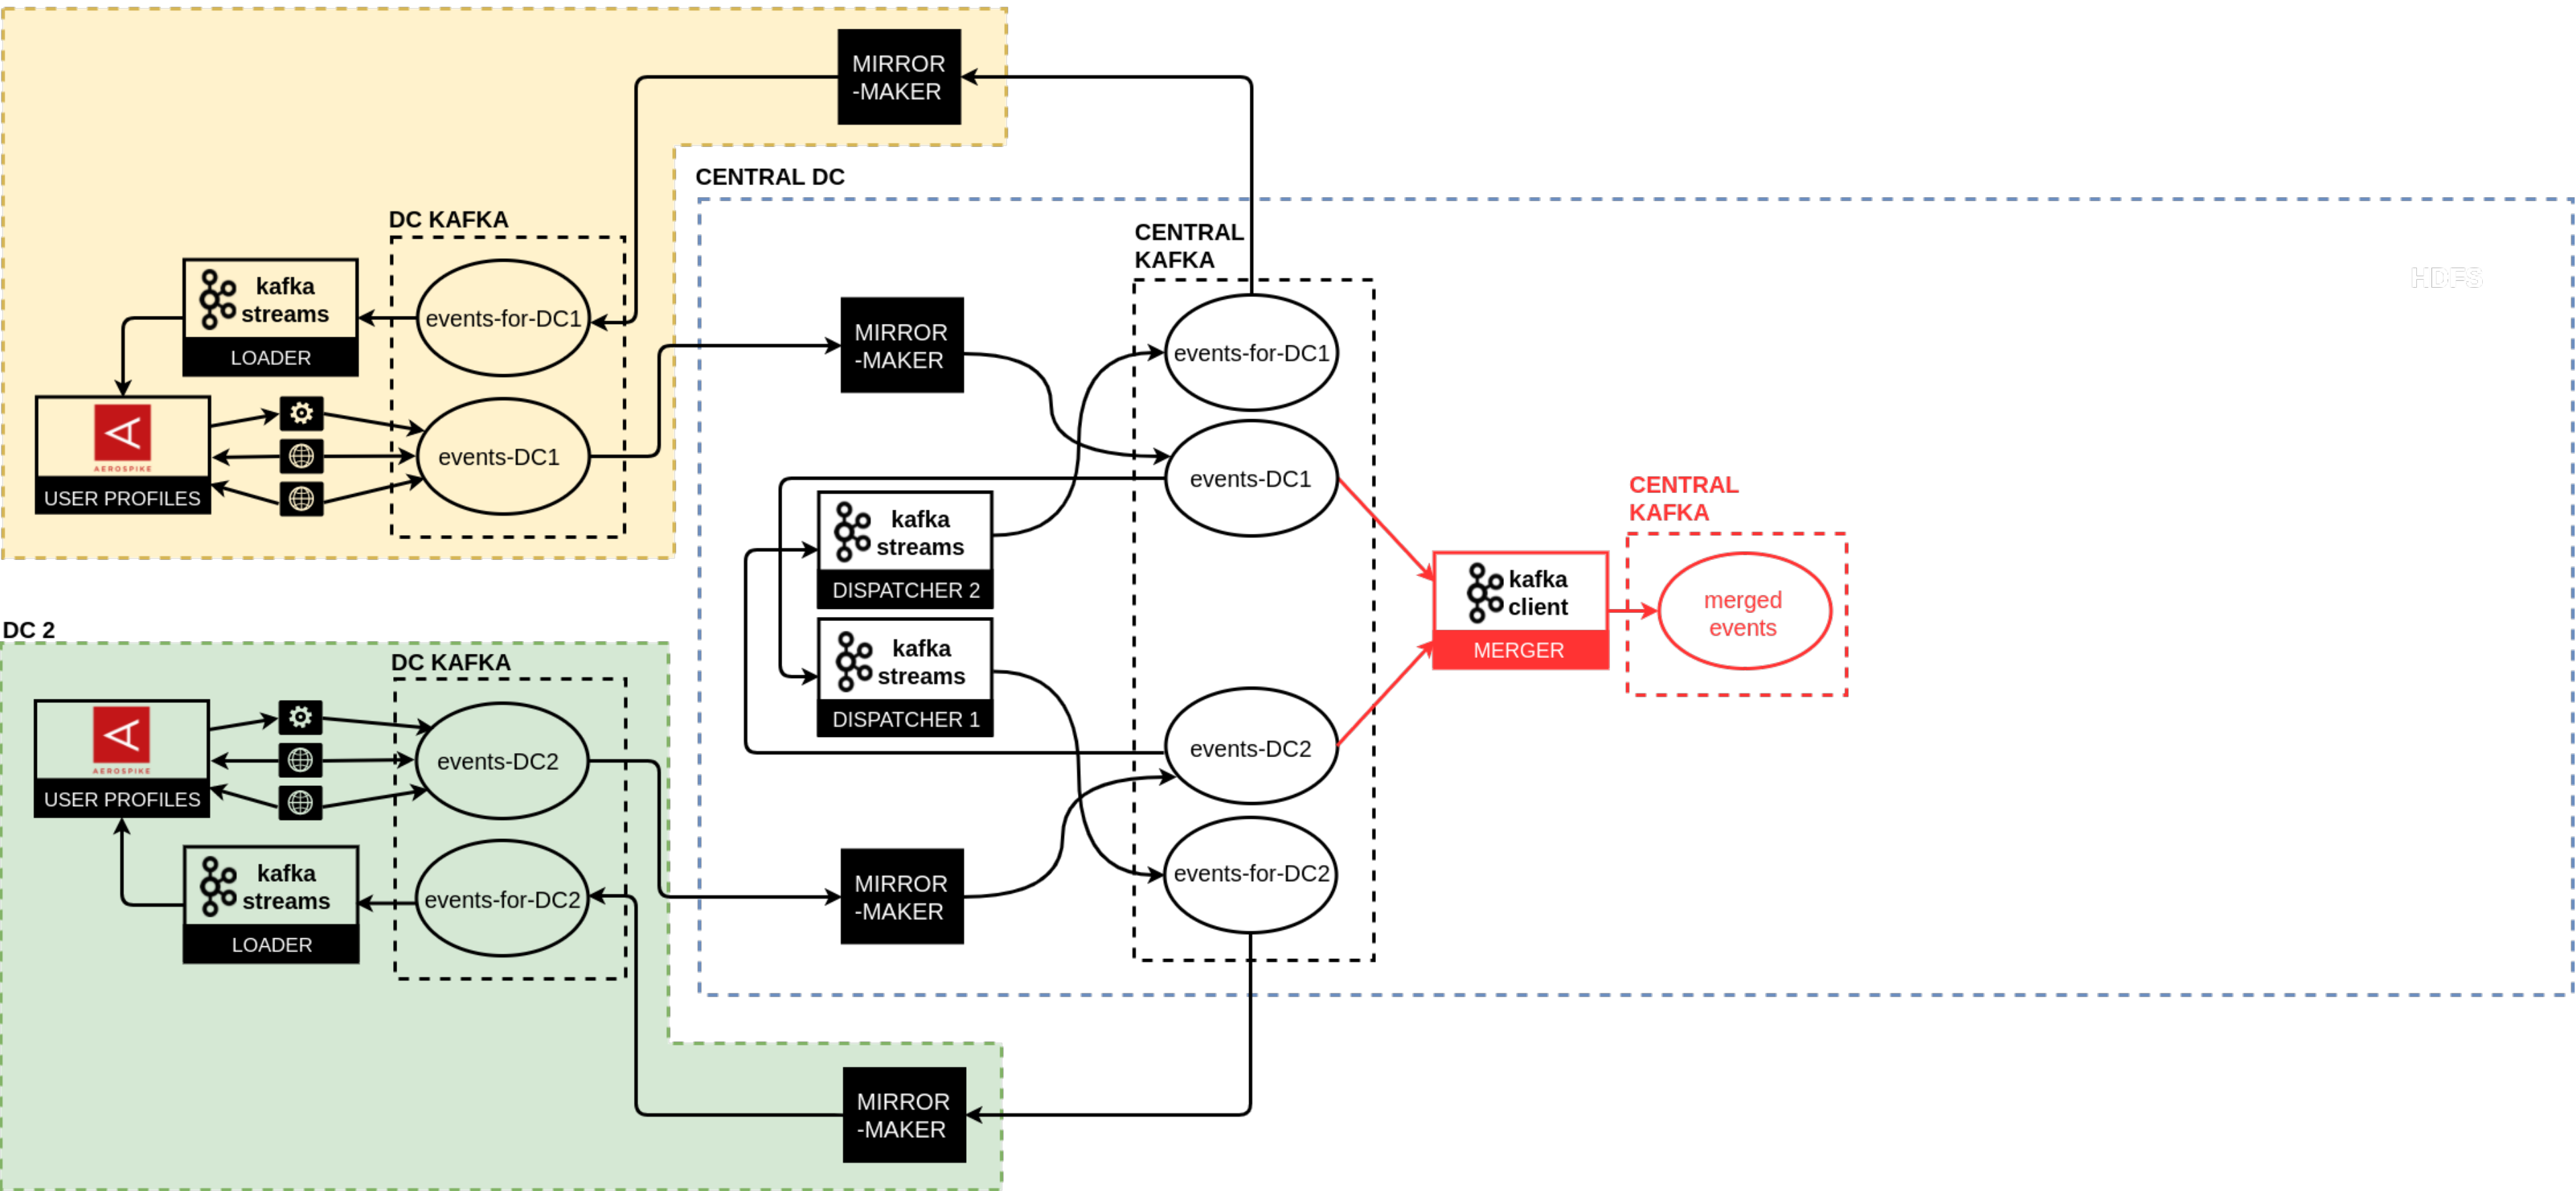




# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

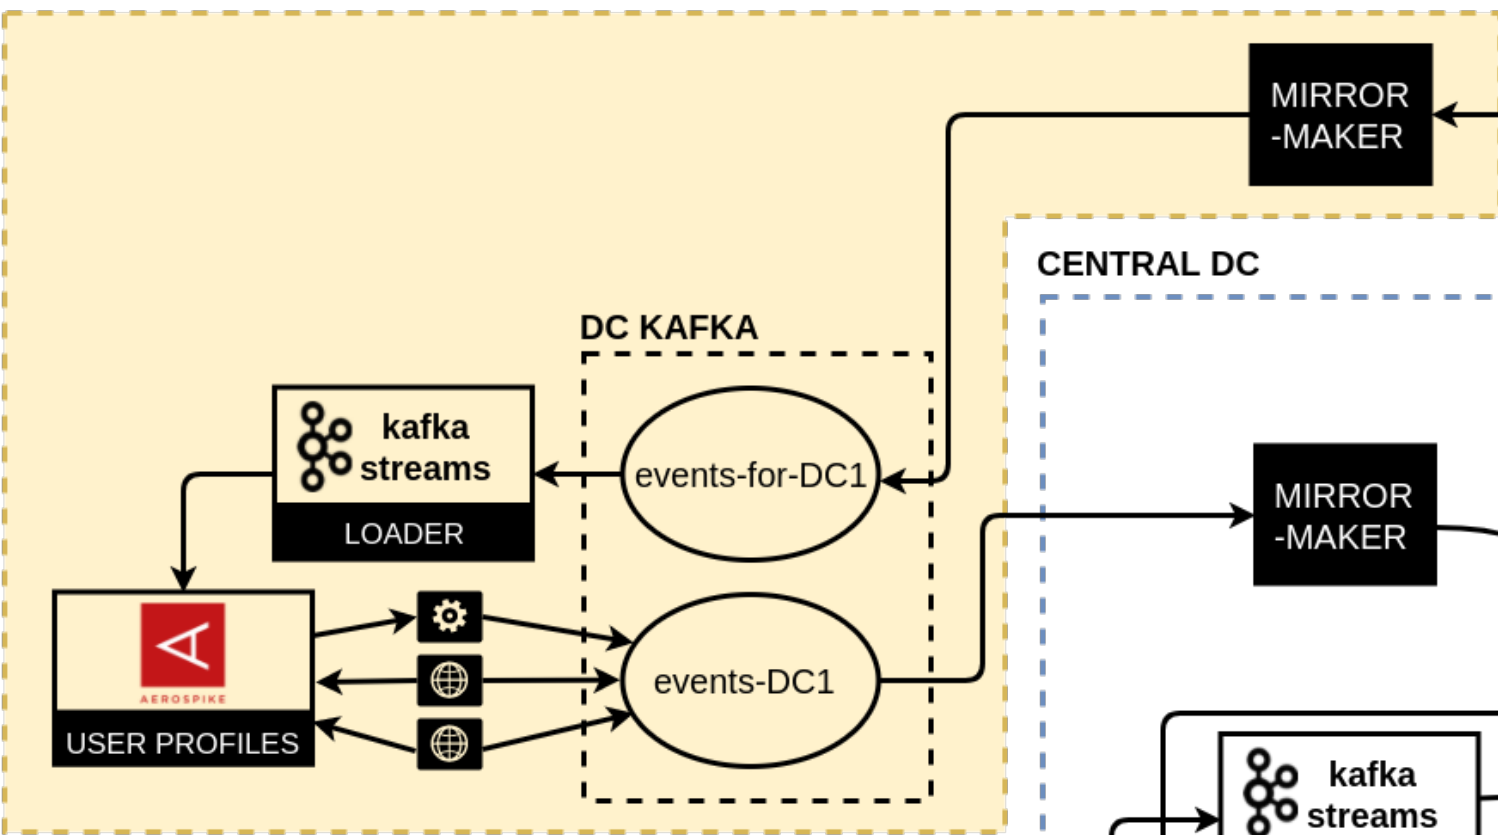


# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

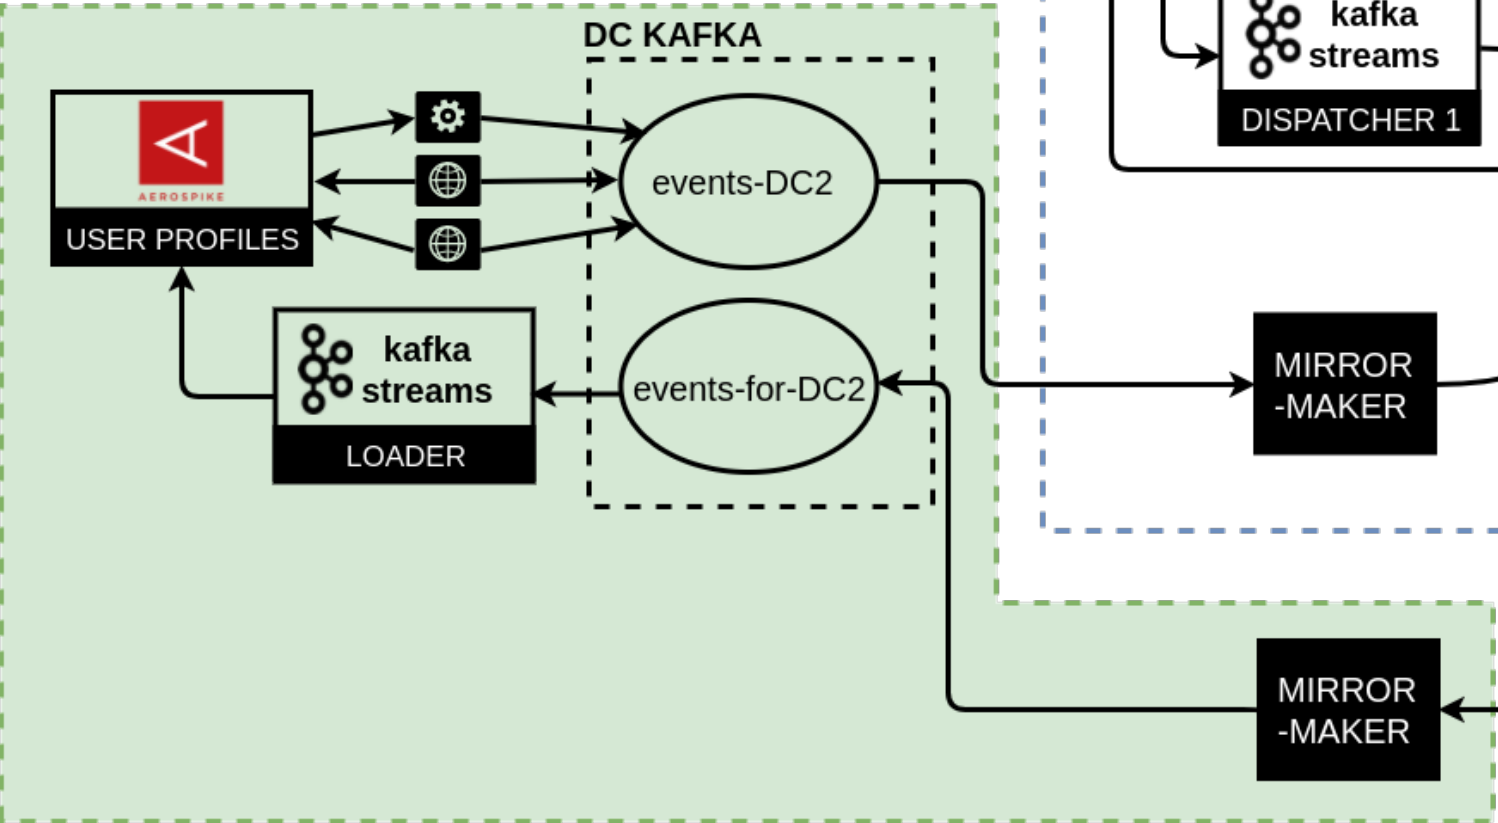


# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

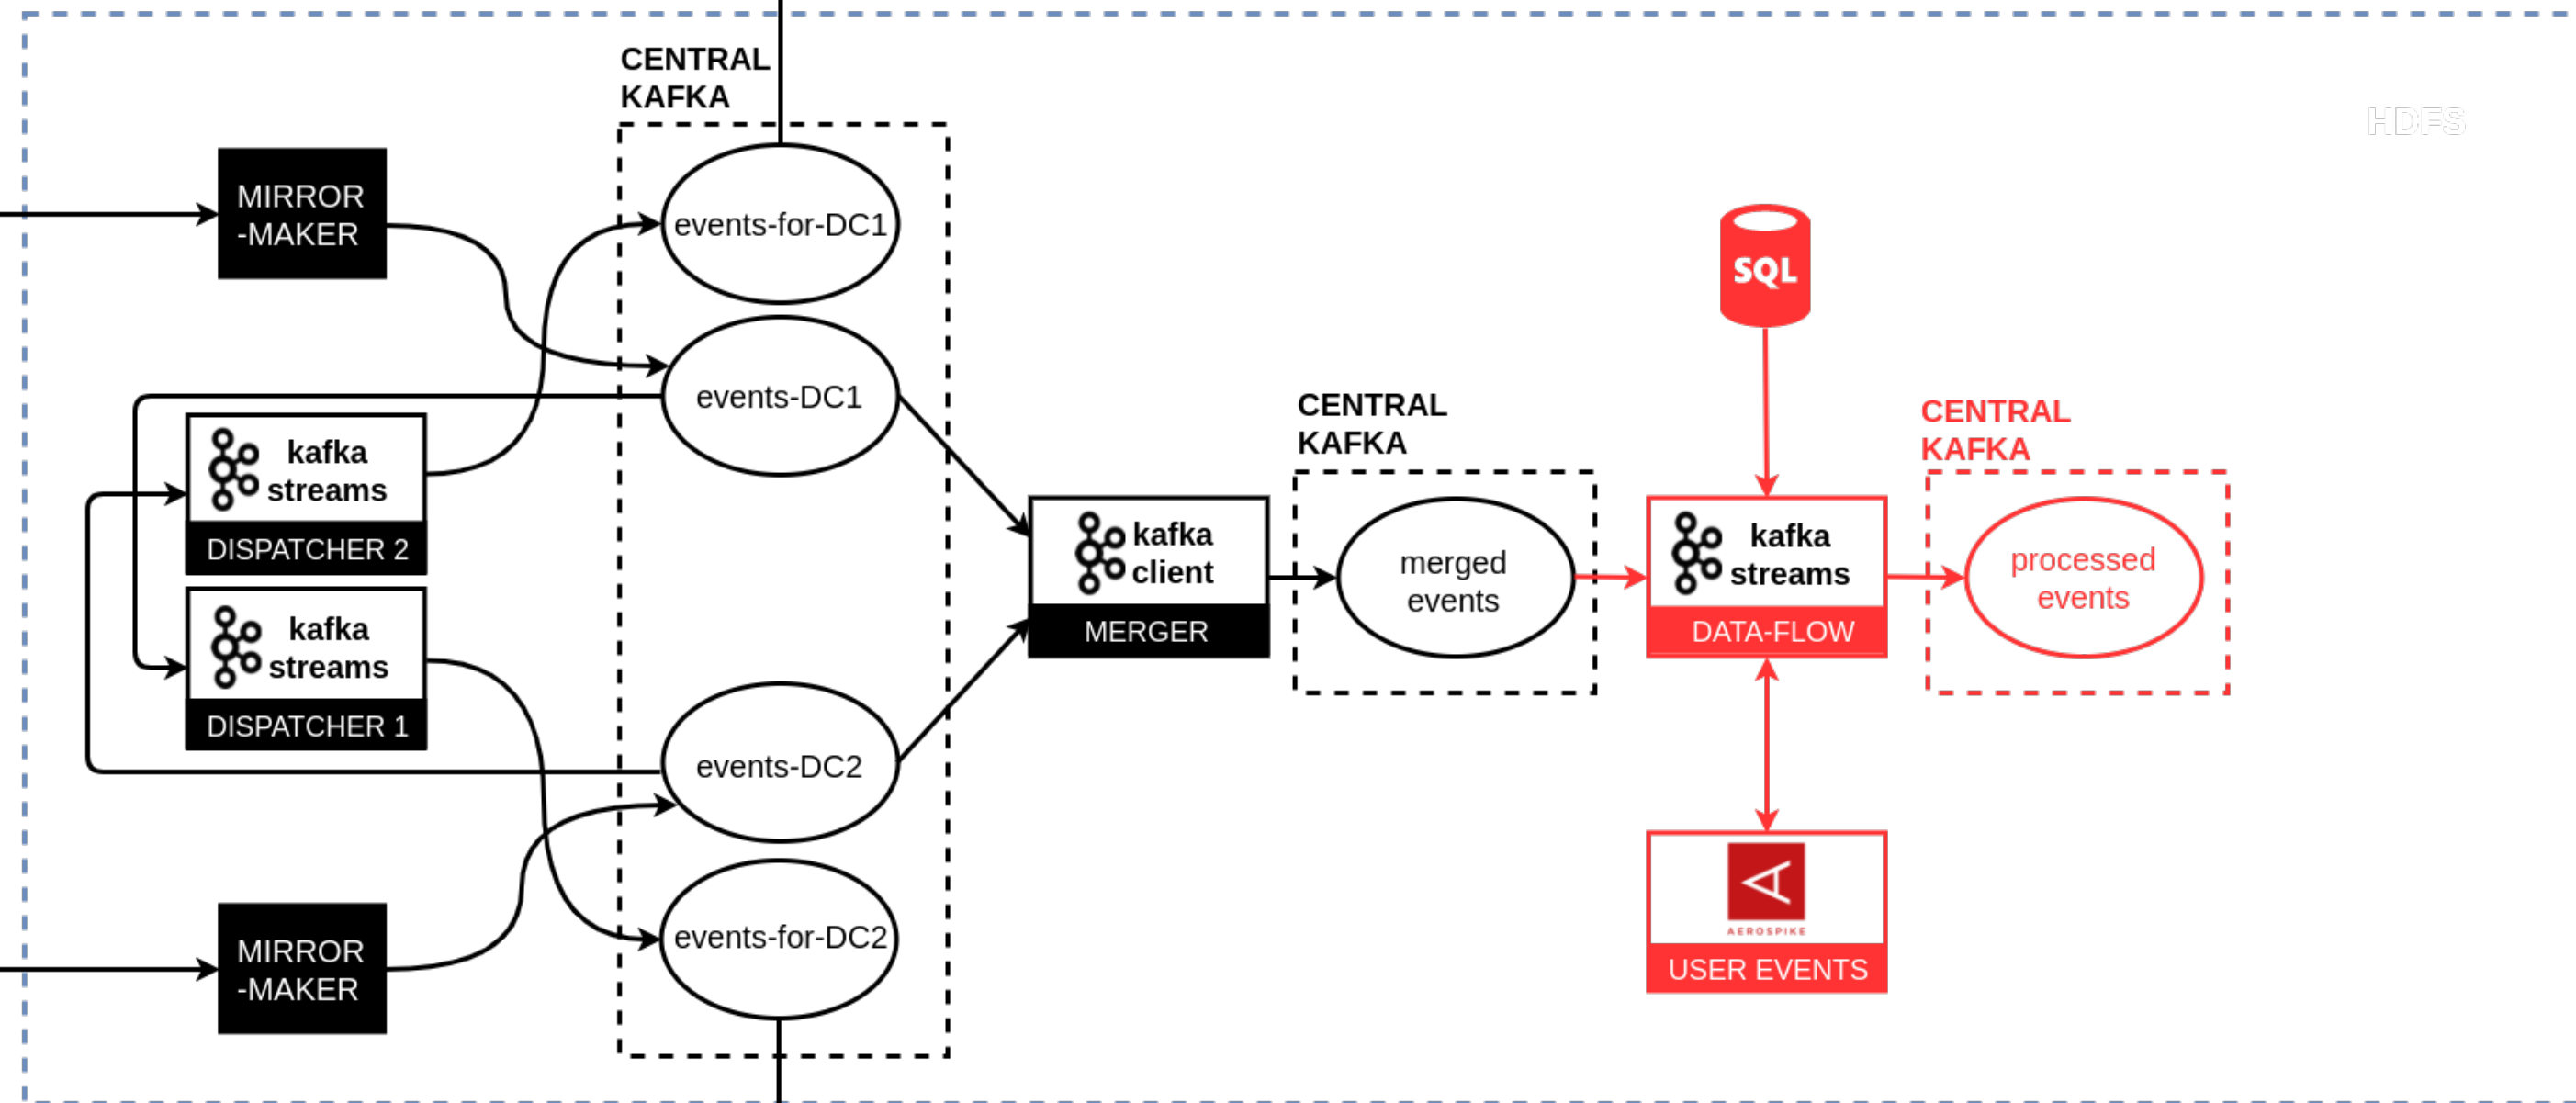
DC 1



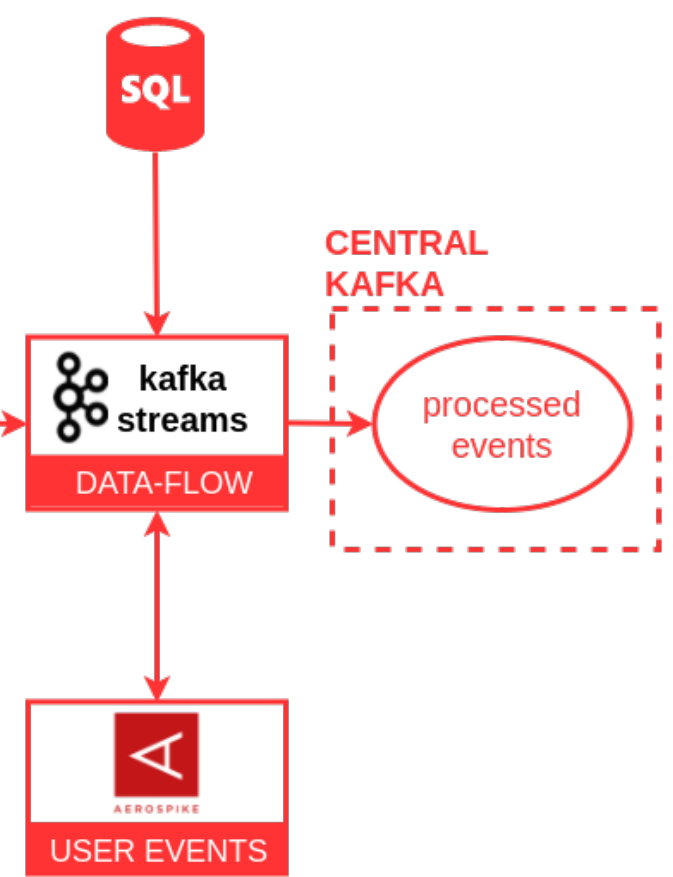
DC 2



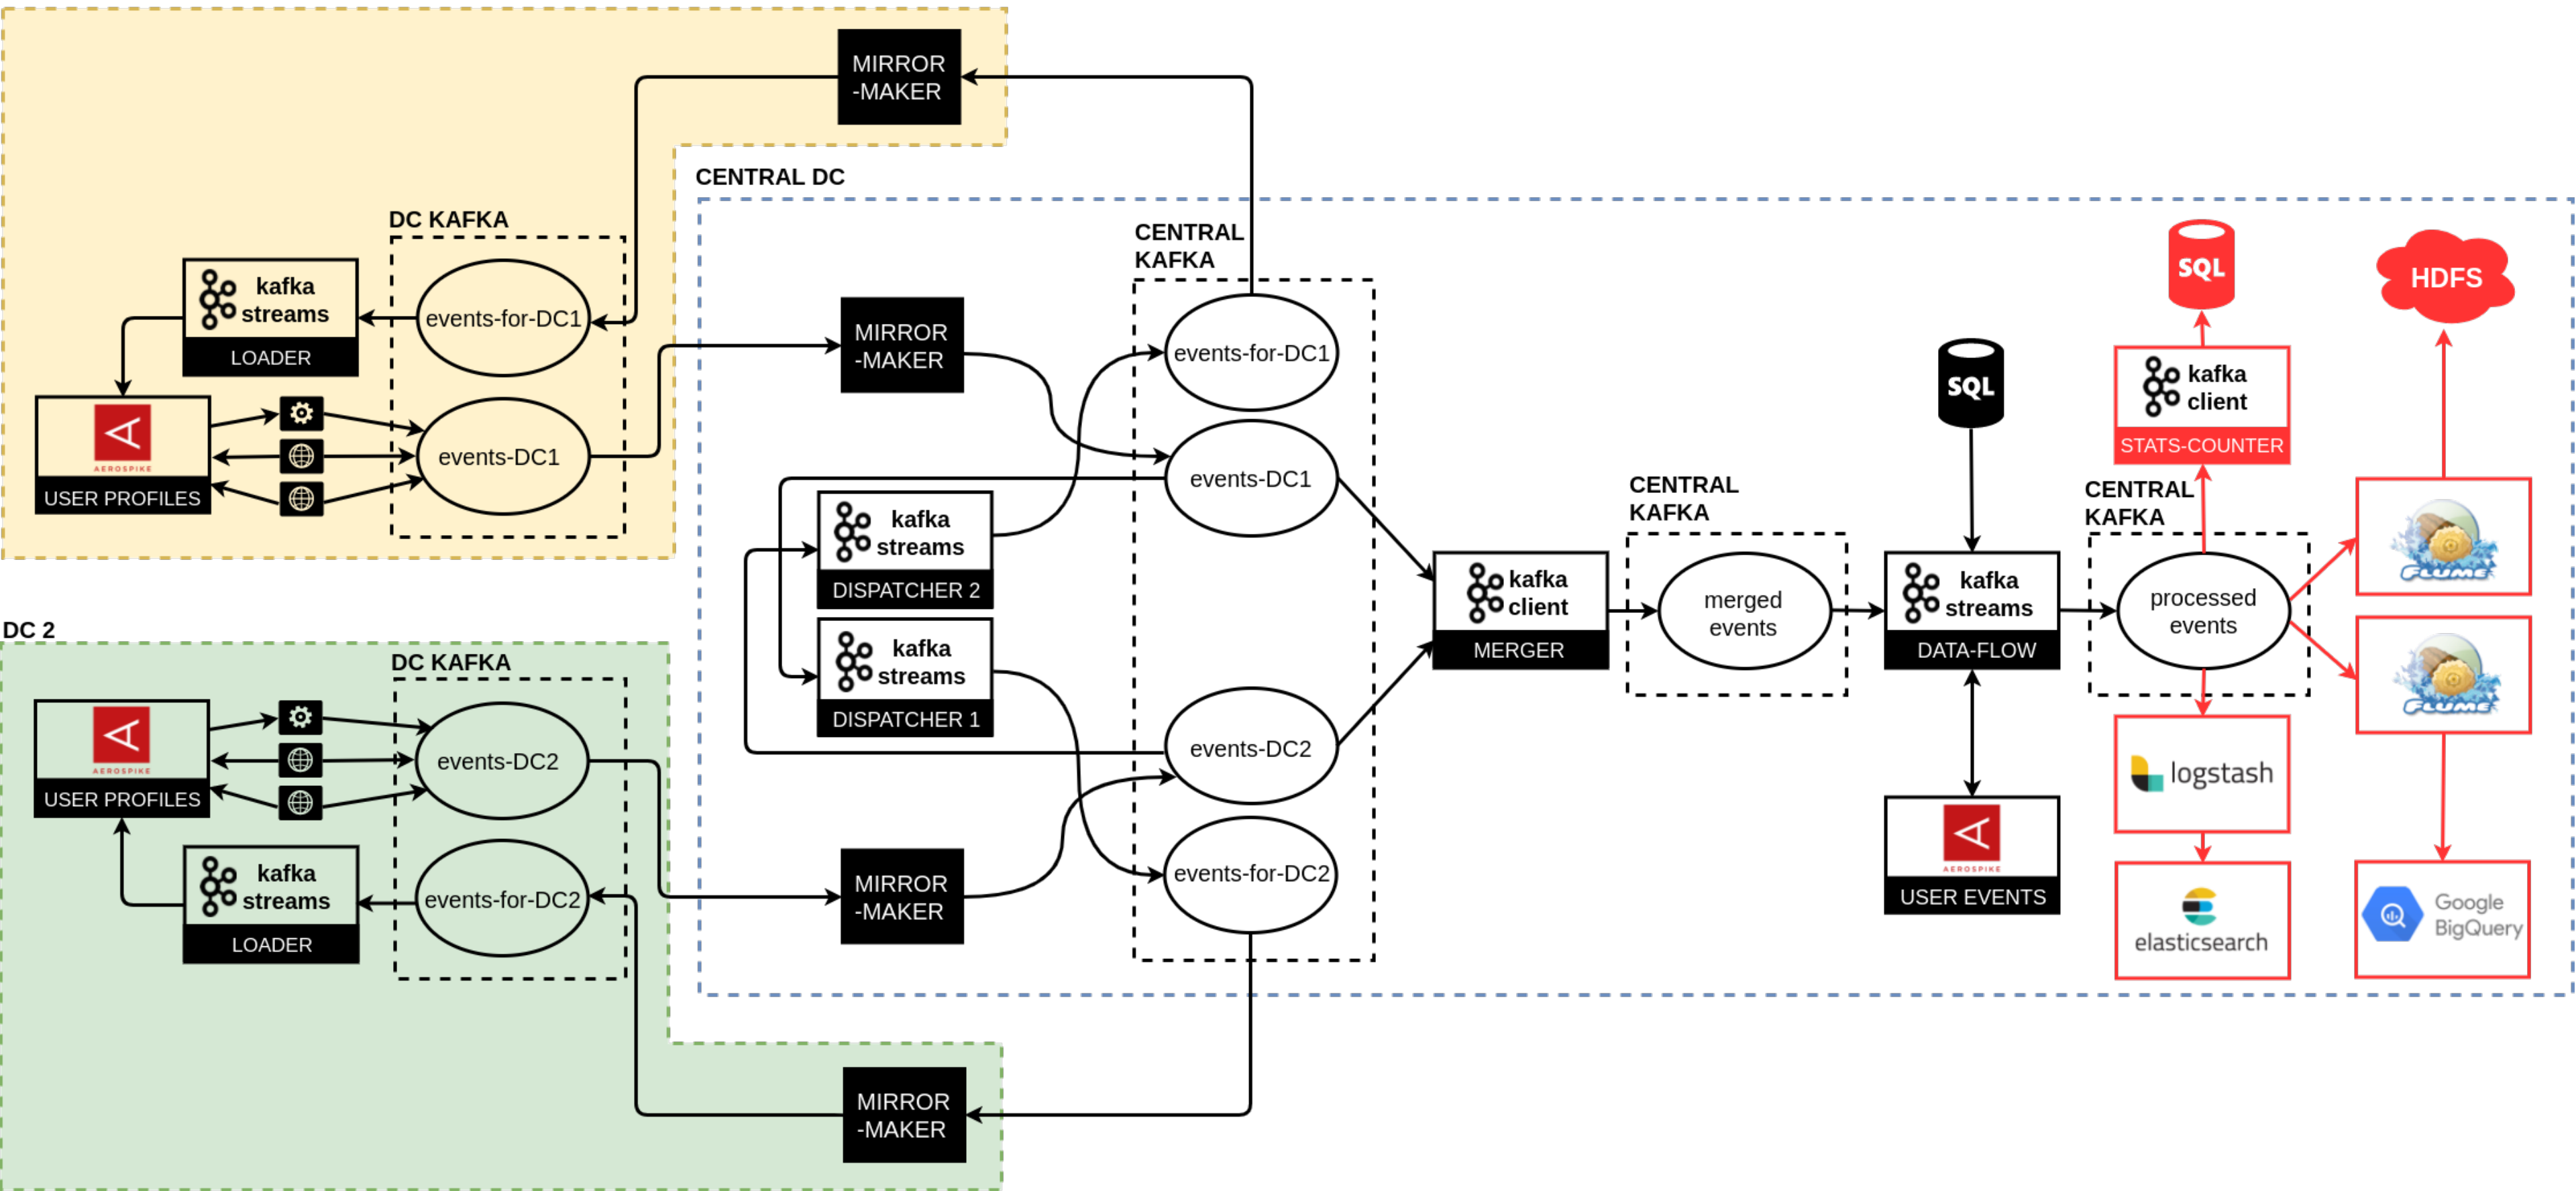
CENTRAL DC



HDFS

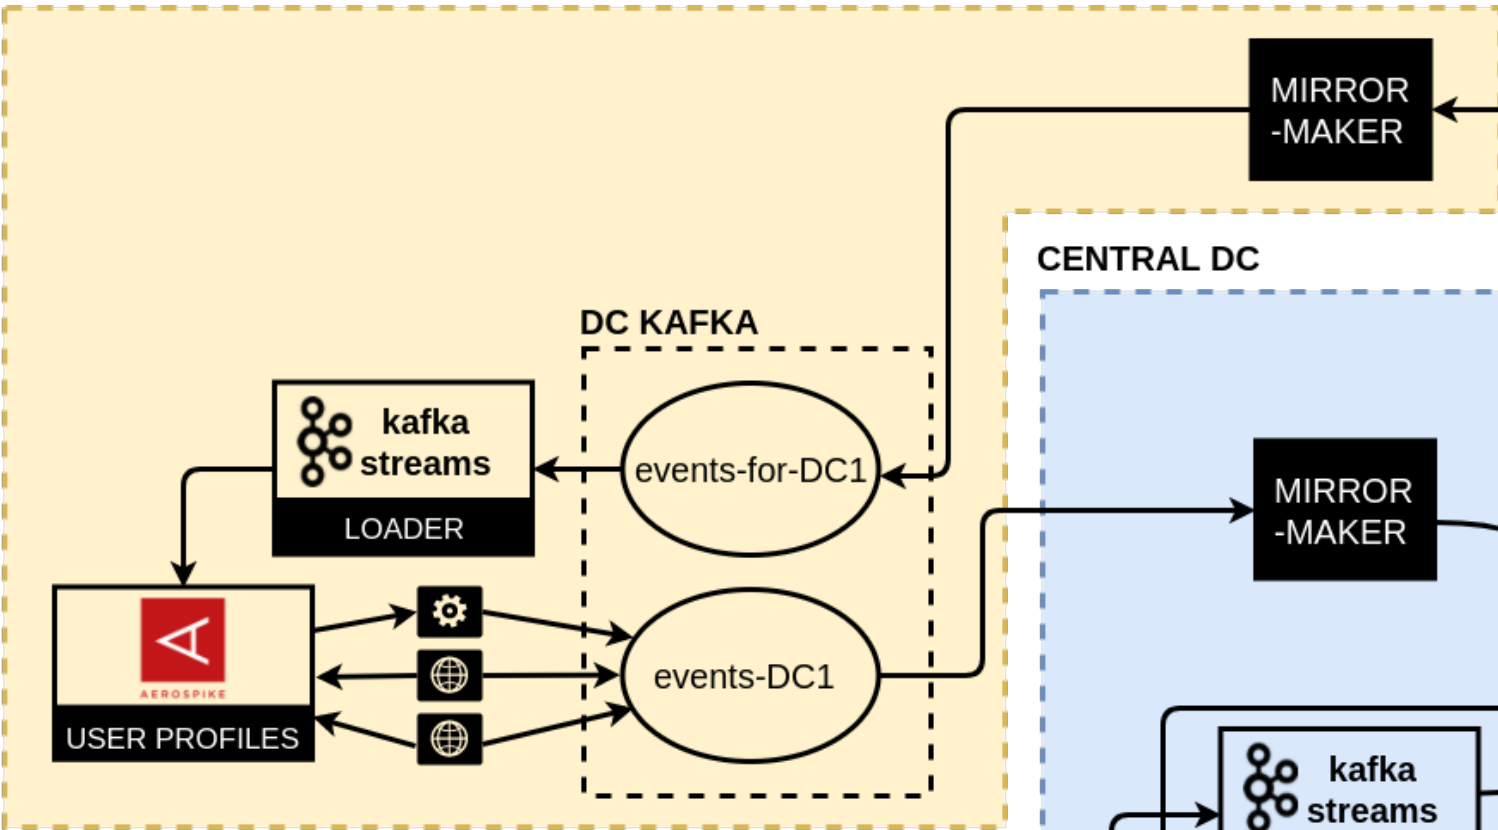


# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

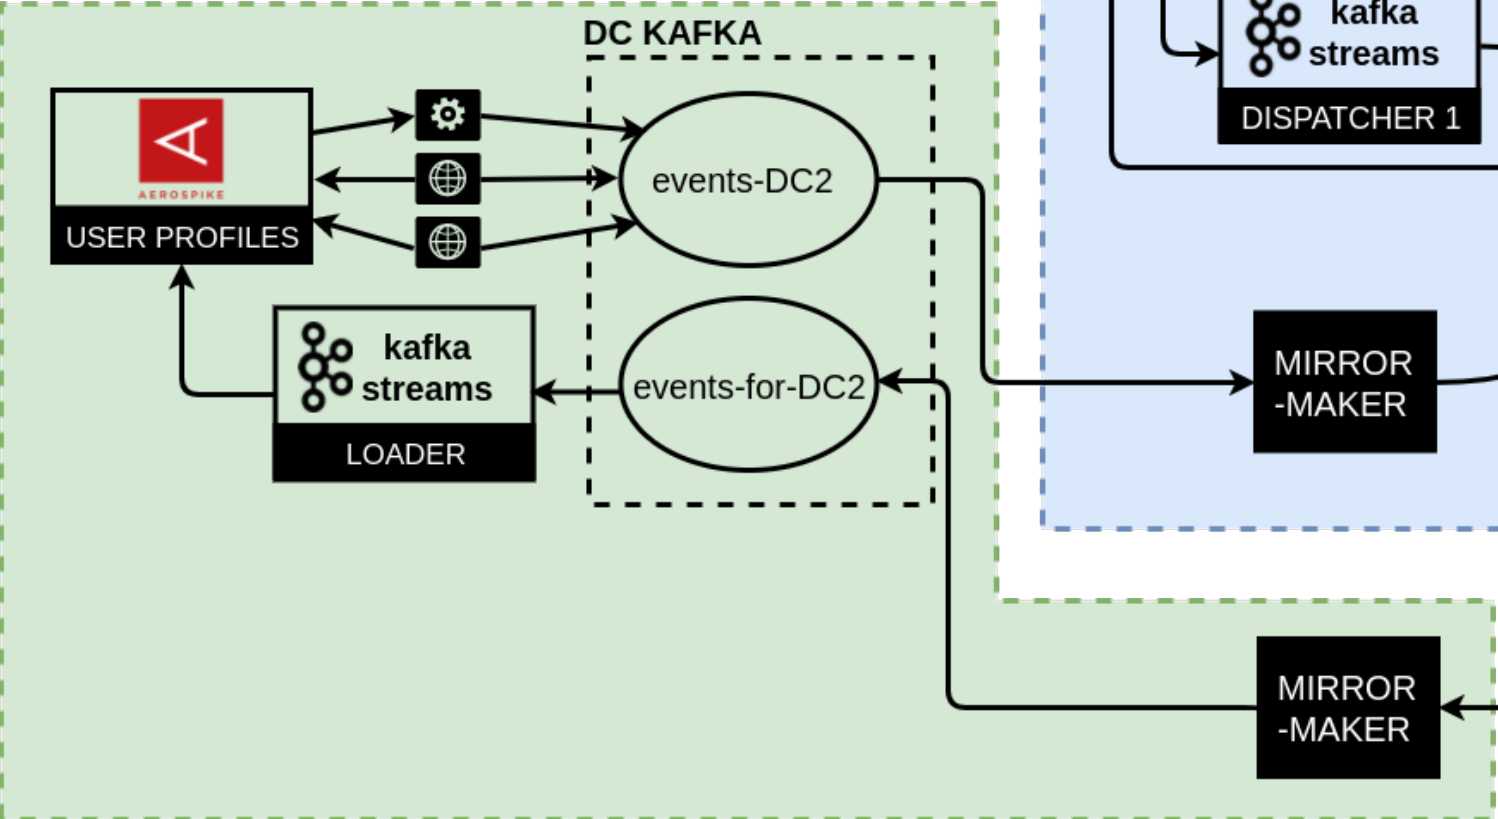


# THE 4TH ITERATION: MULTI-DC ARCHITECTURE

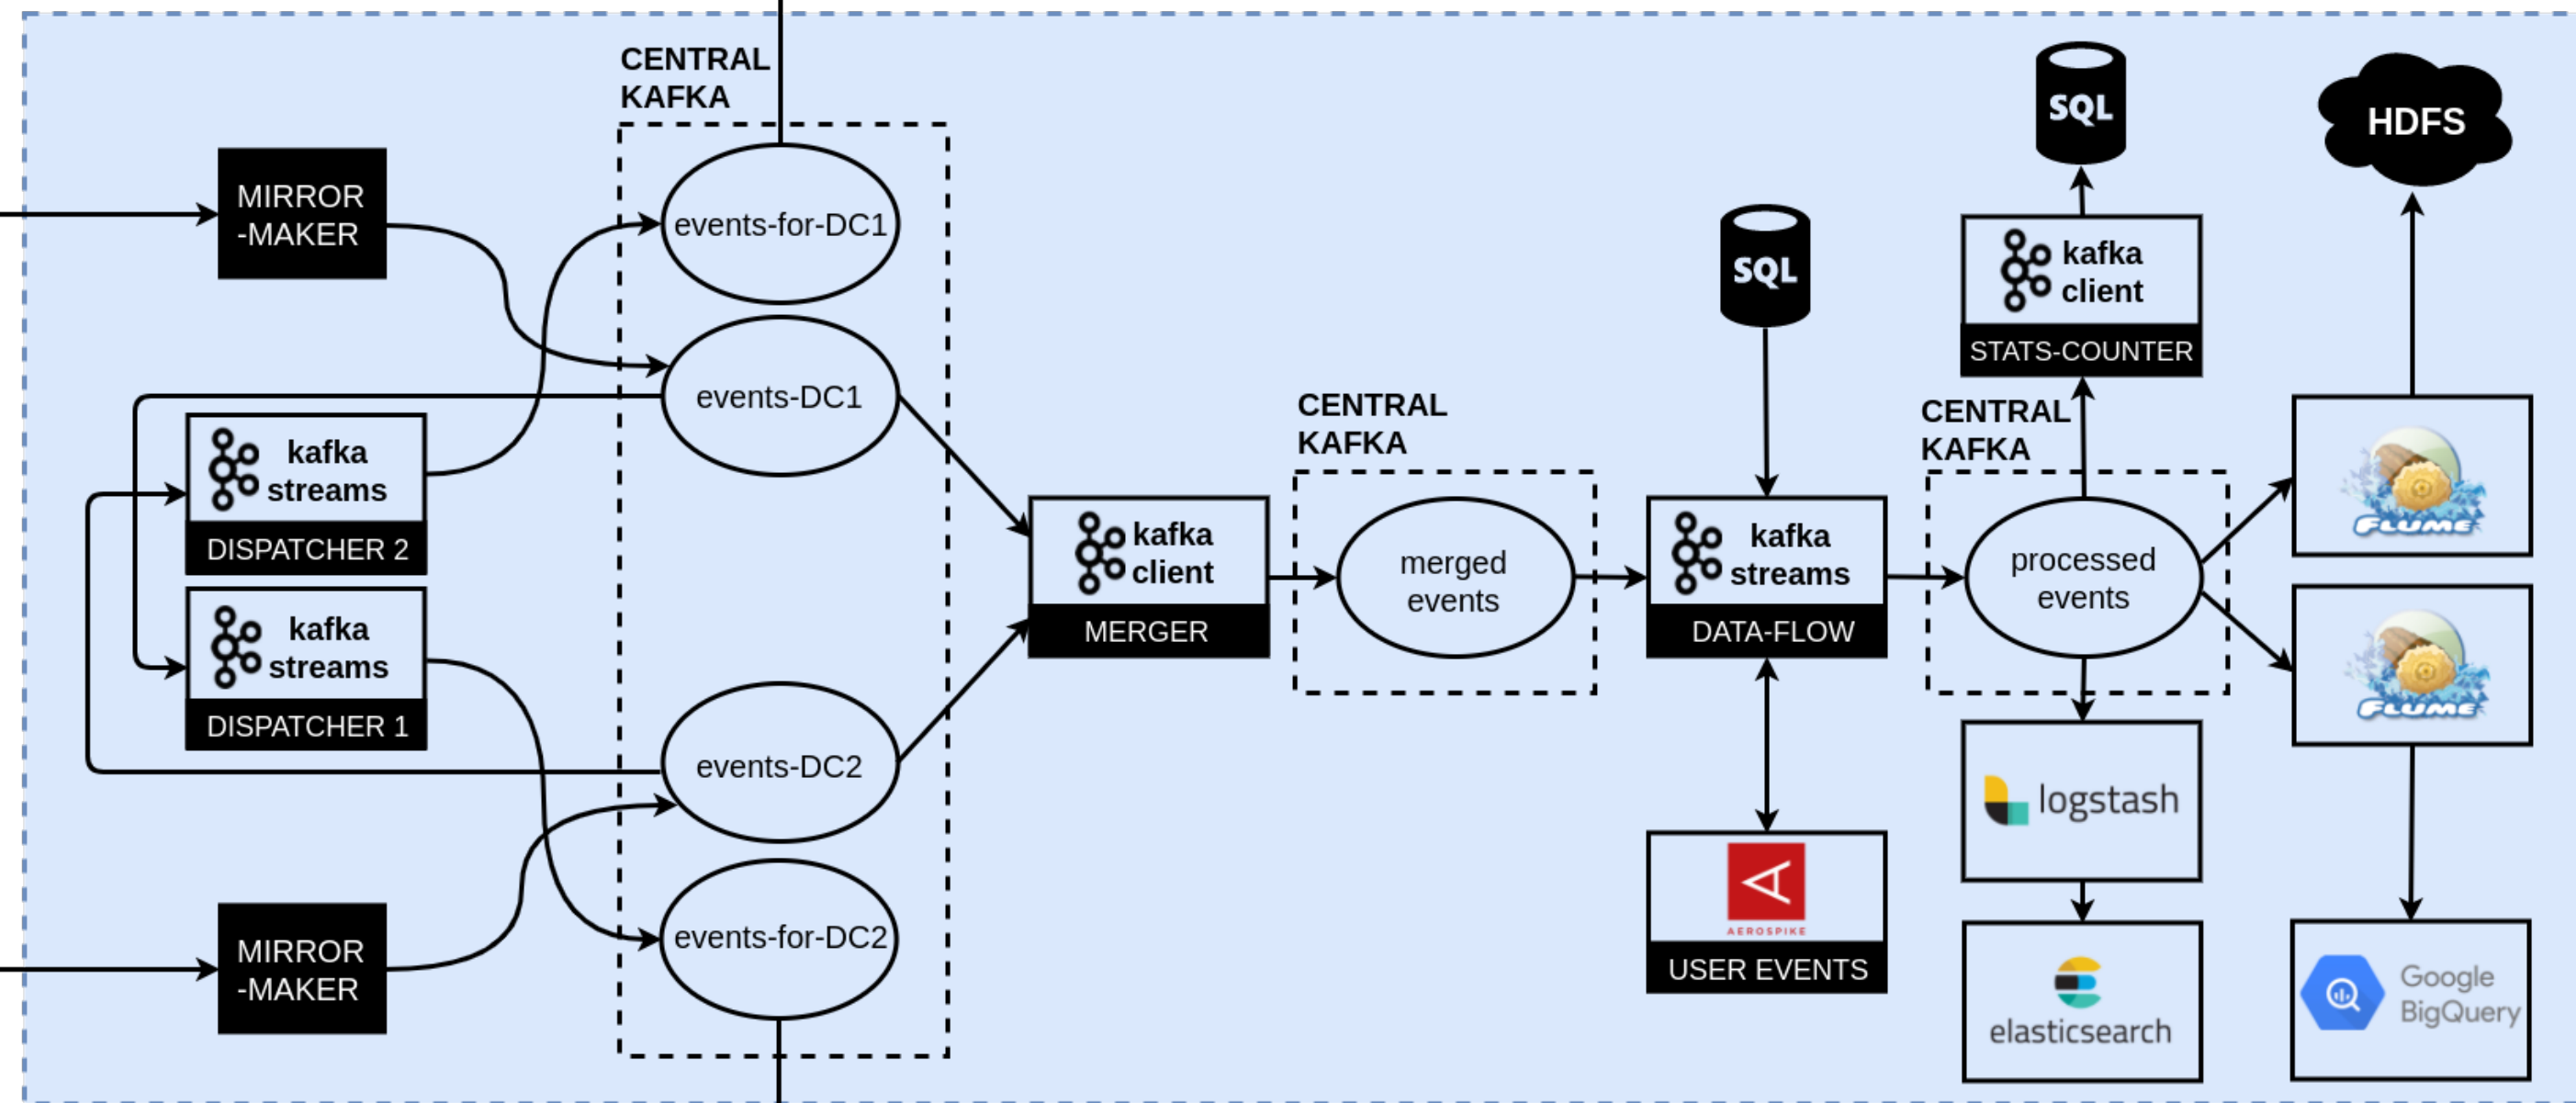
DC 1



DC 2



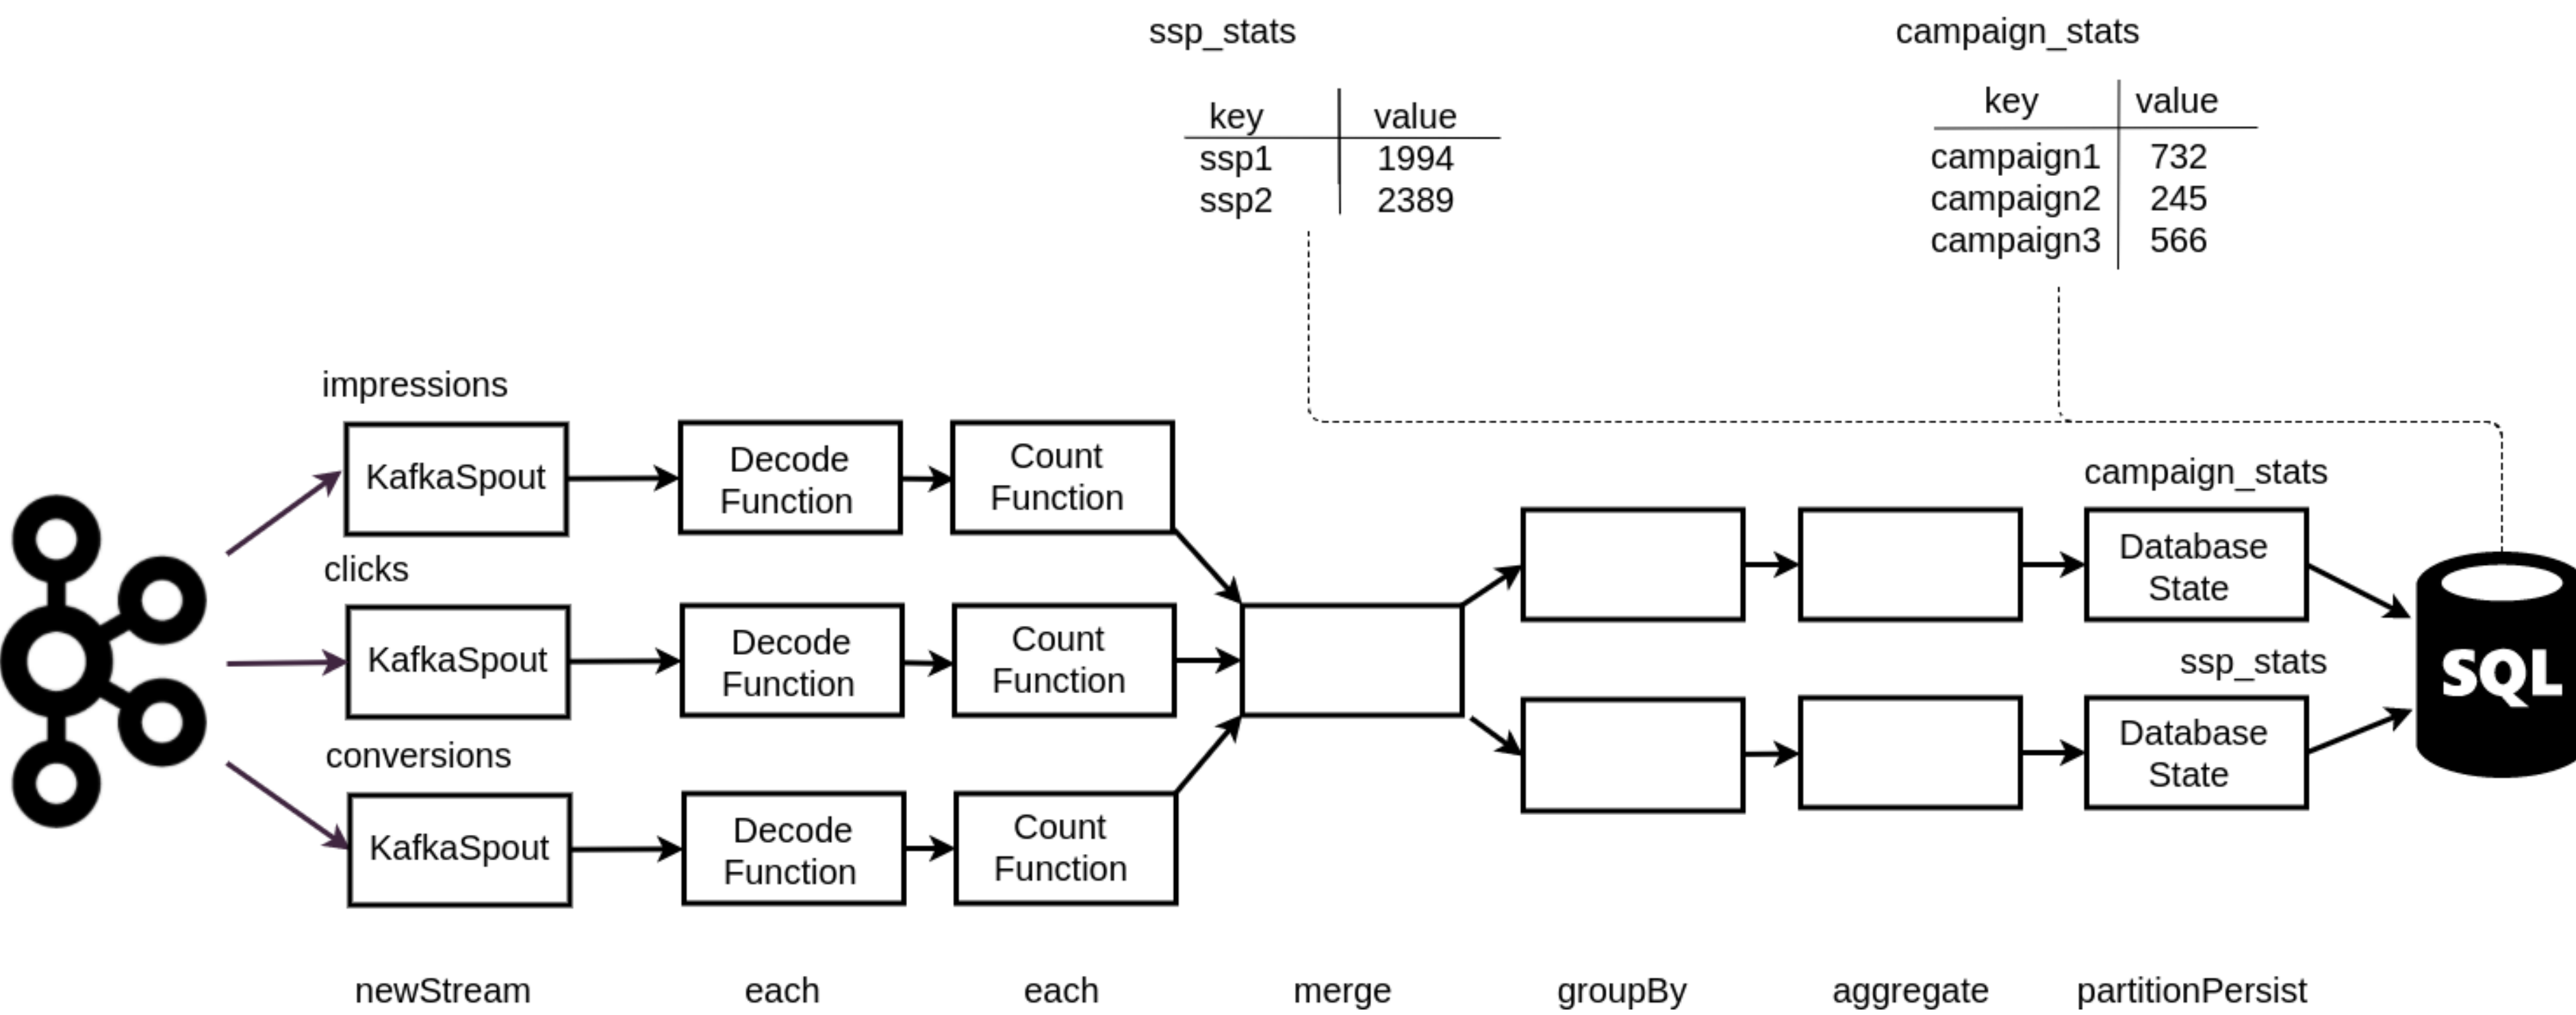
CENTRAL DC



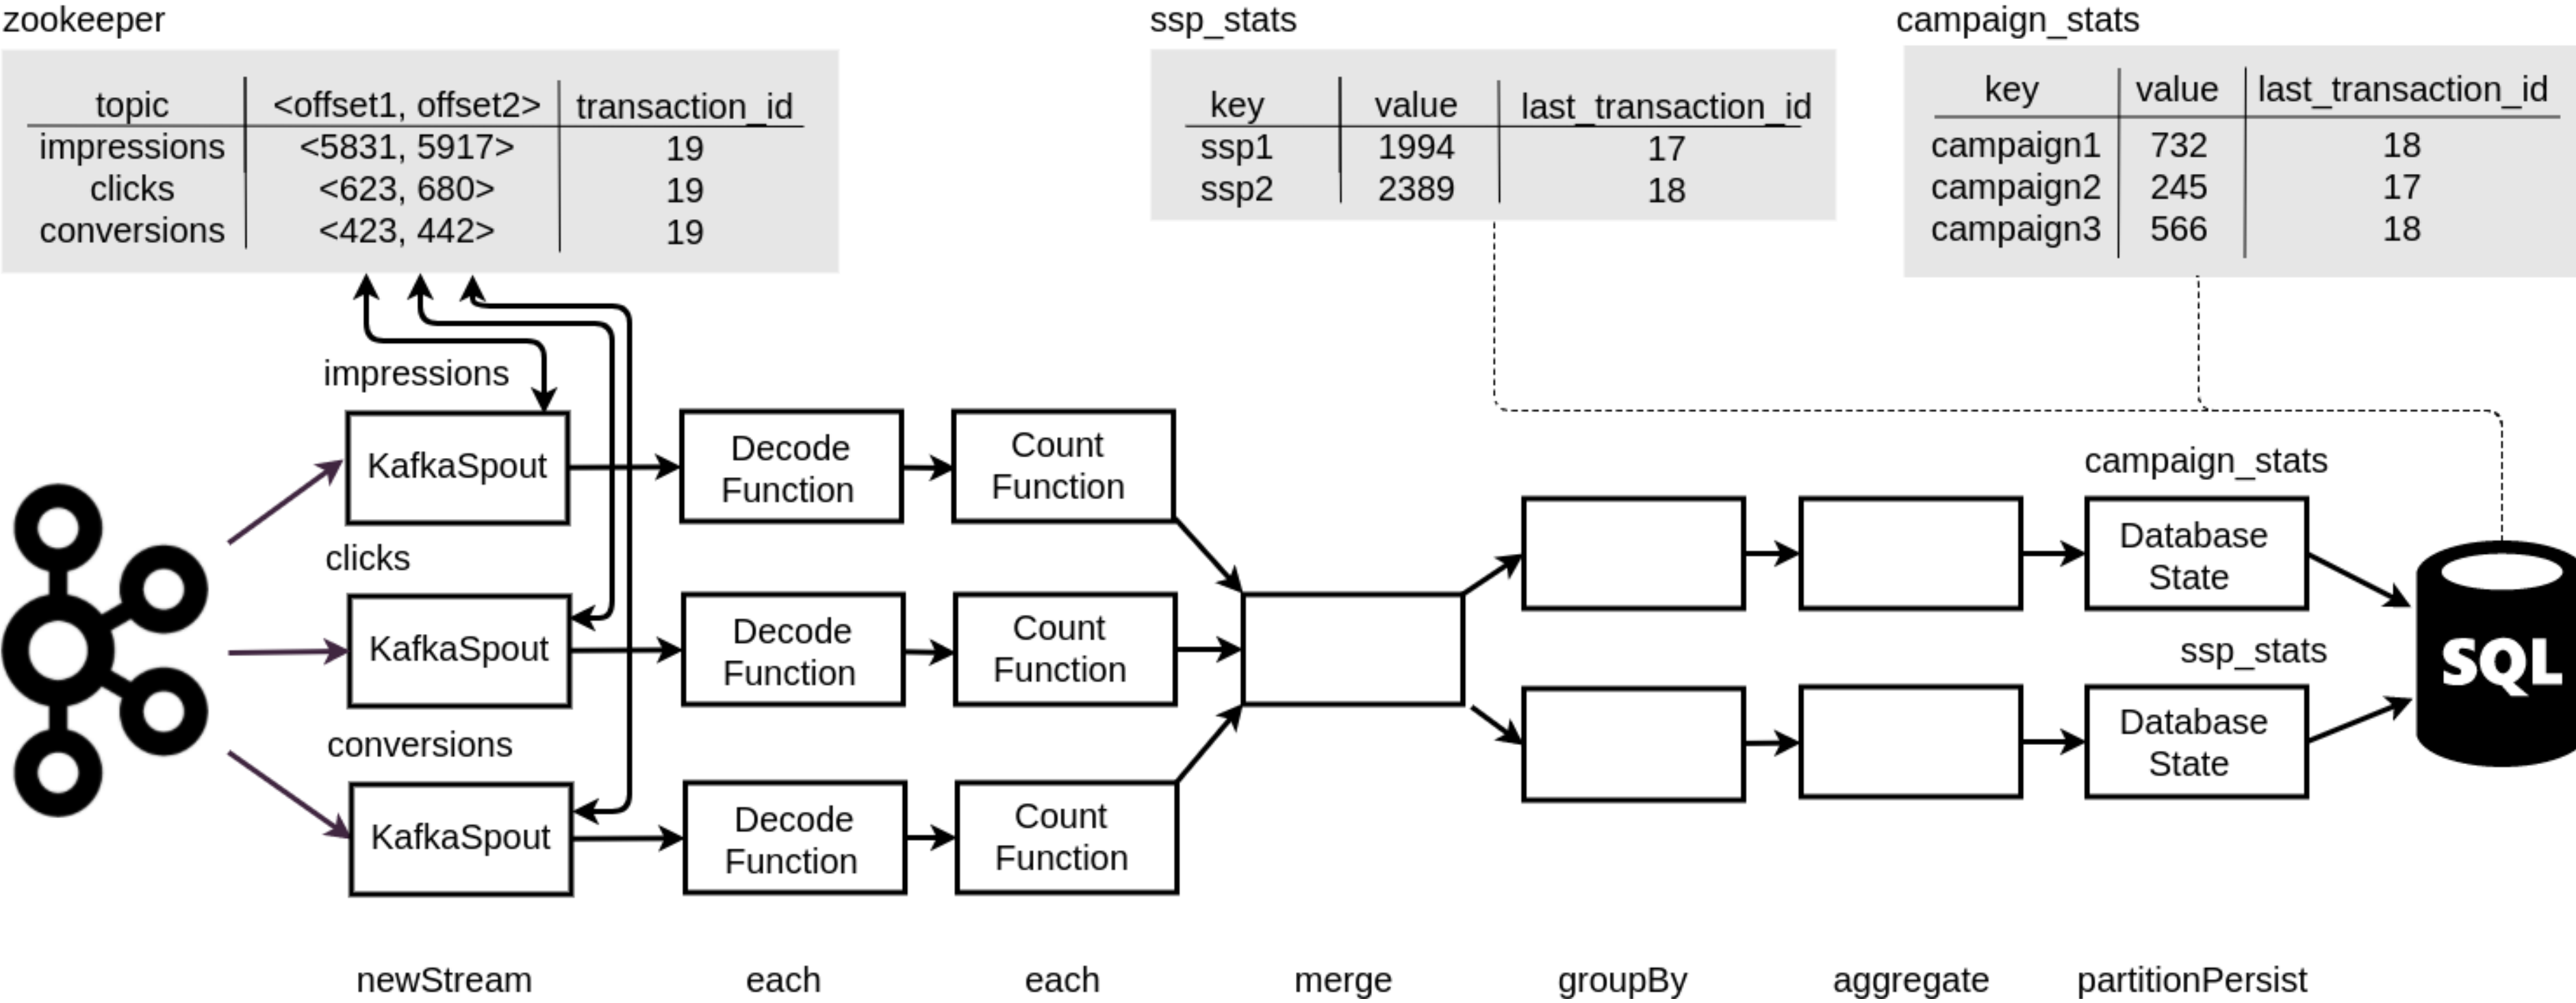


# OUR USE CASES

# STATS-COUNTER: STORM TOPOLOGY (THE 2ND ITERATION)

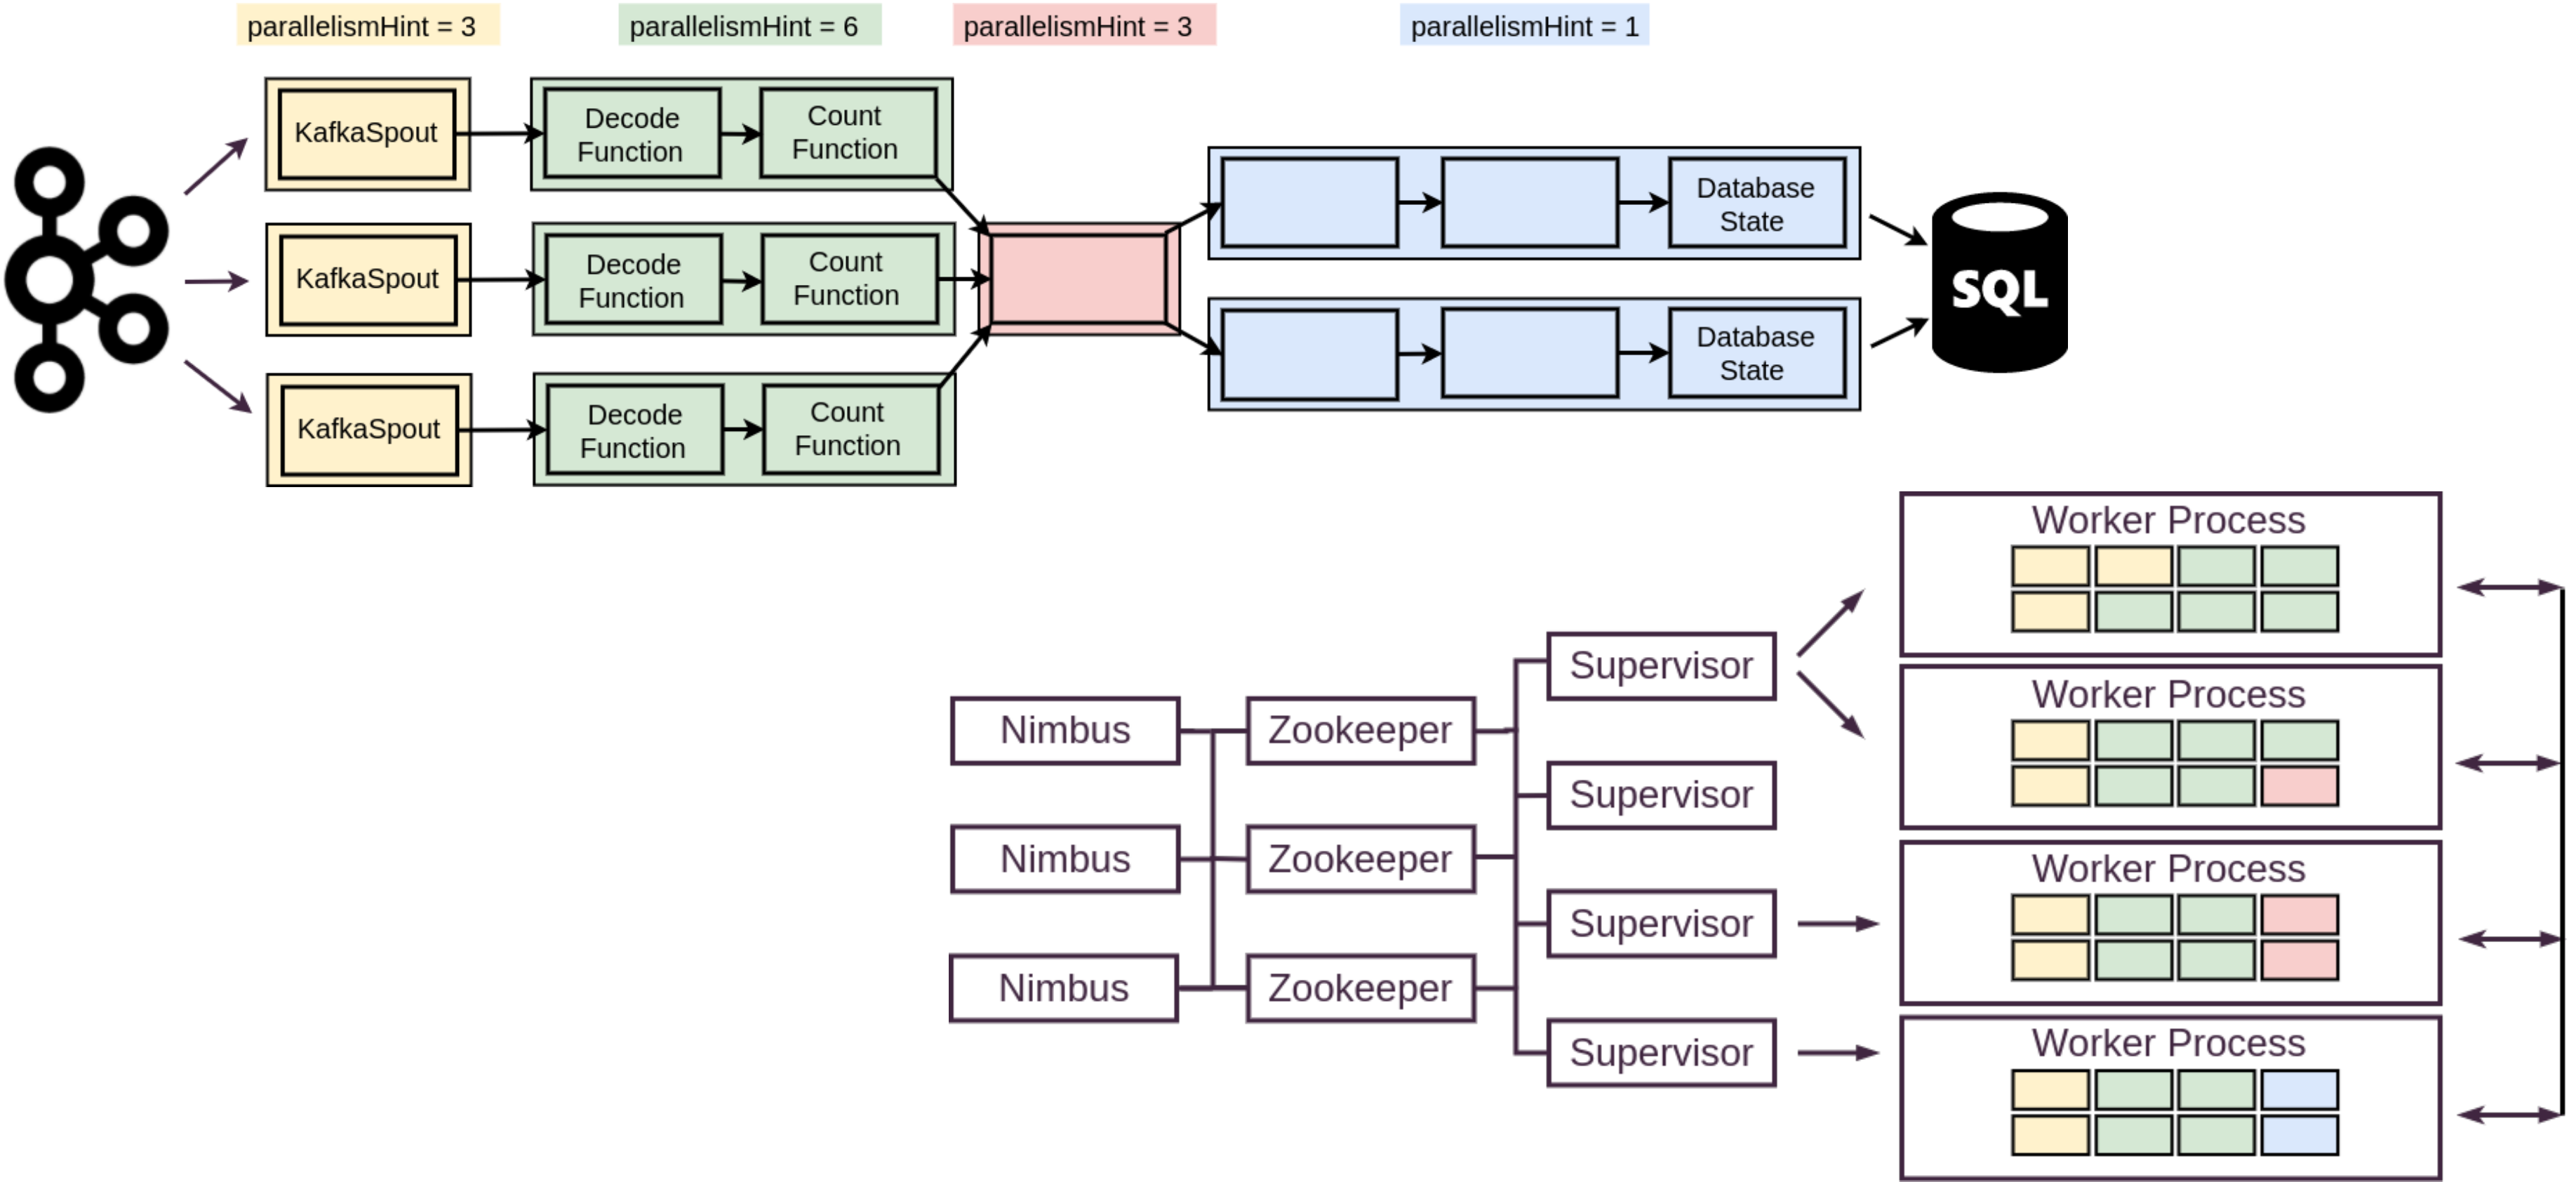


# APACHE STORM: TRIDENT + EXACTLY-ONCE STATE

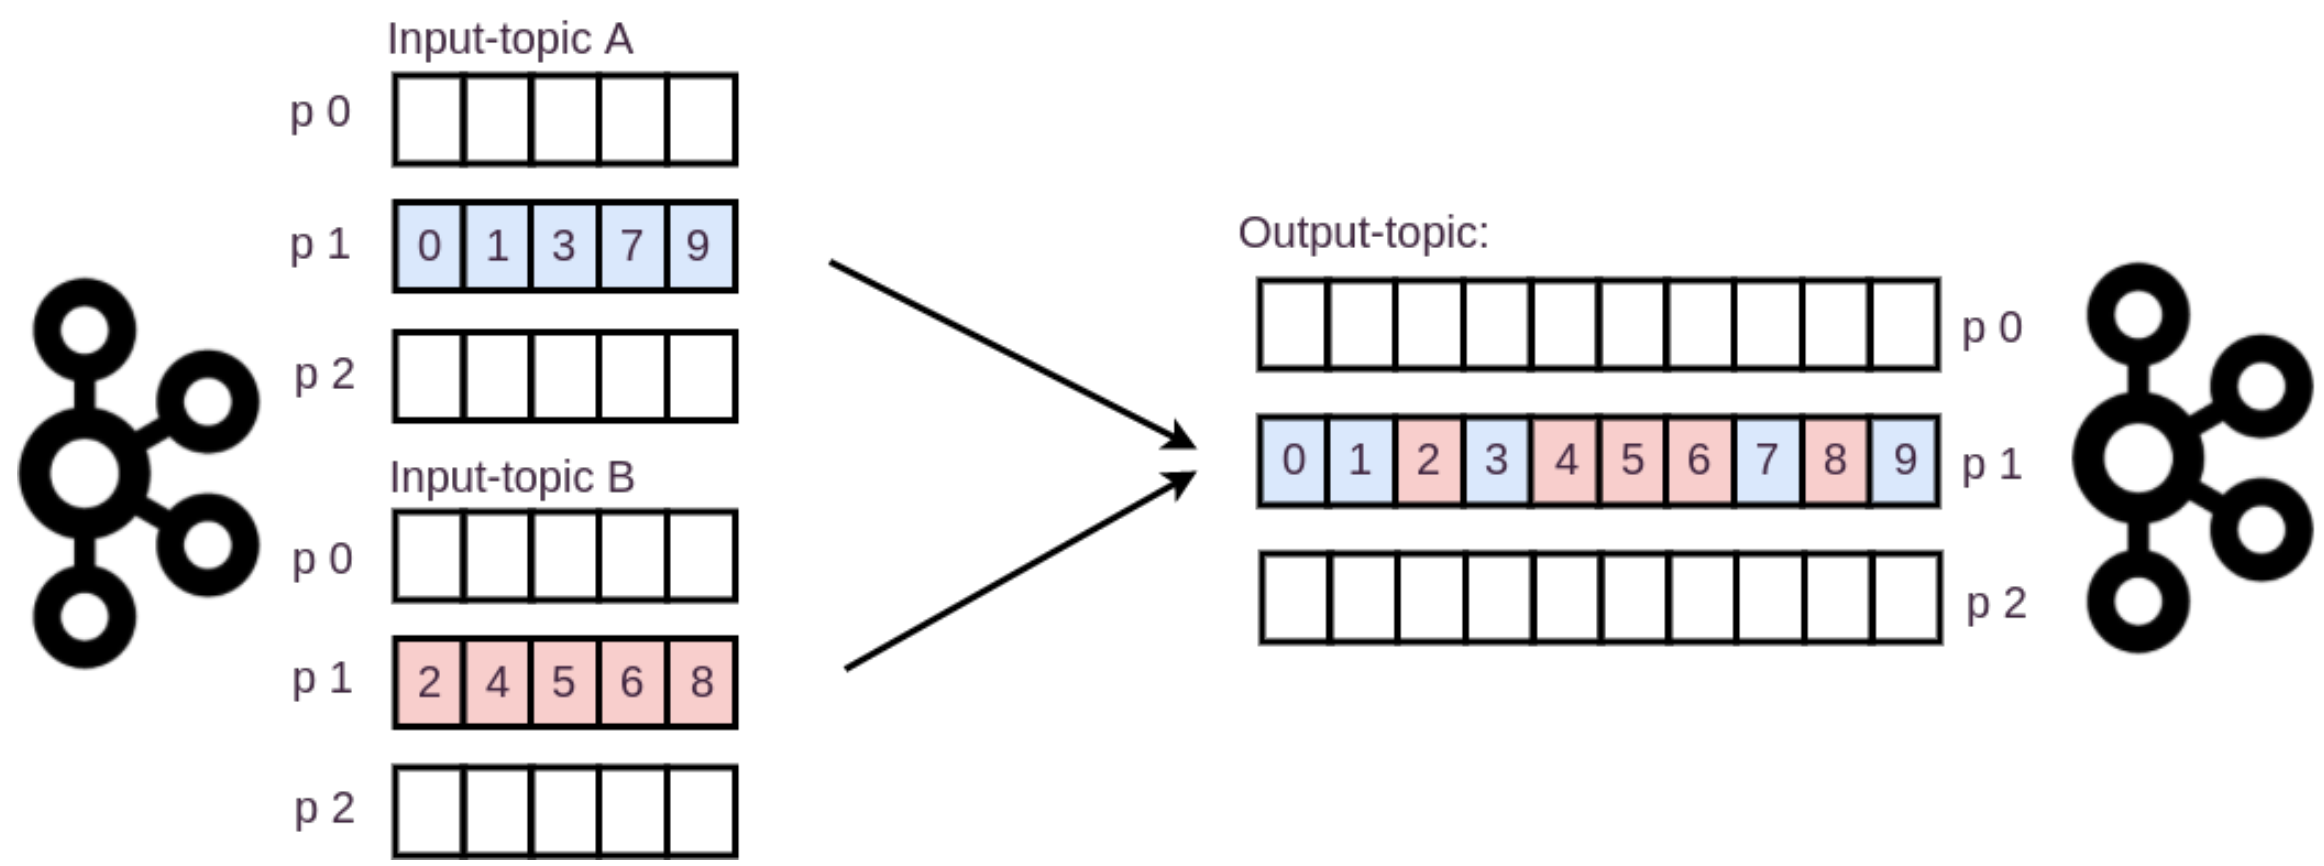




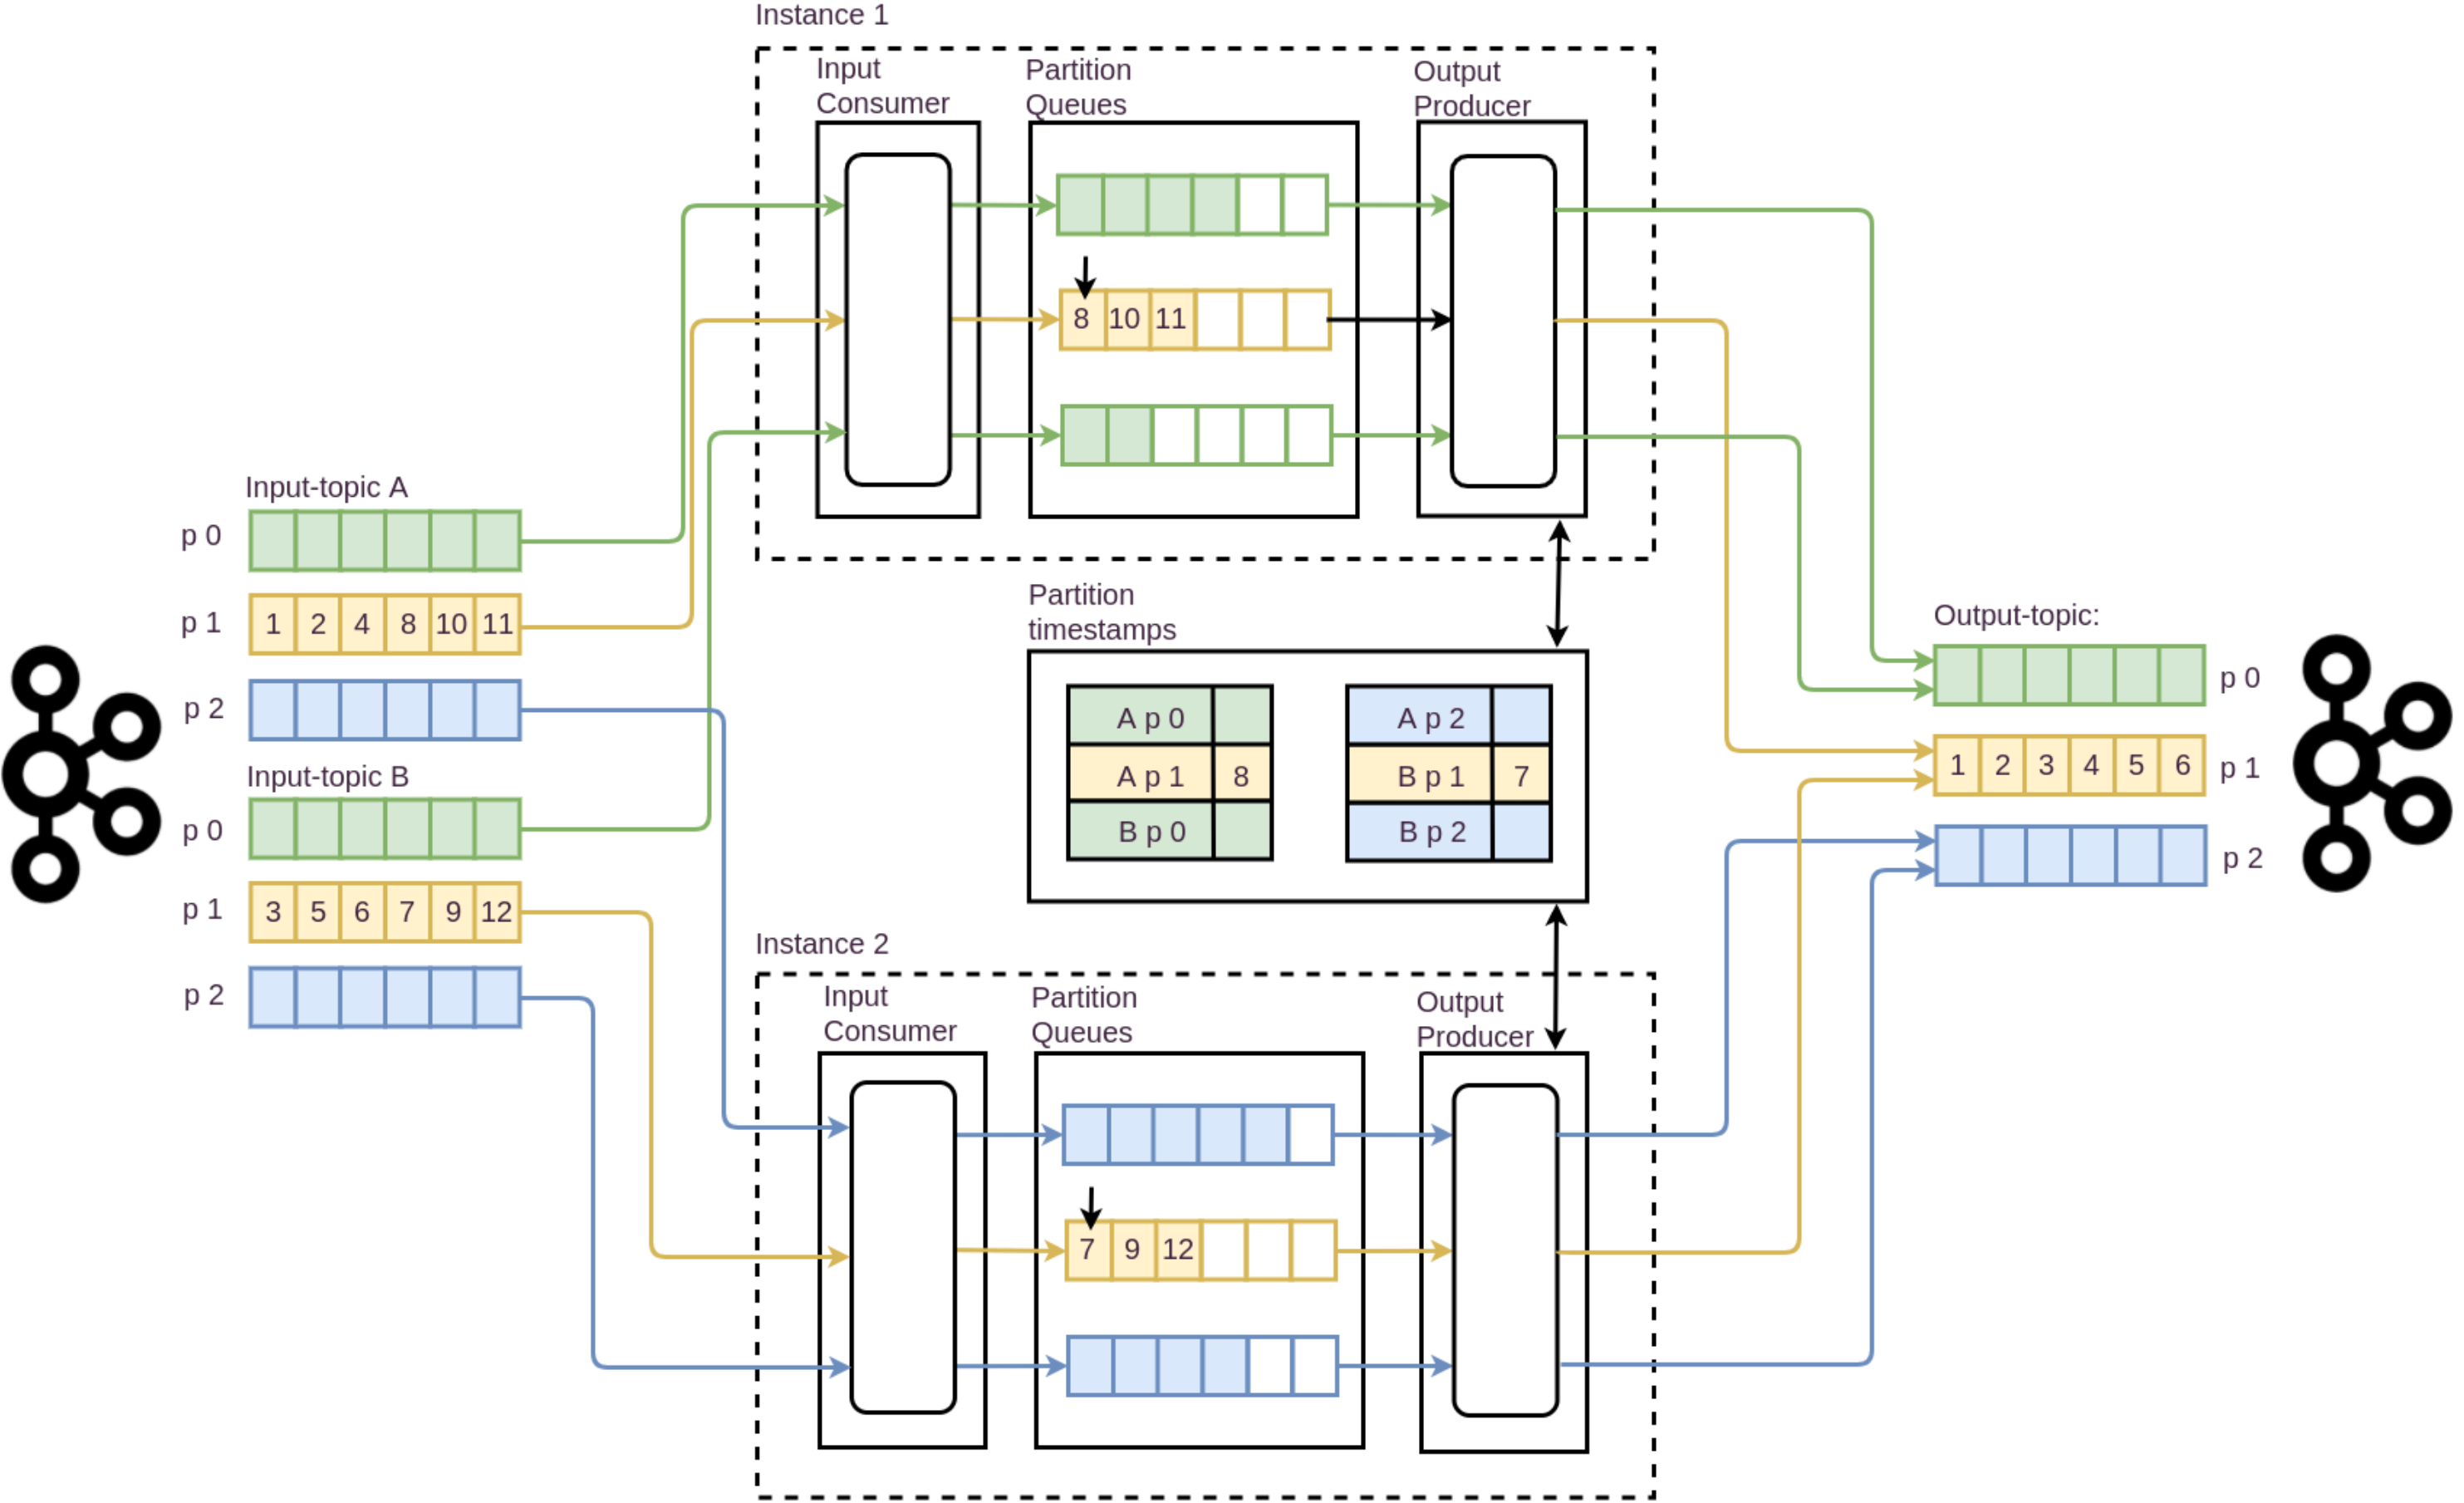
# APACHE STORM: PARALLELISM MODEL



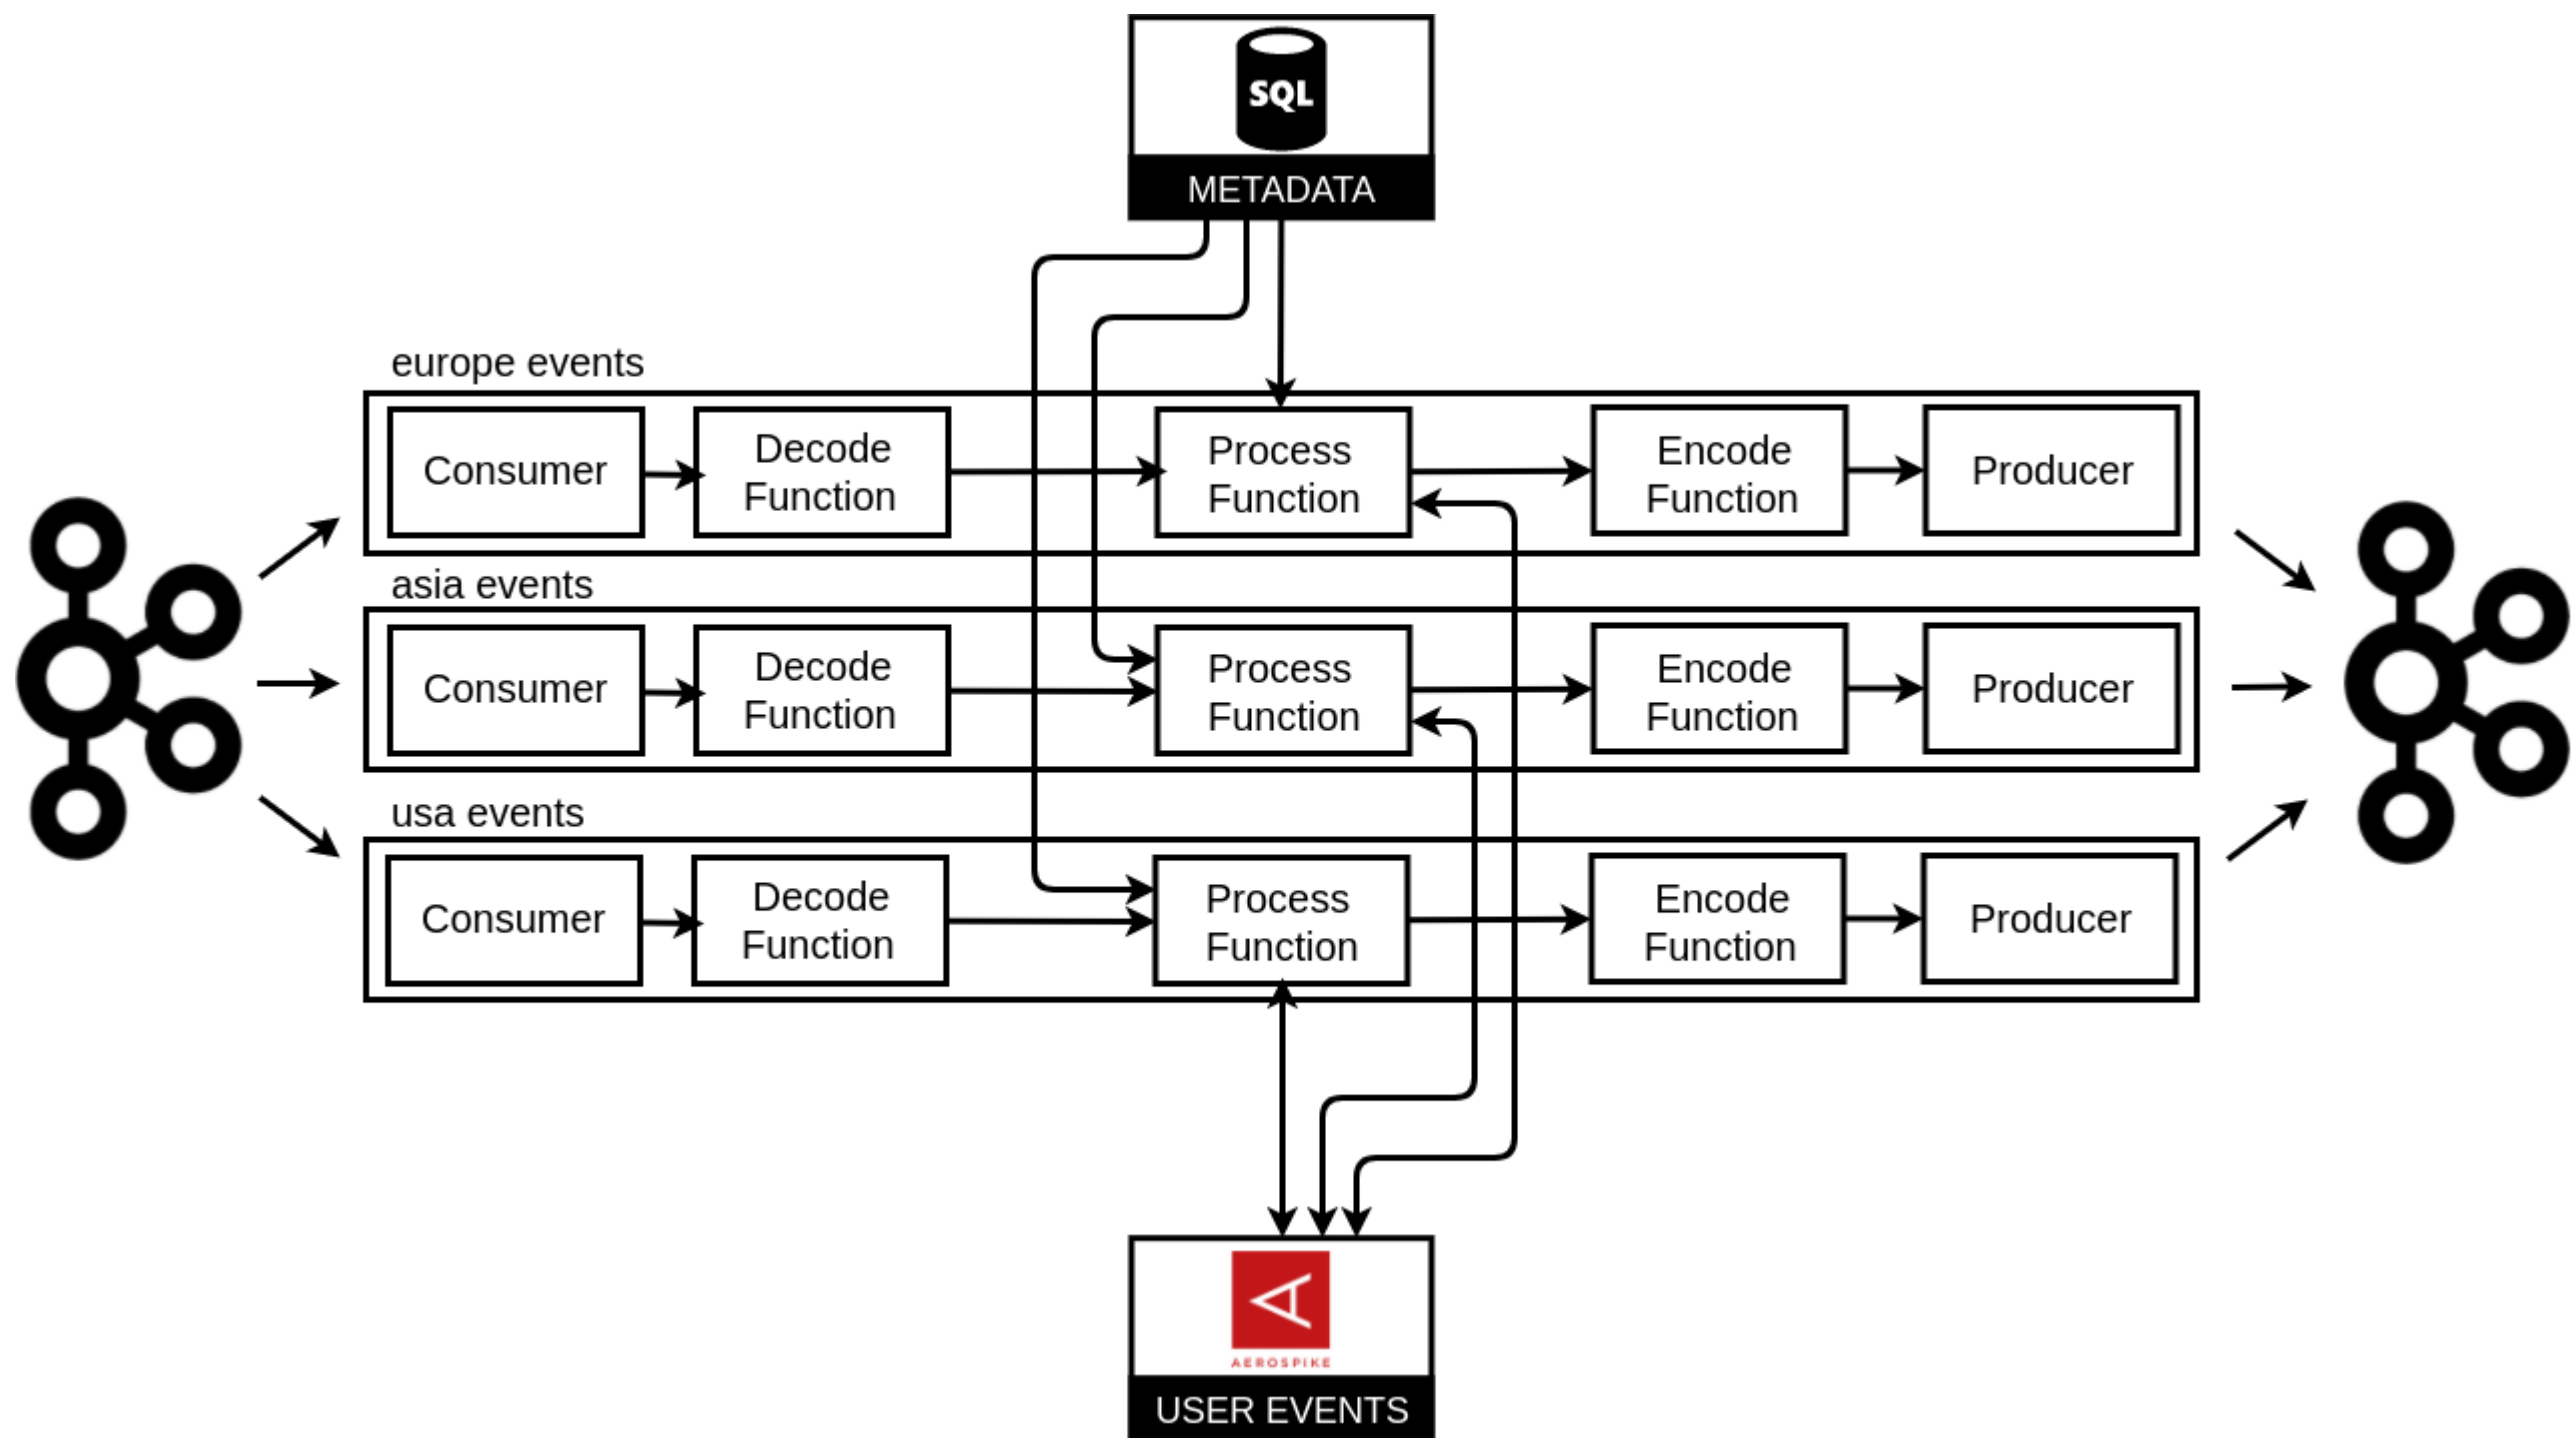
# MERGER (THE 4TH ITERATION)



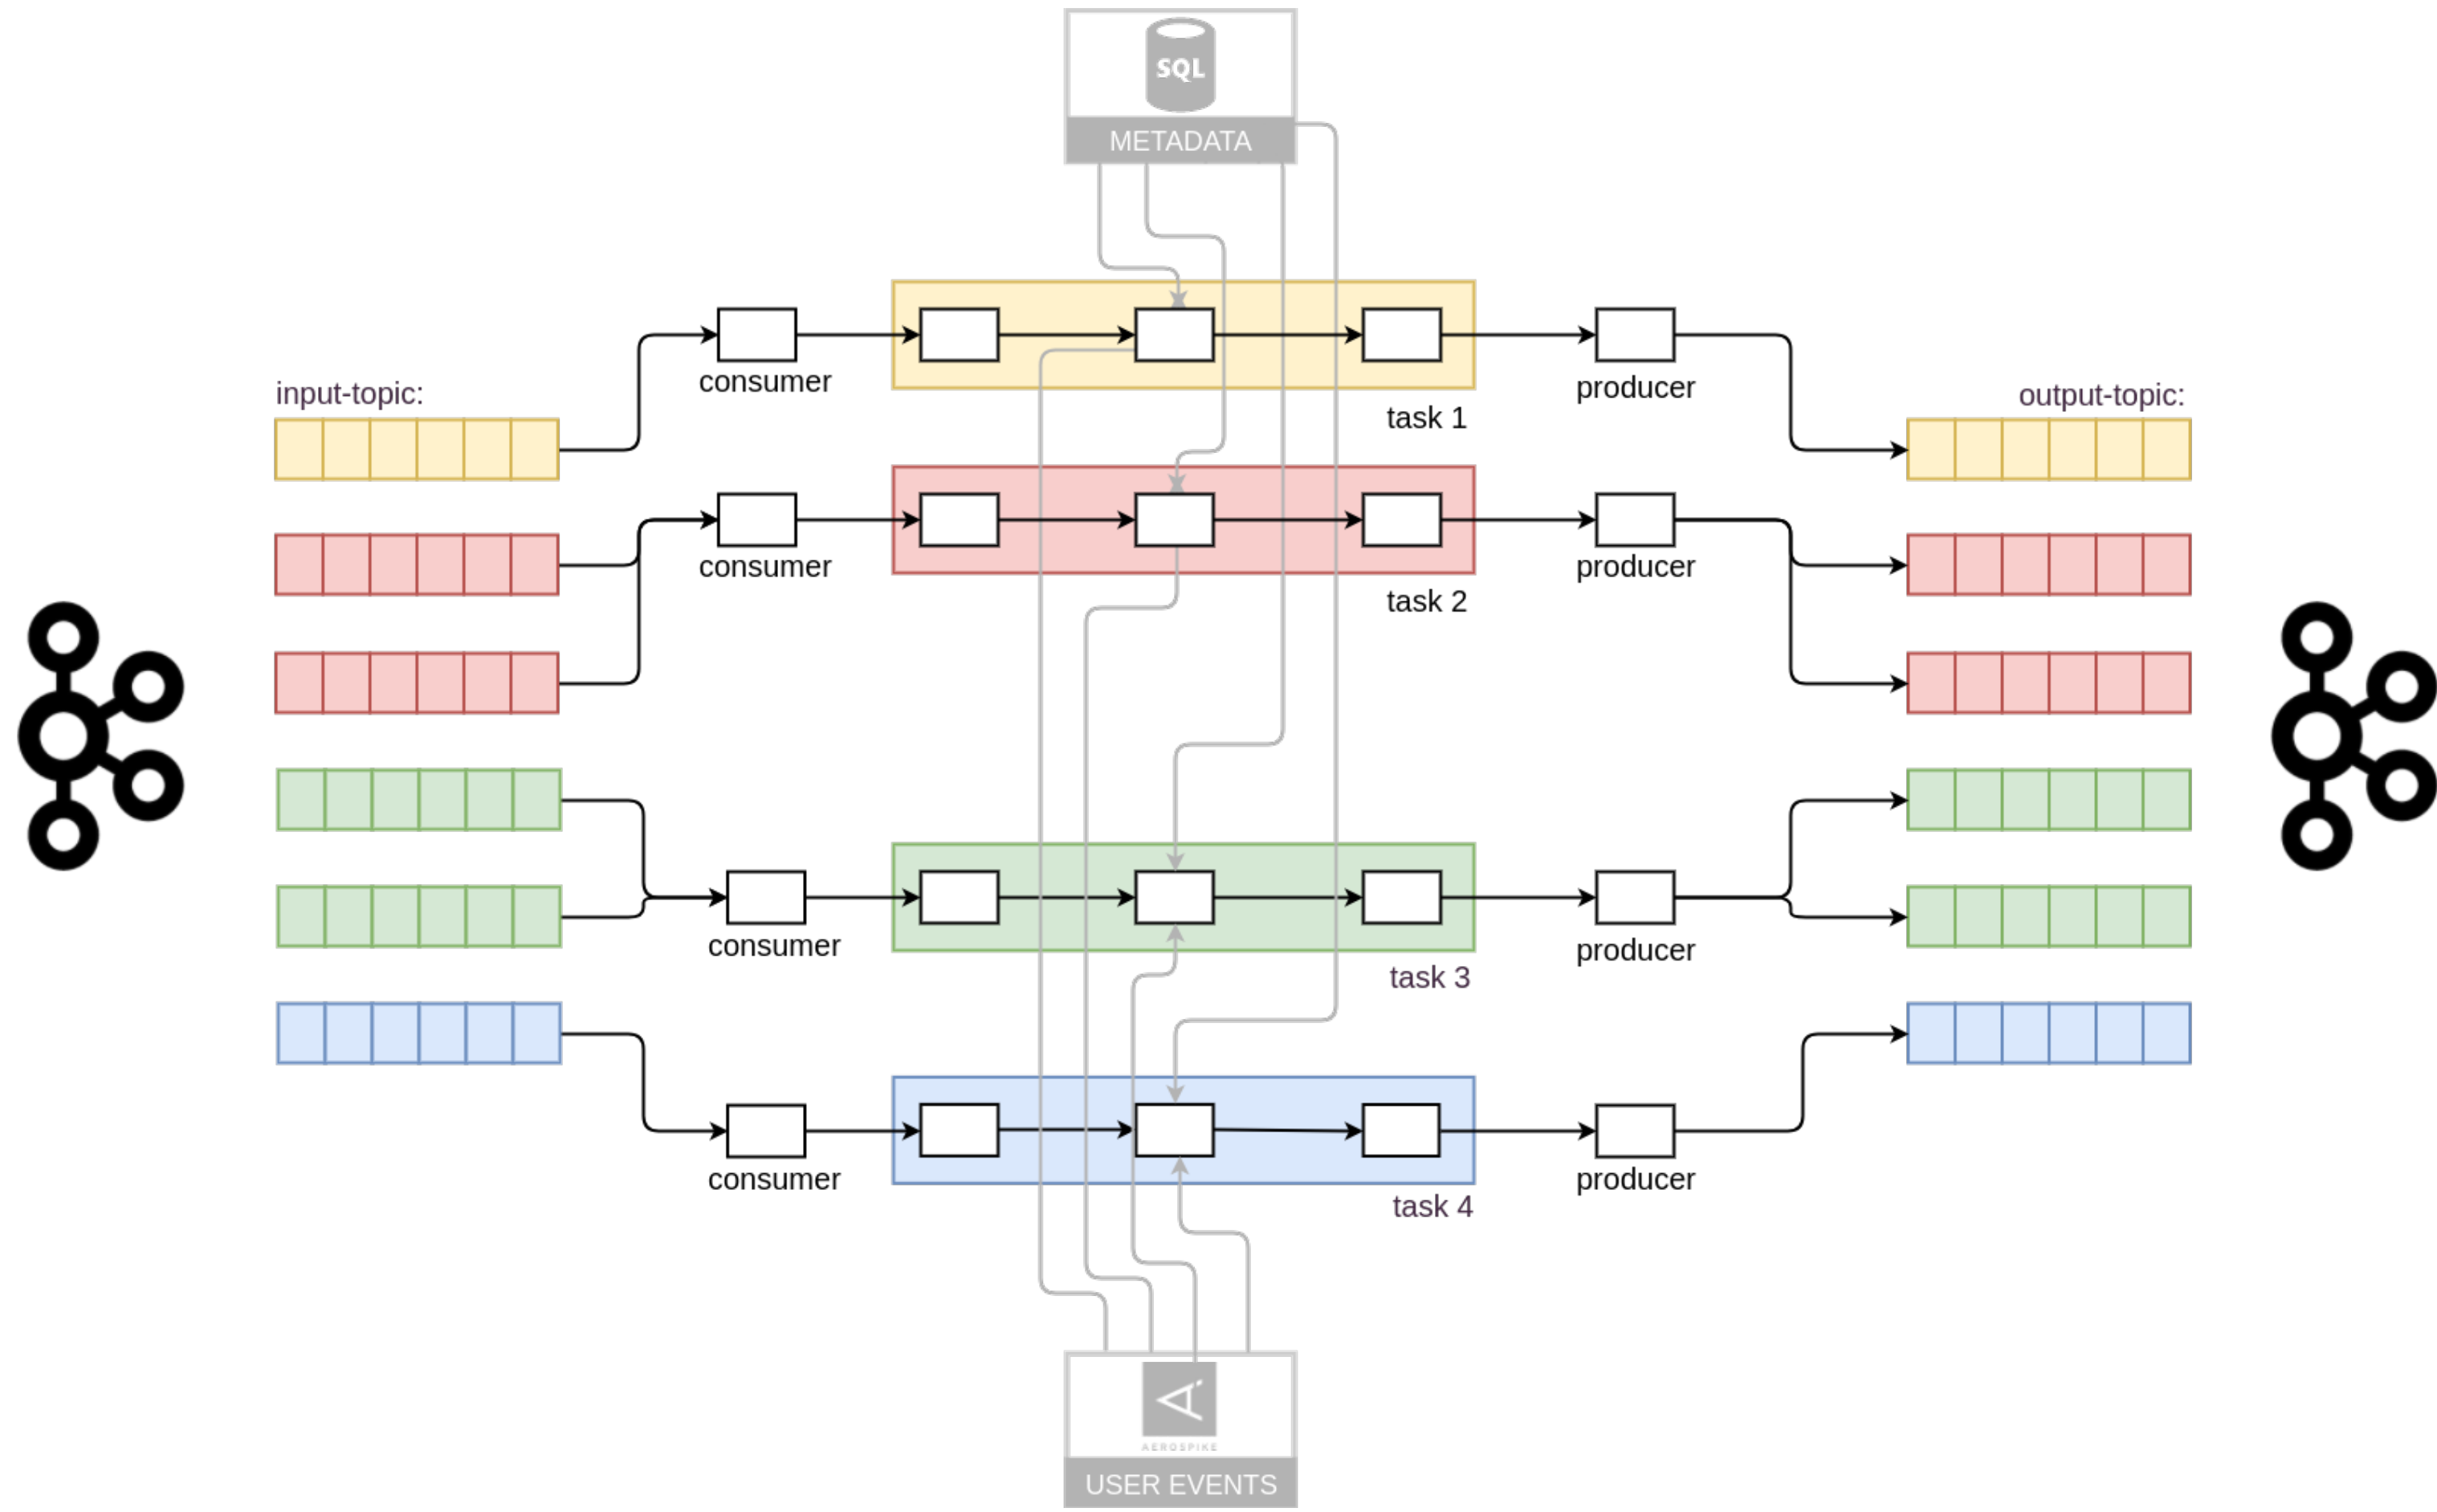
# MERGER: KAFKA CONSUMER API



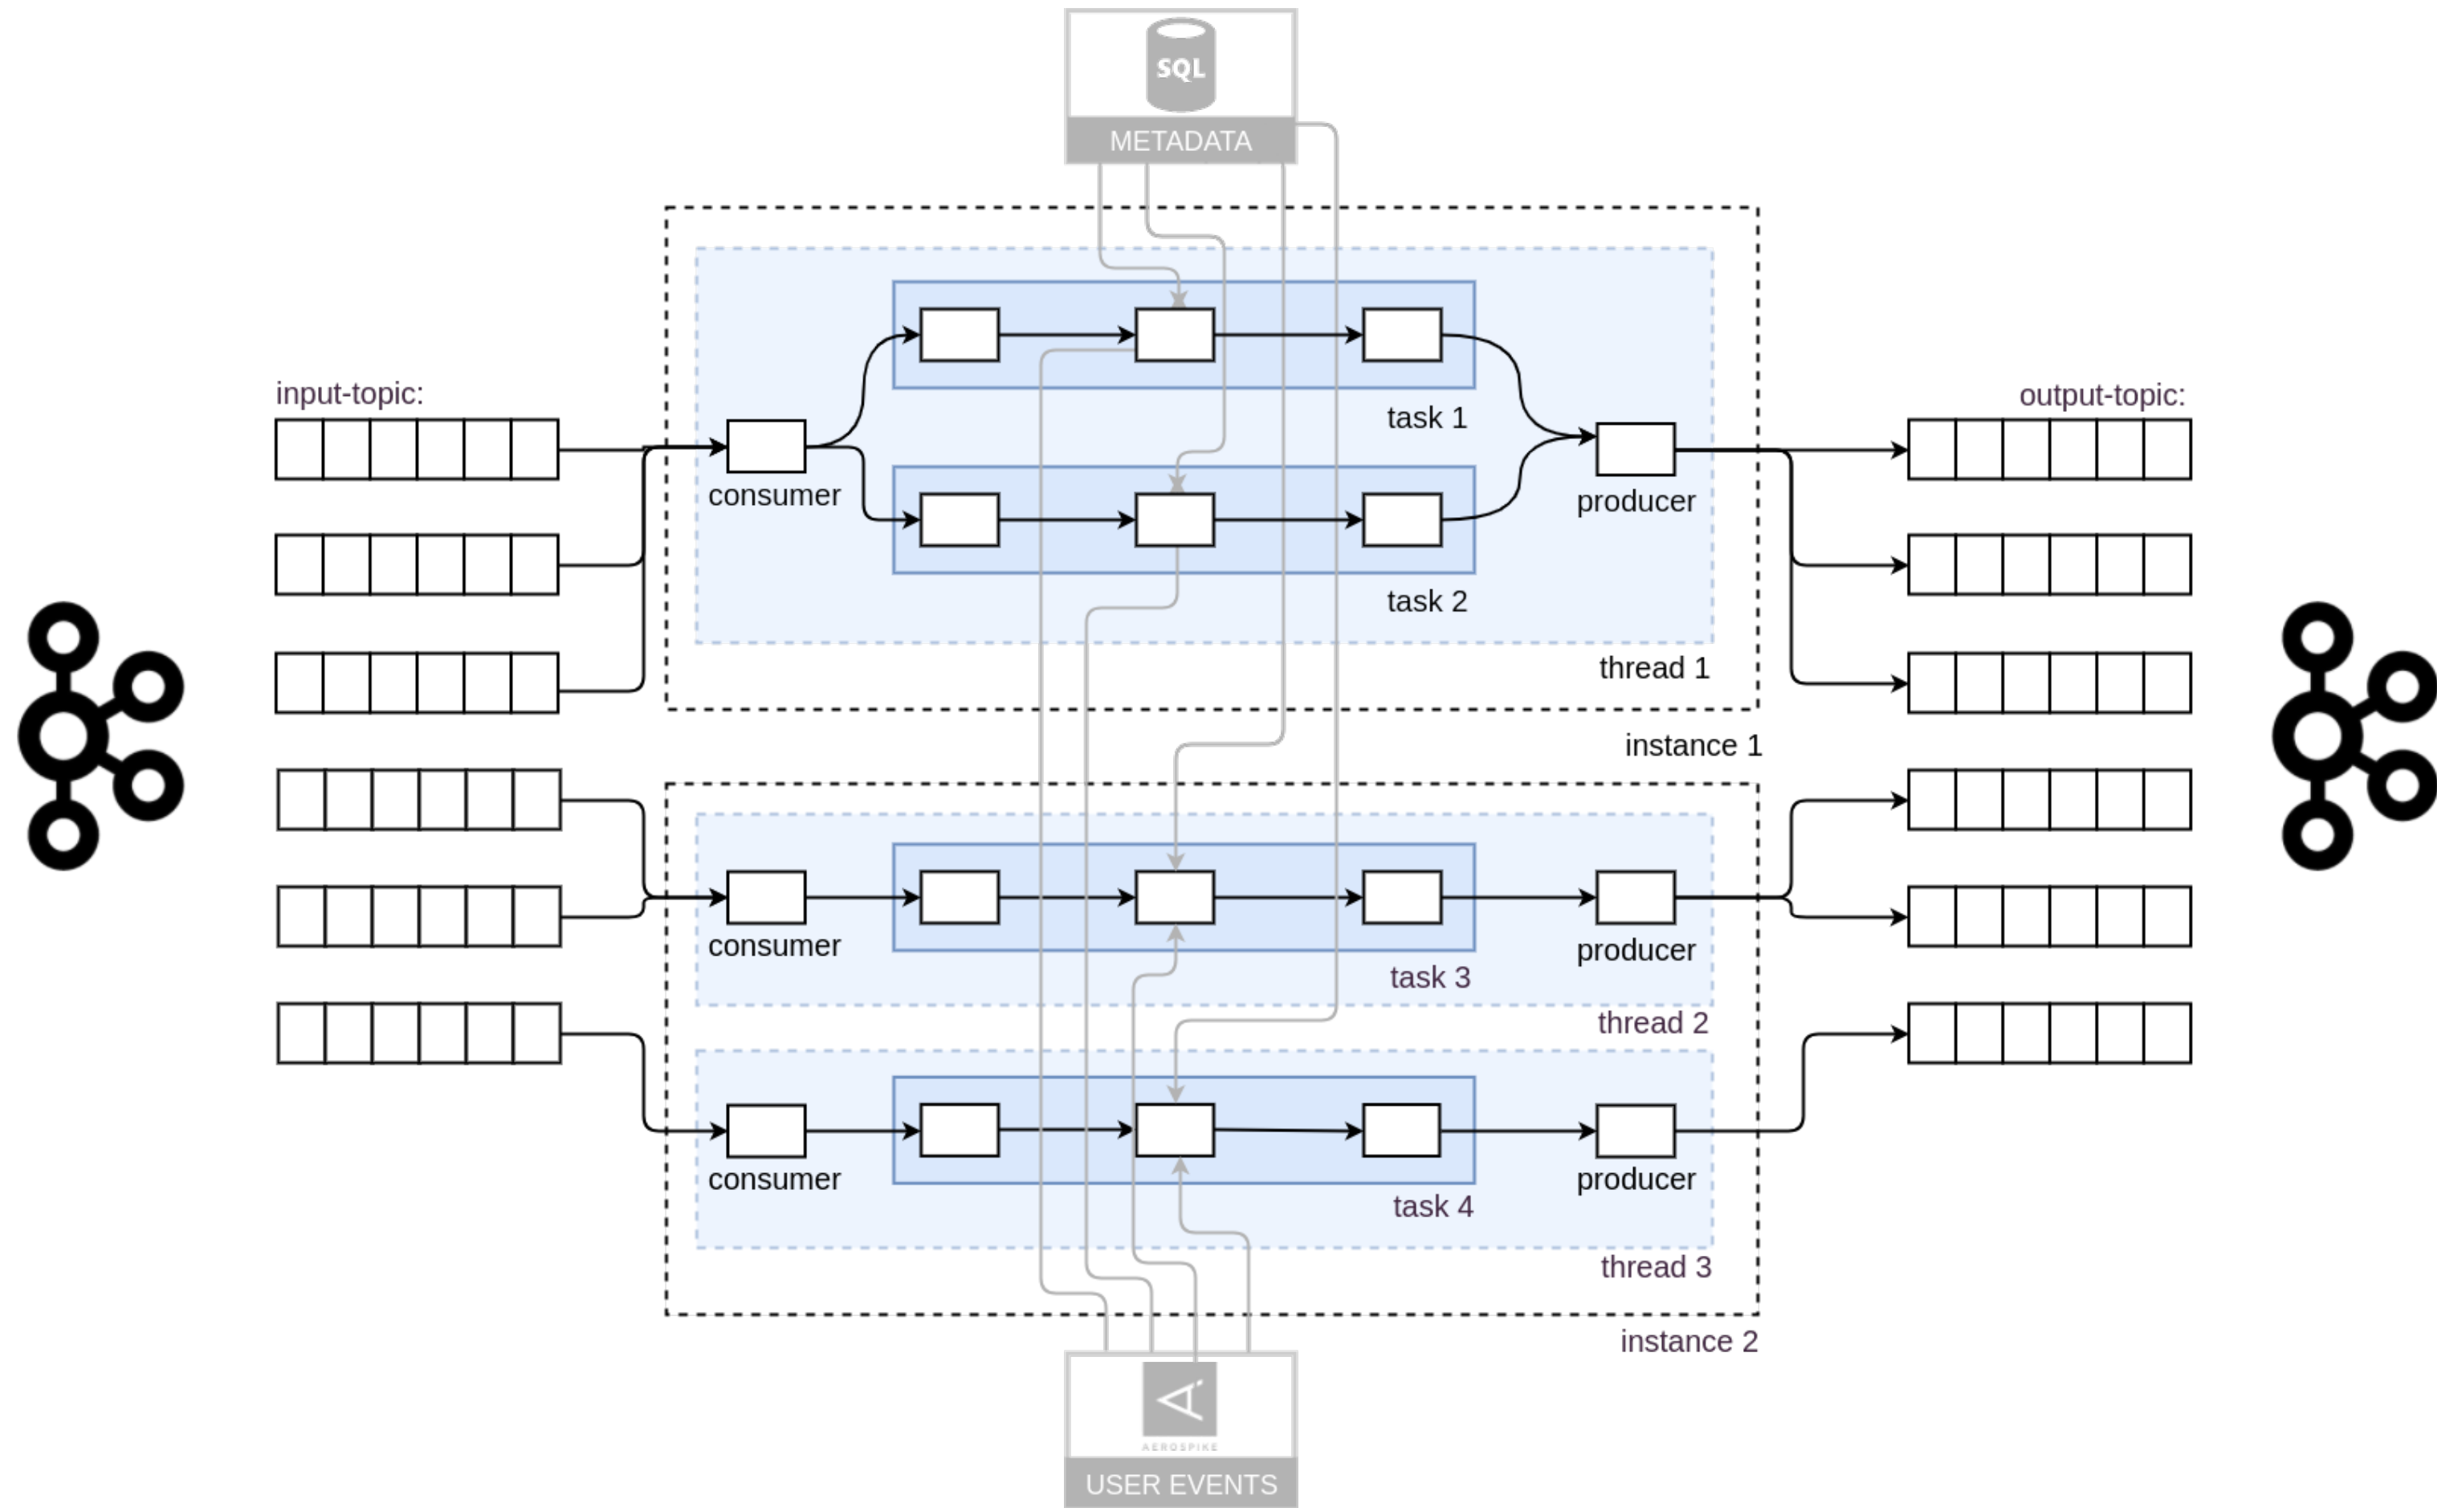
DATA-FLOW: **KAFKA STREAMS (THE 4TH ITERATION)**



# KAFKA STREAMS: **PARALLELISM MODEL**



# KAFKA STREAMS: **PARALLELISM MODEL**



# KAFKA STREAMS: **EXACTLY-ONCE DELIVERY**

Kafka Streams:

- `processing.guarantee = exactly-once`

# KAFKA STREAMS: **EXACTLY-ONCE DELIVERY**

Kafka Streams:

- `processing.guarantee = exactly-once`

Producer:

- `transactions`
- `enable.idempotence = true`



# KAFKA STREAMS: **EXACTLY-ONCE DELIVERY**

Kafka Streams:

- `processing.guarantee = exactly-once`

Producer:

- `transactions`
- `enable.idempotence = true`

Consumer:

- `isolation.level = read_committed`

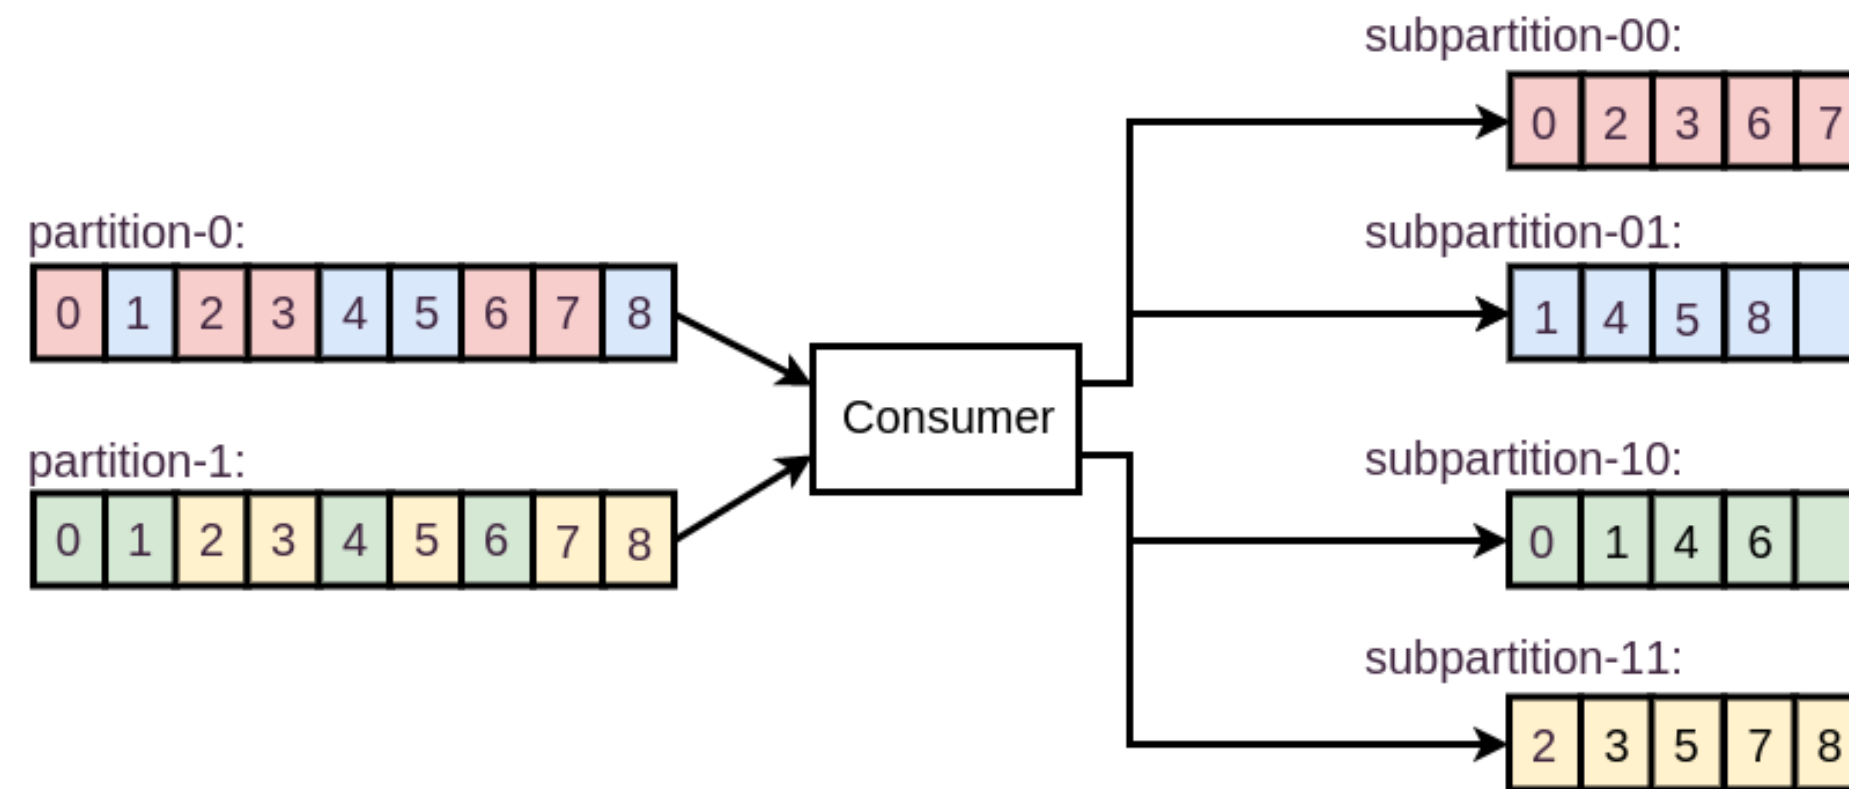
**KAFKA WORKERS**

## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution

# KAFKA WORKERS: MAIN FEATURES

- higher level of distribution



## KAFKA WORKERS: MAIN FEATURES

- higher level of distribution

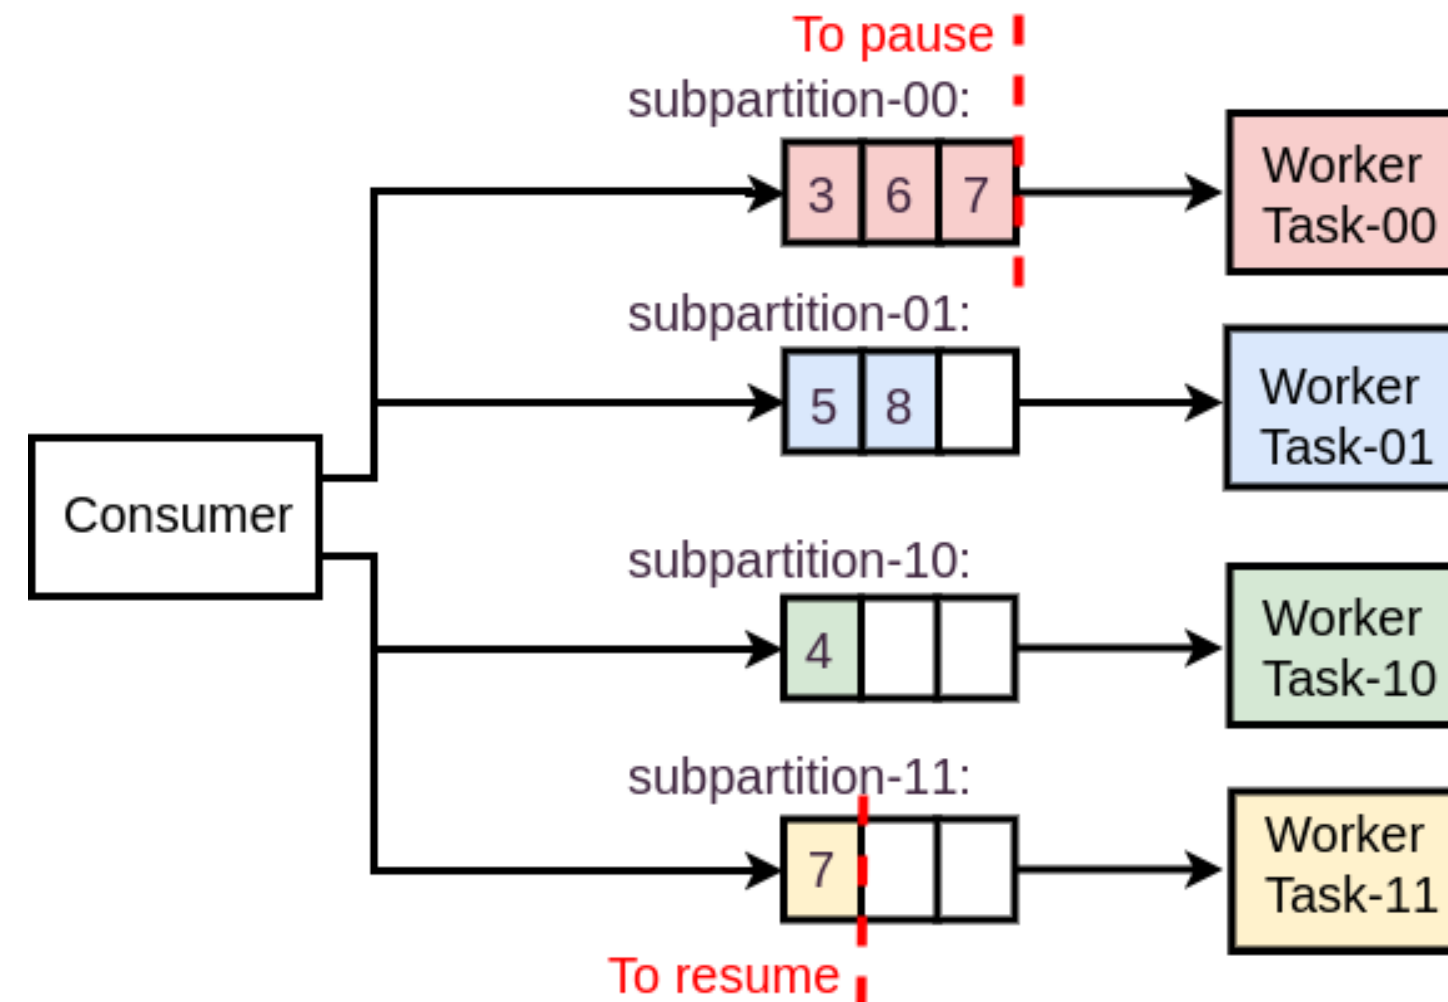
```
public interface WorkerPartitioner<K, V> {  
    int subpartition(ConsumerRecord<K, V> consumerRecord);  
}
```

## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution
- possibility to pause and resume processing for given partition

## KAFKA WORKERS: MAIN FEATURES

- higher level of distribution
- possibility to pause and resume processing for given partition



## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution
- possibility to pause and resume processing for given partition

```
public interface WorkerTask<K, V> {  
    boolean accept(WorkerRecord<K, V> record);  
    void process(WorkerRecord<K, V> record, RecordStatusObserver observer);  
}
```



## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts


## KAFKA WORKERS: MAIN FEATURES

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts

```
public interface RecordStatusObserver {  
    void onSuccess();  
    void onFailure(Exception exception);  
}
```

## KAFKA WORKERS: MAIN FEATURES

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts



## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts
- at-least-once semantics

## KAFKA WORKERS: **MAIN FEATURES**

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts
- at-least-once semantics
- handling failures

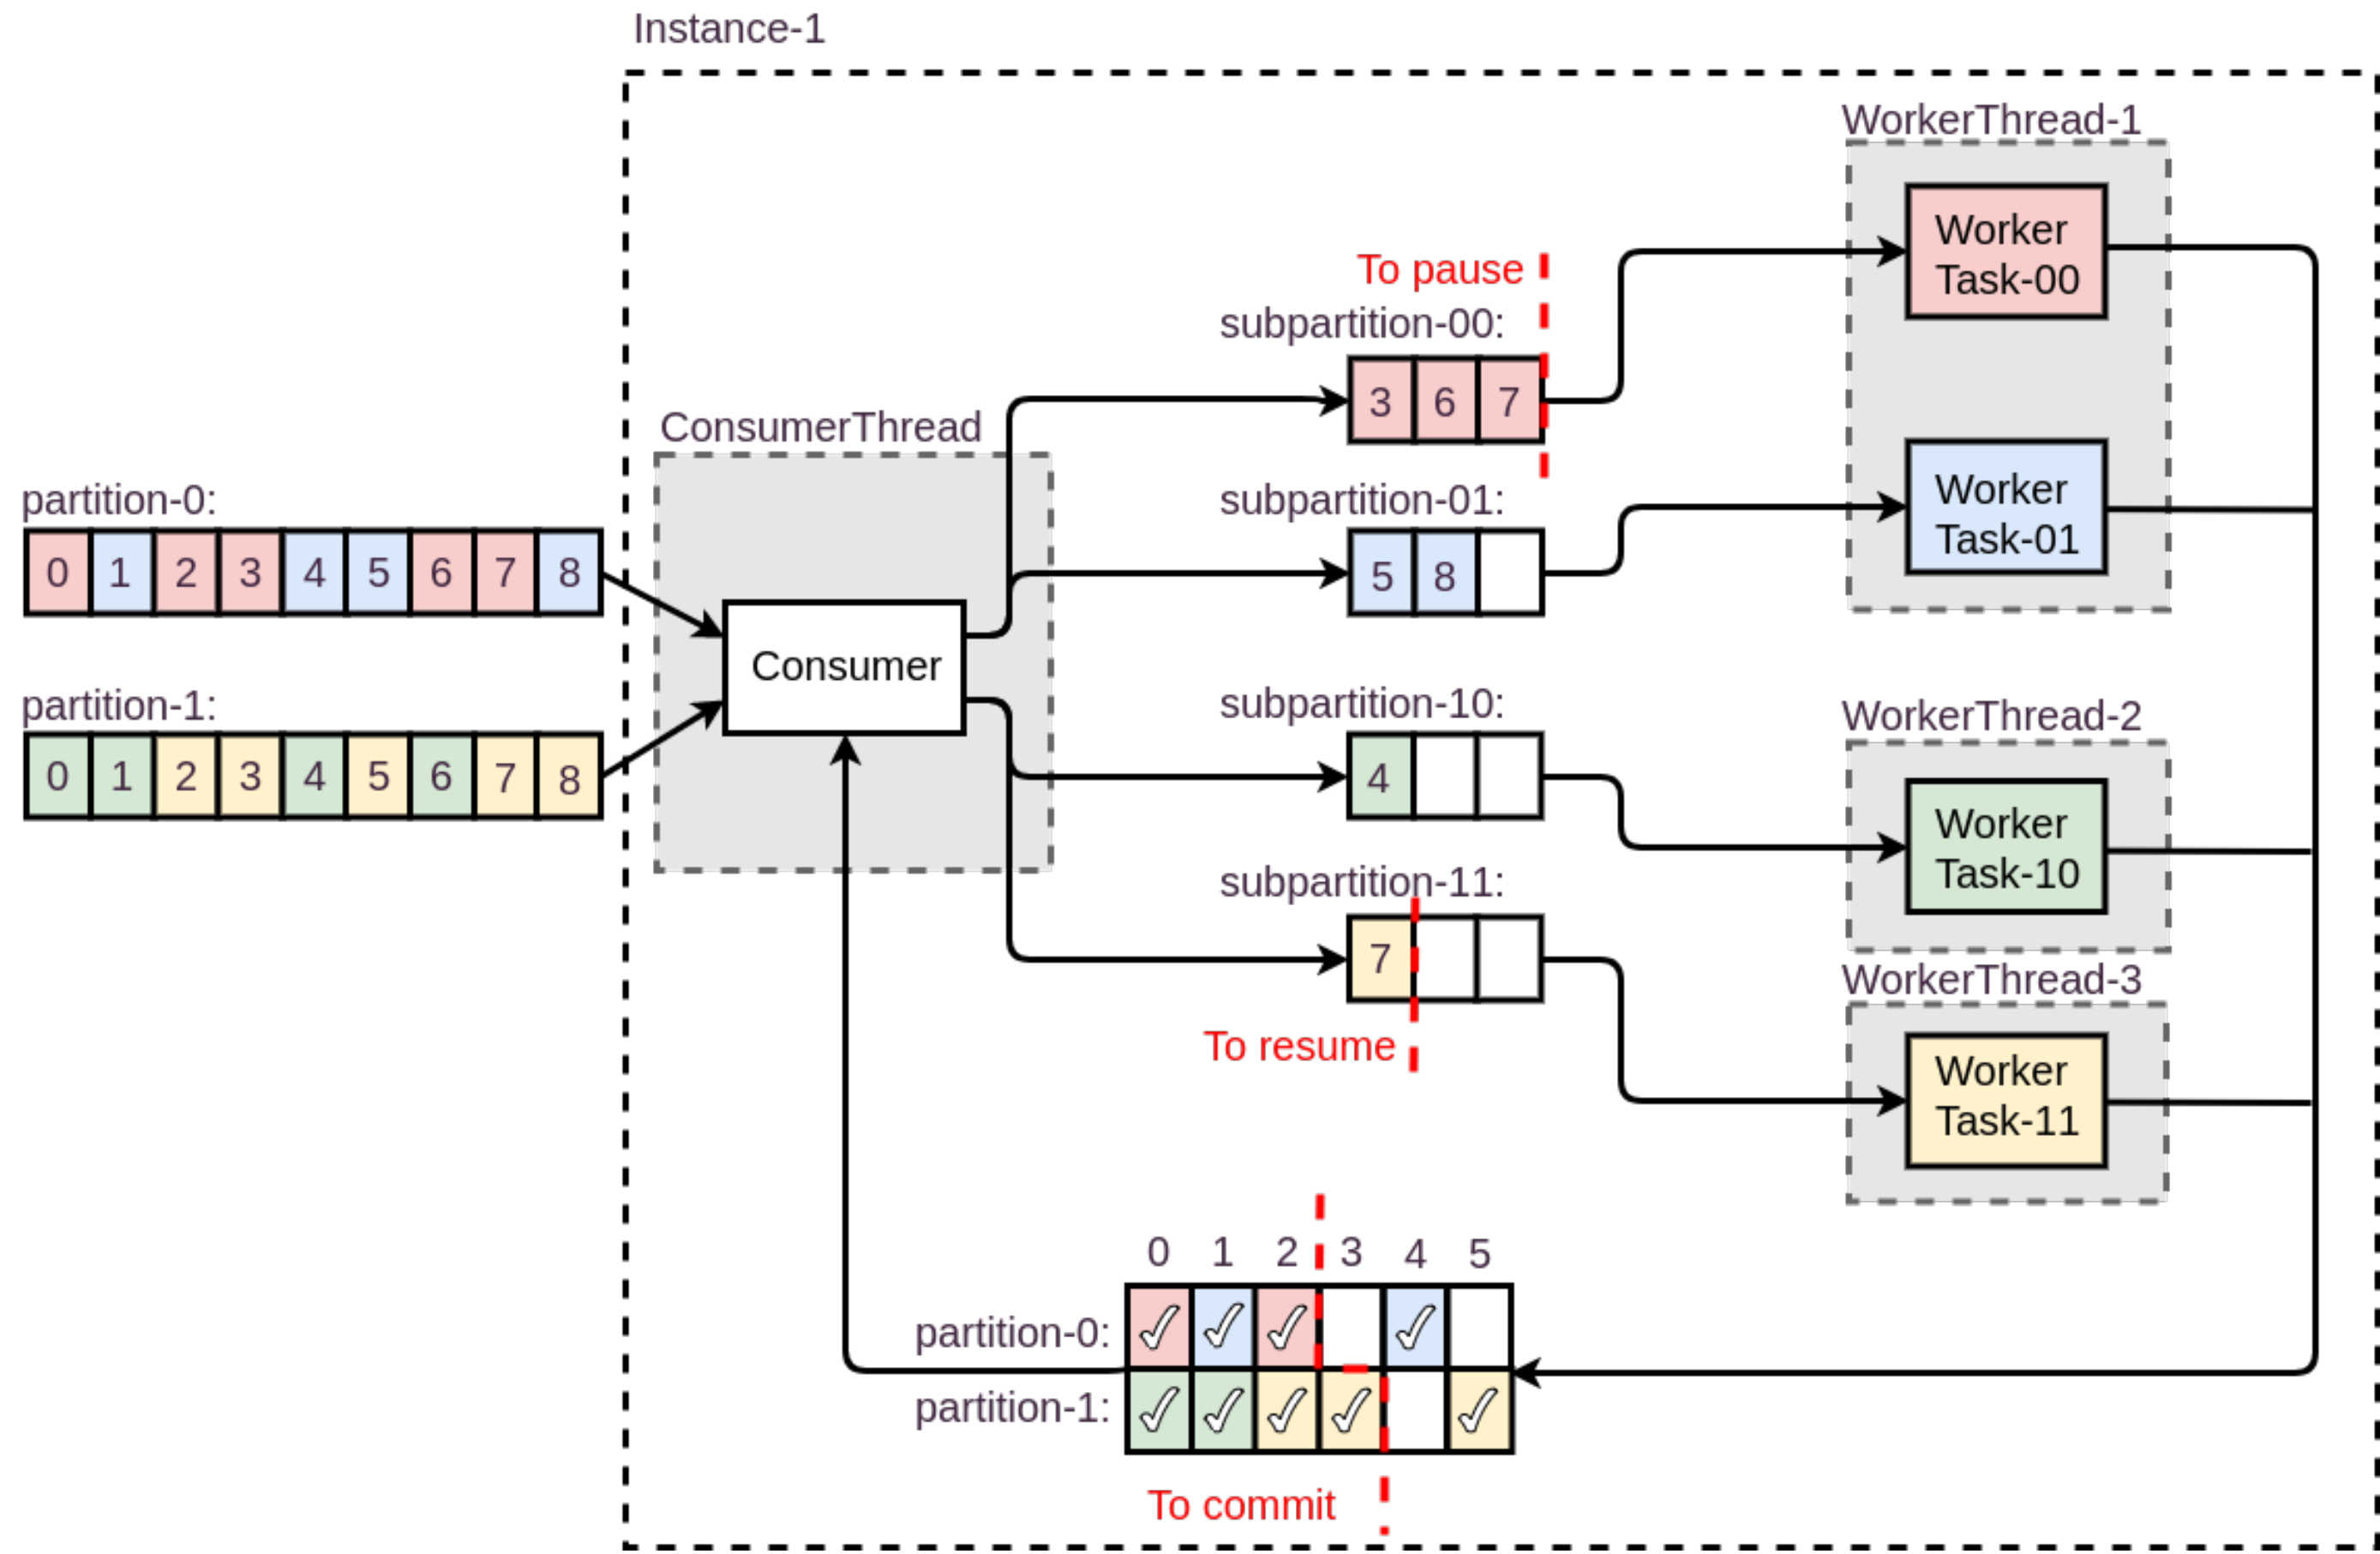
## KAFKA WORKERS: MAIN FEATURES

- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts
- at-least-once semantics
- handling failures
- kafka-to-kafka, hdfs, bigquery, elasticsearch connectors

## KAFKA WORKERS: **MAIN FEATURES**

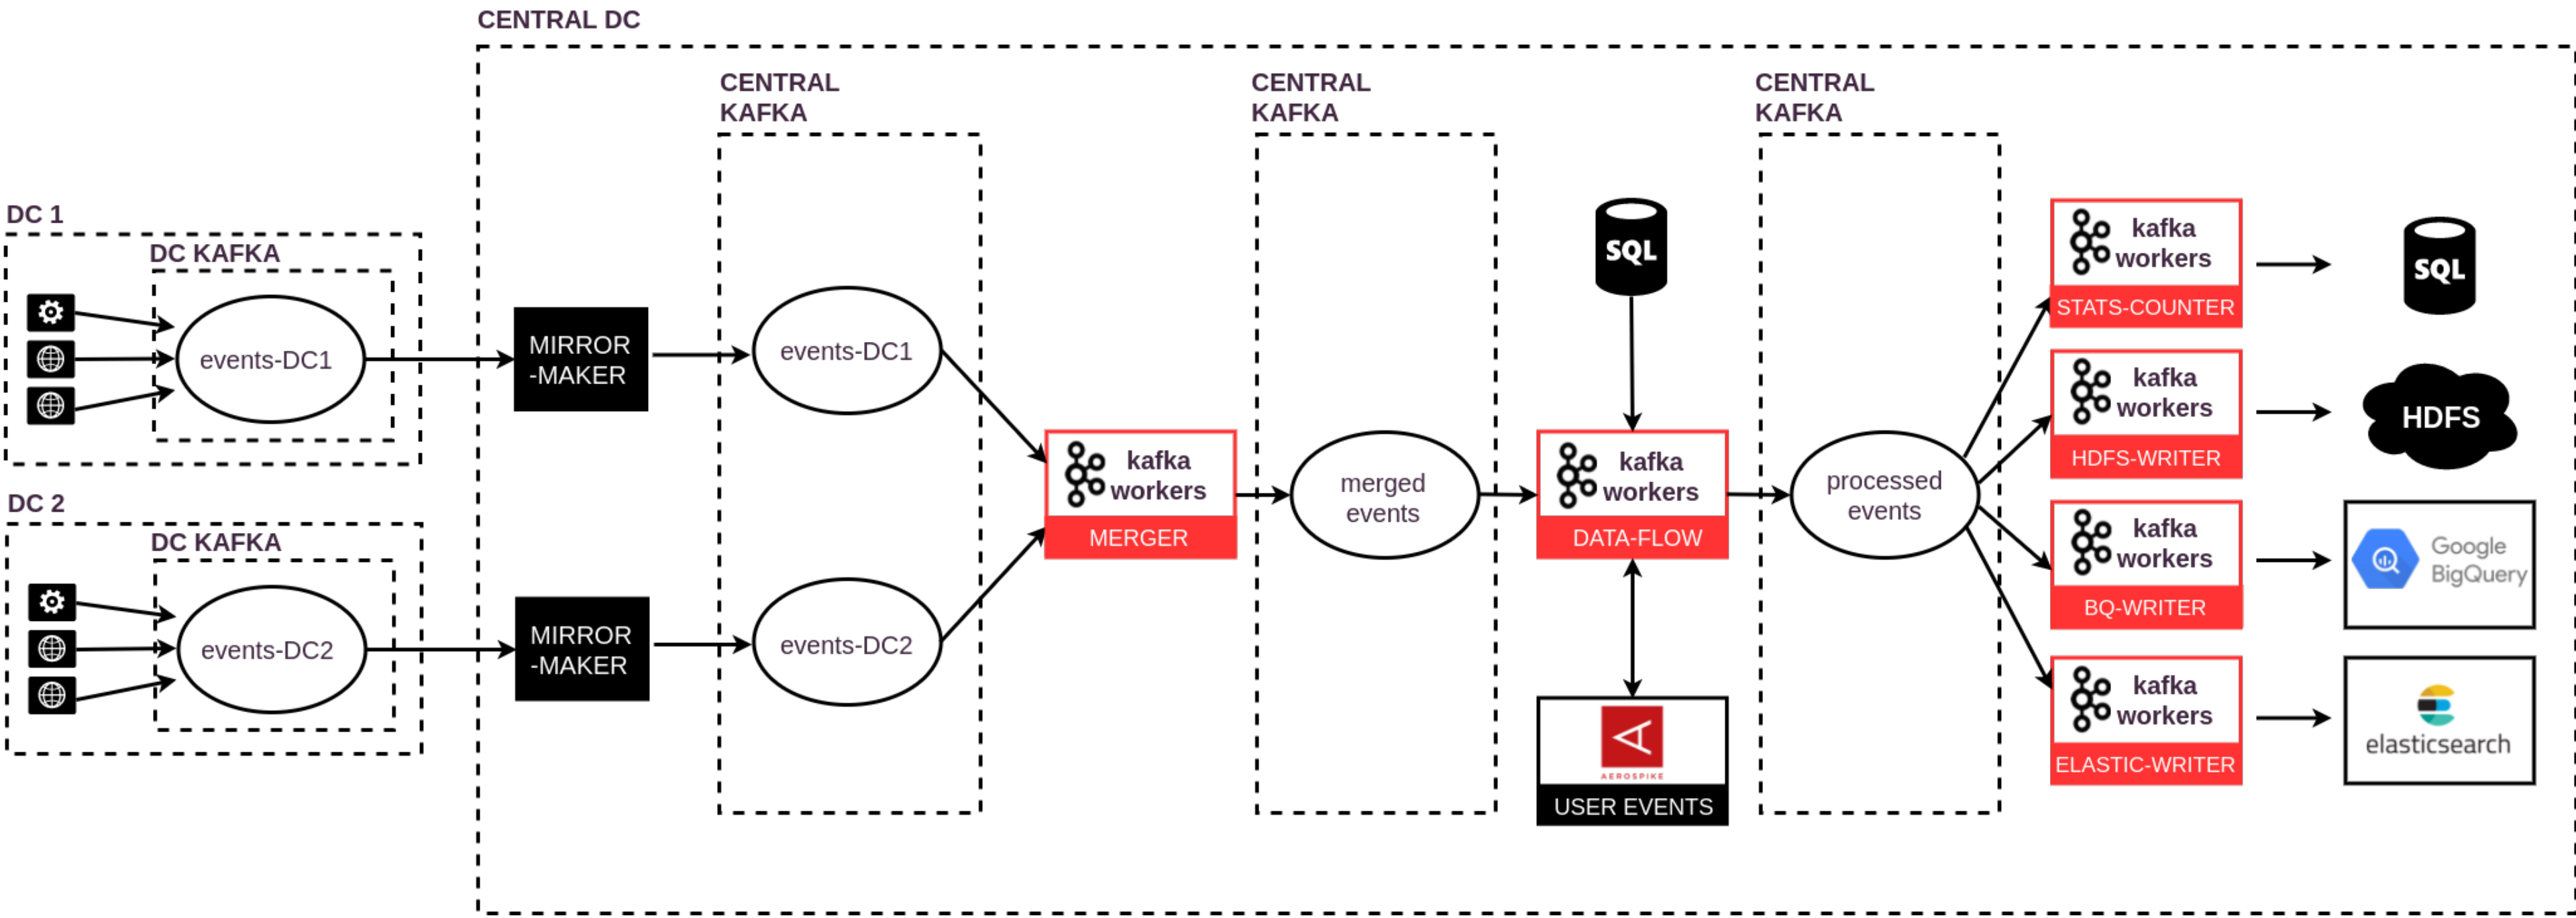
- higher level of distribution
- possibility to pause and resume processing for given partition
- asynchronous processing
  - tighter control of offsets commits
  - backpressure
  - processing timeouts
- at-least-once semantics
- handling failures
- kafka-to-kafka, hdfs, bigquery, elasticsearch connectors
- [github.com/RTBHOUSE/kafka-workers](https://github.com/RTBHOUSE/kafka-workers)

# KAFKA WORKERS: PARALLELISM MODEL





# THE 5TH ITERATION: KAFKA WORKERS



RTBHOUSE =

WE ARE  
HIRING

[techblog.rtbhouse.com/jobs](https://techblog.rtbhouse.com/jobs)